

Spark

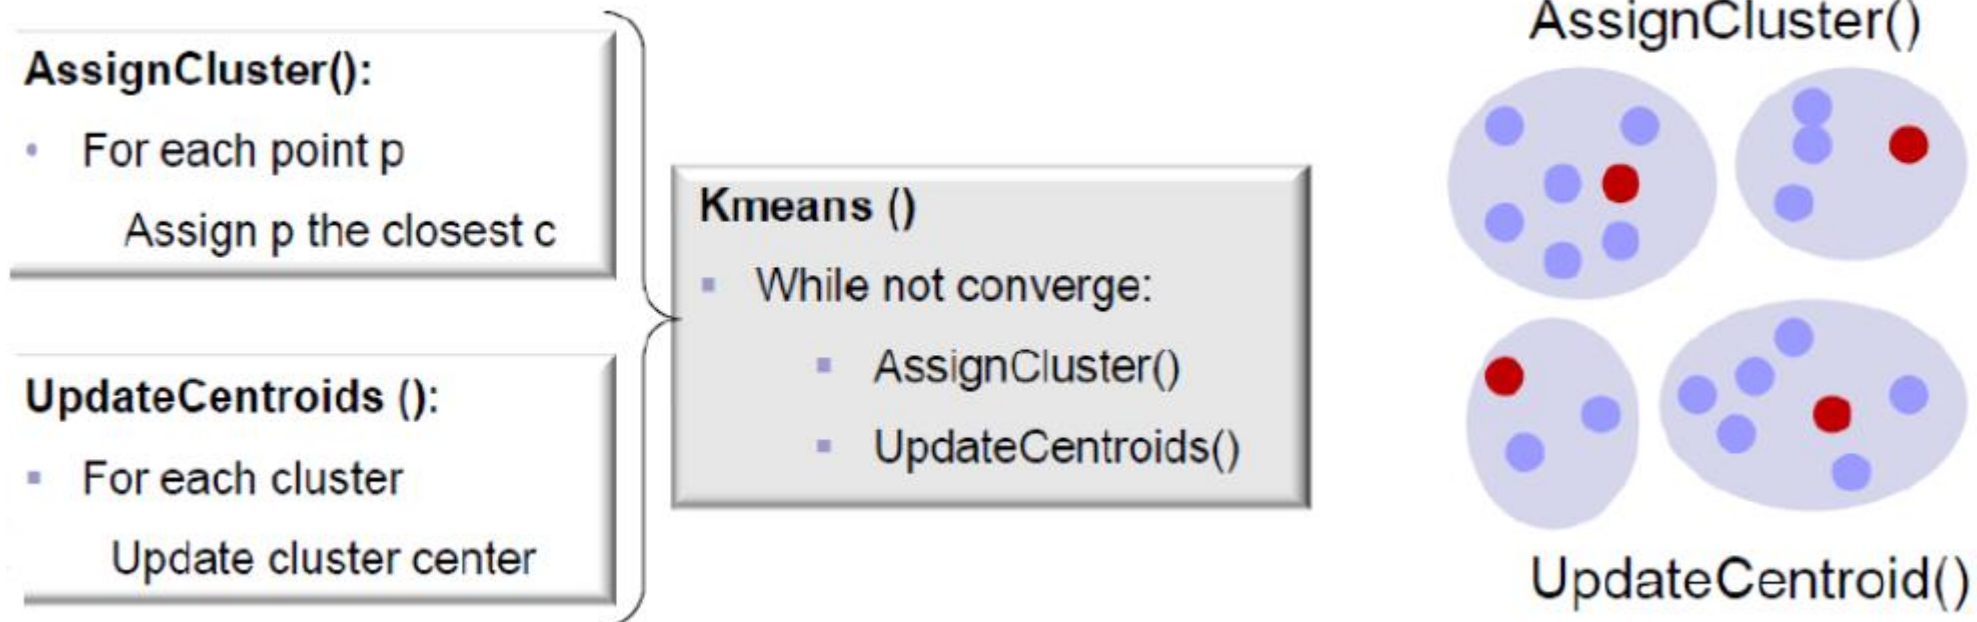
Lecture 6

Recap

- MapReduce introduced by Google
 - Simple programming model for building distributed applications that process vast amounts of data
 - Runtime for executing jobs on large clusters in a reliable, fault-tolerant manner
- Hadoop makes MapReduce broadly available
 - HDFS becomes central data repository
 - Becomes Defacto standard for batch processing
- Stonebraker & Dewitt: Mapreduce a major step backwards

New applications, new workloads

- Iterative computations
 - Ex: More and more people aiming to get insights from data
 - Apache Mahout becomes popular framework for ML over Hadoop
- How would we implement K-Means with MapReduce?
- Traditional k-means



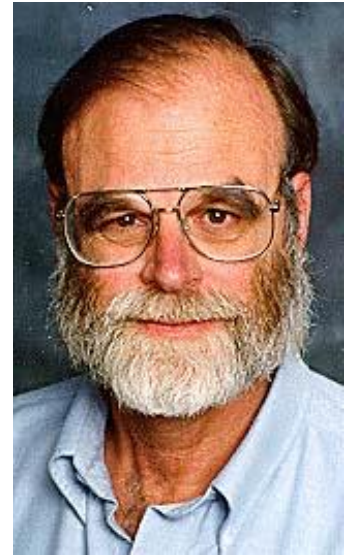
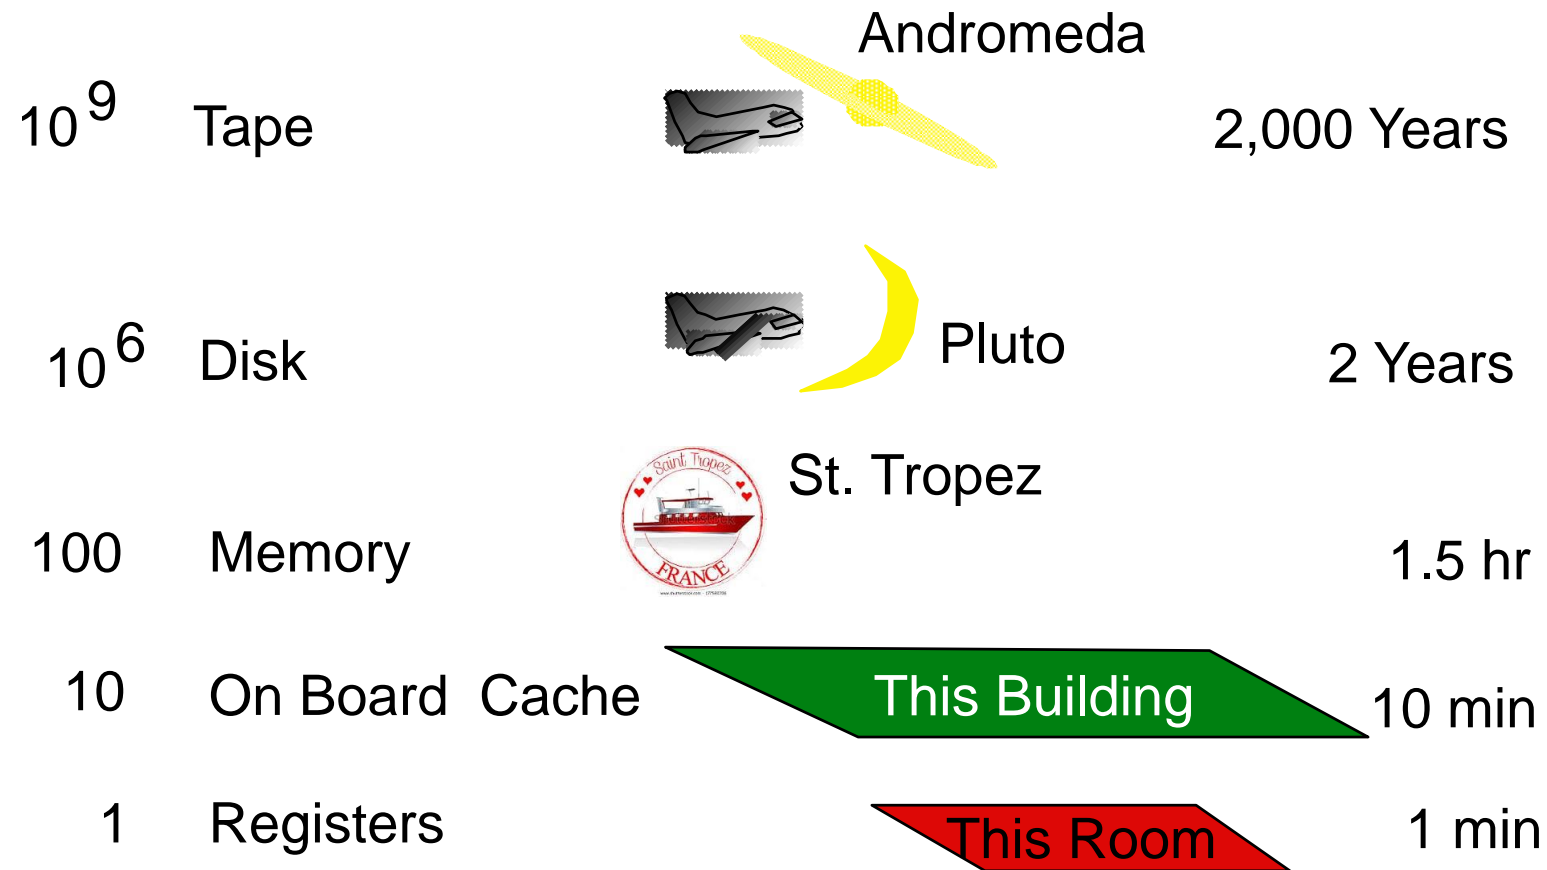
K-means MapReduce algorithm

- **Configure:** A single file containing cluster centers
- **Mapper**
 - Input: Input data points
 - Compute: Distance of a point from each centroid to identify a cluster
 - Output: (cluster id, data id)
- **Reducer**
 - Input: (cluster id, data id)
 - Compute: New cluster centroid based on data points assigned
 - Output: (cluster id, cluster centroid)
- **Driver**
 - Each iteration produces new cluster centroids
 - Run multiple iteration jobs using mapper + reducer until convergence: High overhead leading to poor performance

MapReduce & Iterative Computations

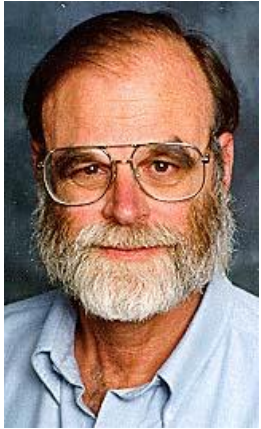
- MapReduce is built for batch processing
 - Entirely disk based: Input and output sit on HDFS
- Let us look at **k-means algorithm, 1 iteration**
 - **HDFS Read**
 - **Map**(Assign sample to closest centroid)
 - NETWORK Shuffle
 - **Reduce**(Compute new centroids)
 - **HDFS Write**
- Each iteration reads and writes data from disk-based HDFS
 - To understand why this is bad, let us look at the memory hierarchy

Understanding Memory Hierarchy: How Far Away is the Data?



Jim Gray

1980s Database Administrators Dilemma

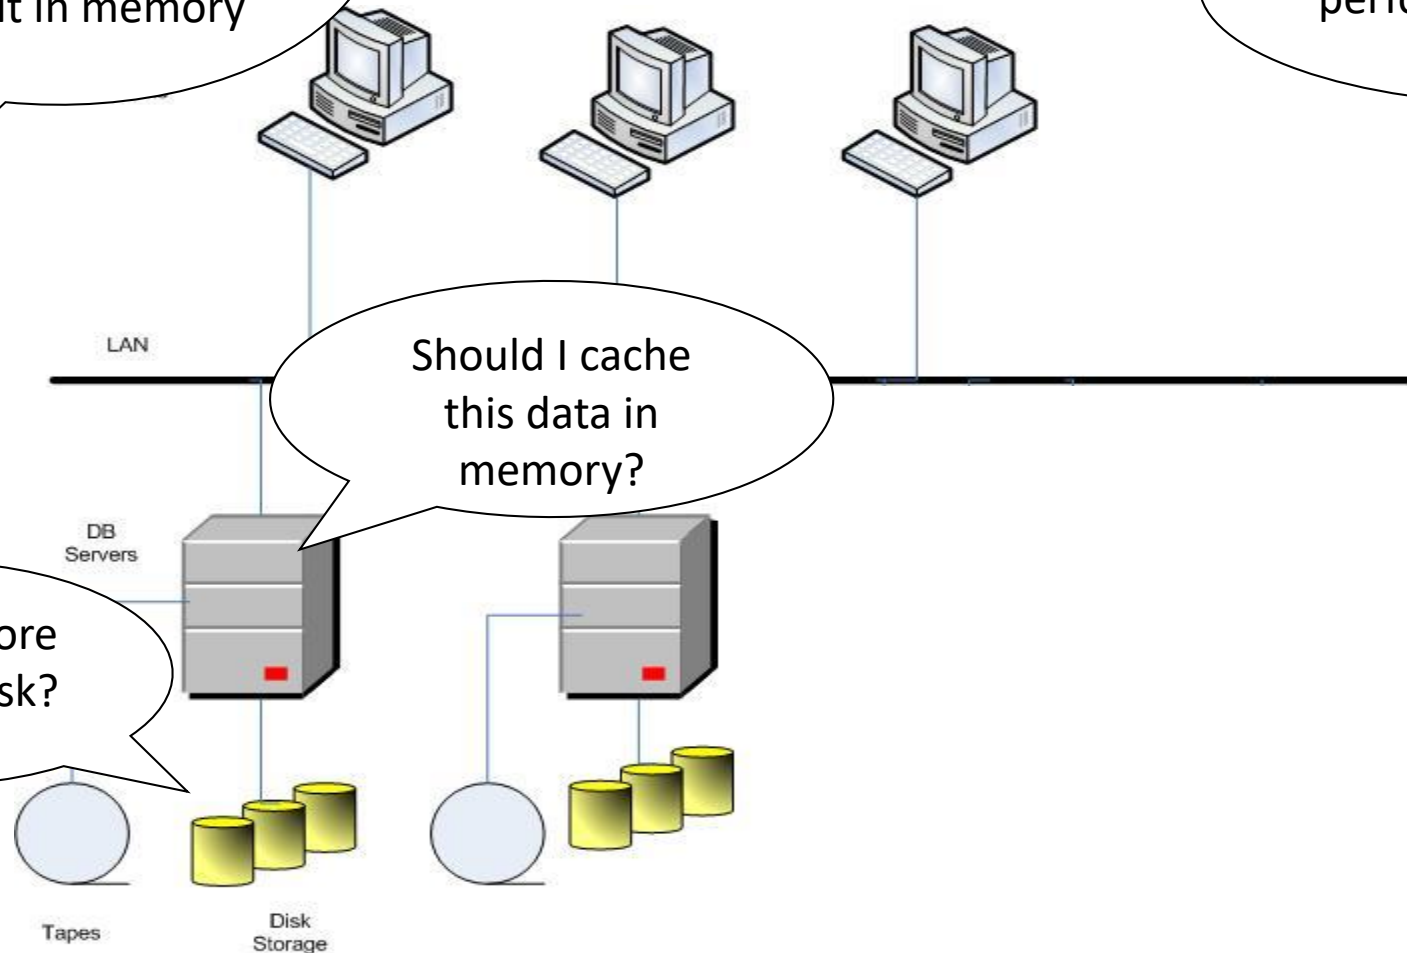


If a data is accessed more than once within 5 minutes, cache it in memory

How do I improve the performance of my DB server?

Should I cache this data in memory?

Should I store data on disk?



Tandem Computers: Price/performance

- Tandem disk: **\$1k/access**
 - cost: \$15k for 180MB
 - performance: 15 accesses / second
- Tandem CPU + supporting hardware: **\$1k/access**
 - Cost: \$15,000
- Cost of accessing data from disk: **\$2k/access**
- Memory cost: \$5k for 1MB => **\$5/KB**



Five-minute rule

- Cost of accessing data from disk: **\$2k/access**, Memory cost: **\$5/KB**
- If we keep 1KB in memory, assuming we have 1 access/sec
 - We save \$2k of disk i/o by paying \$5 for memory
- If we have 1 access every 10 secs => 0.1 access/sec
 - We save \$200 of disk i/o by paying \$5 for memory
- .
- .
- .
- Break even point : 1 access every 400 secs

400 seconds ~ 5 minutes

Five-minute rule: then and now

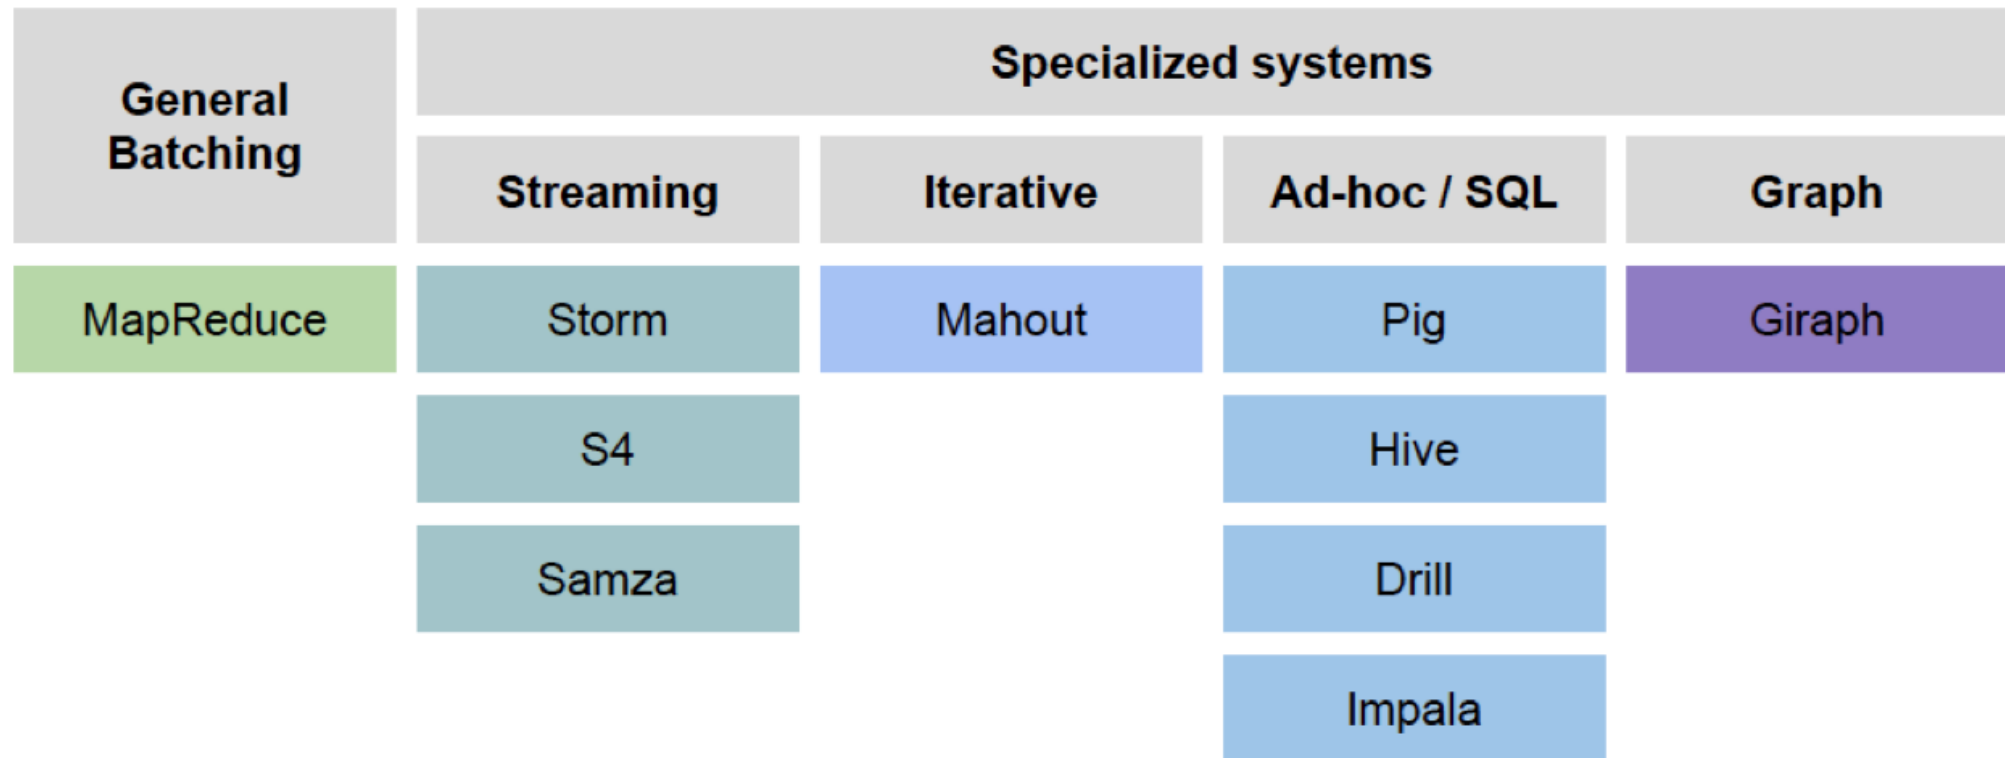
Page size (4KB)	1987	Now
RAM-HDD	5 mins	5 hours

- RAM-HDD break-even 60x higher due to drop in DRAM price
 - Take away: Never ever go to disk!
- See “Five minute rule” CACM paper for more details
 - <https://cacm.acm.org/magazines/2019/11/240388-the-five-minute-rule-30-years-later-and-its-impact-on-the-storage-hierarchy/fulltext>

MapReduce/Hadoop and memory hierarchy

- Hadoop misaligned with five-minute rule
 - All data is stored in disk
 - Does not cache data in memory even if workload can fit
- Hadoop unfit for new classes of workloads
 - Interactive and iterative applications are bottlenecked by disk
- MapReduce was also too simple a computational model
 - Algorithm design with just map and reduce functions is non trivial

Hadoop: Fractured ecosystem



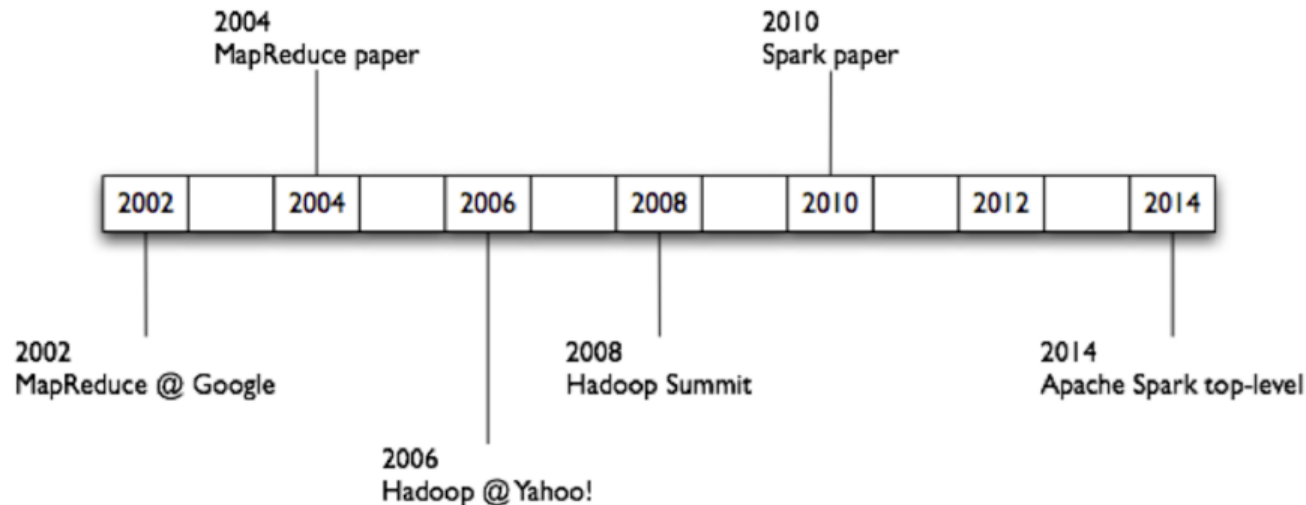
- Specialized systems emerged with no unified vision
 - Diverse APIs, sparse modules, high operational costs
 - MapReduce runtime replaced with more optimized ones

Lighting a Spark

Flexible, in-memory data processing framework written in Scala

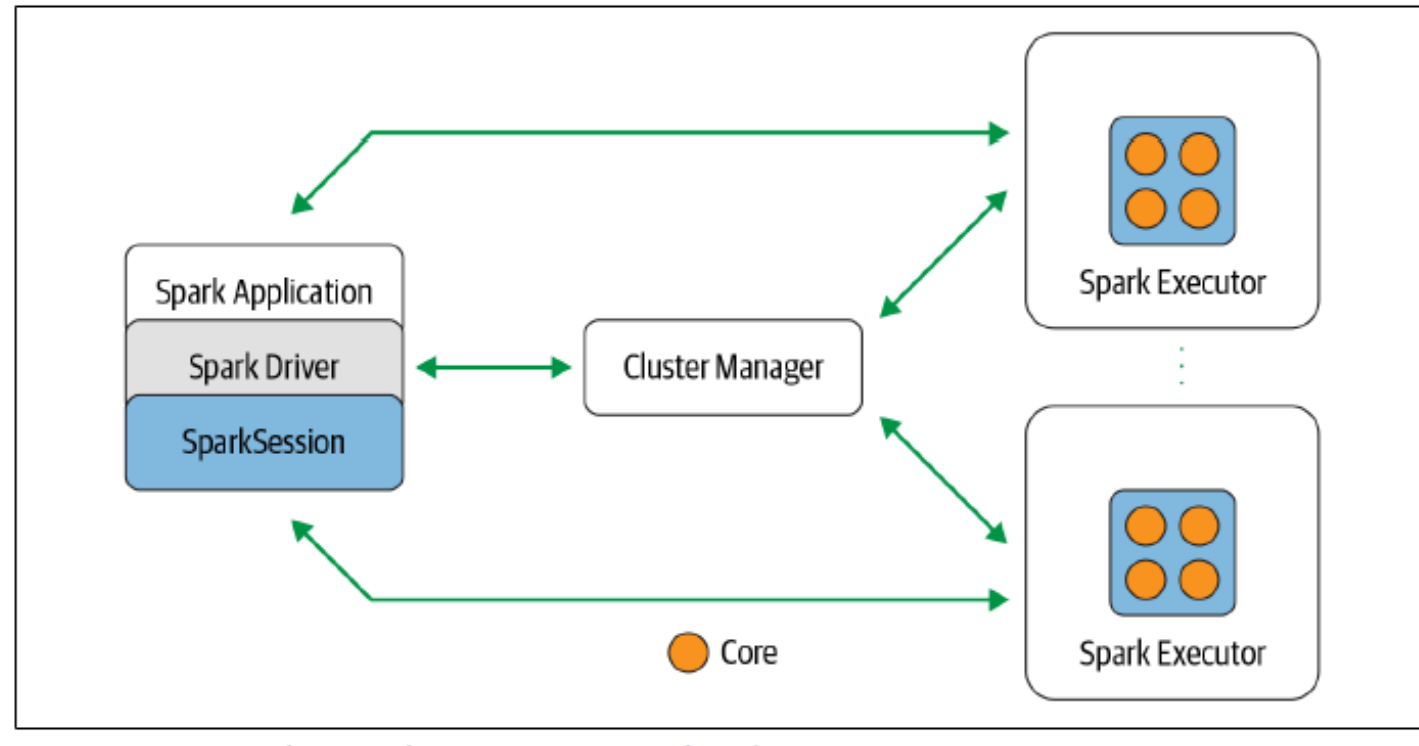
Central Ideas

- Exploit memory by caching data to enable fast data sharing
- Generalize the two-stage computational model of mapreduce to a Directed Acyclic Graph-based one that can support a richer API



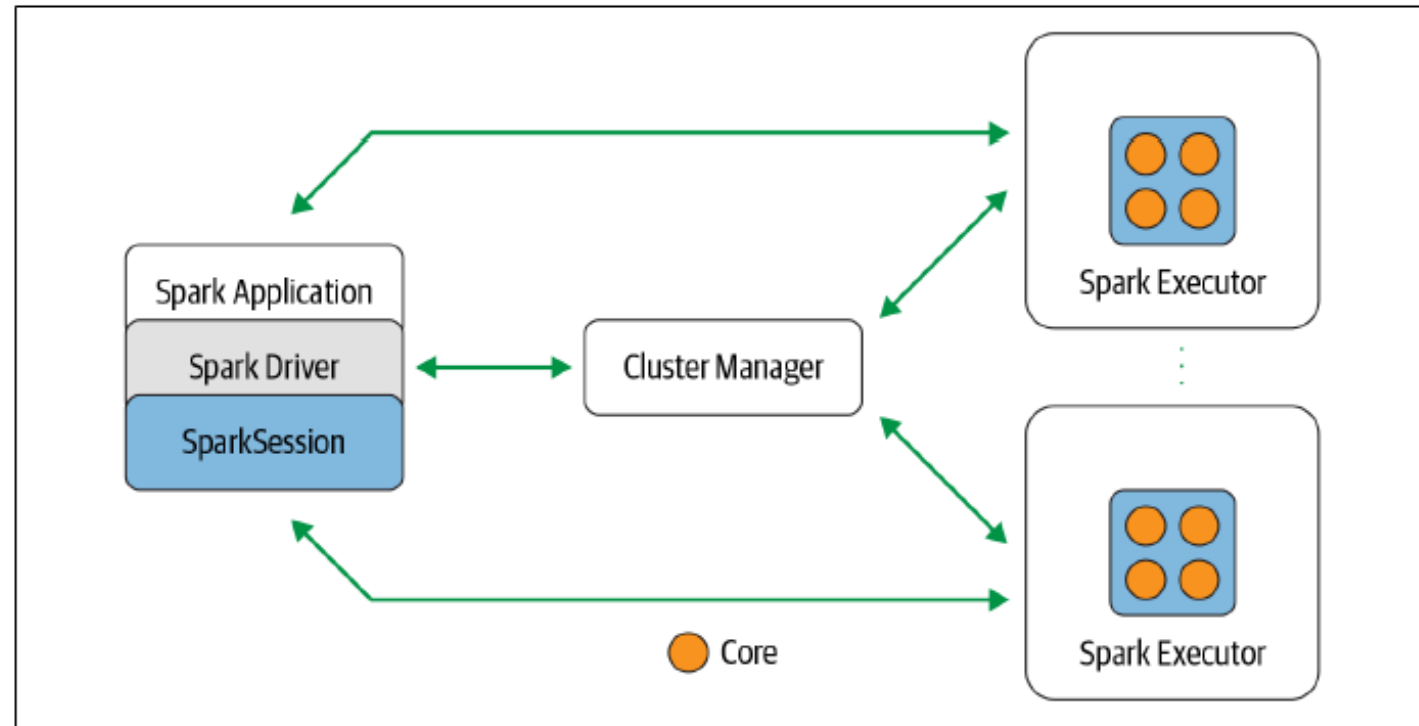
Spark Distributed Architecture

- Spark Application
 - User program built on Spark
 - Contains the Spark driver
- Spark Driver
 - Transforms all the Spark operations into DAG computations
 - Communicates with the cluster manager & requests resources (CPU, memory, etc.) for Spark executors
 - Schedules computations
 - Instantiates SparkSession



Spark Distributed Architecture

- SparkSession
 - A unified conduit to all Spark operations and data
- Cluster Manager
 - Responsible for managing and allocating resources
 - 4 supported: Standalone, YARN, Mesos, Kubernetes
- Executor
 - Responsible for executing tasks
 - Usually one per node, but depends on deployment mode



RDD: Need for a new abstraction

- Need an efficient way to share data stored in memory
- Traditional way: Distributed shared memory abstraction
 - General purpose, extends single-node shared memory to a cluster
 - Applications can make fine-grained updates to any data in memory
 - Can be used to build very efficient applications
- Problem: Fault tolerance
 - Need to replicate data across nodes or log updates which is 10-100x slower than memory write
 - Too expensive for data-intensive apps
- Goal: In-memory abstraction that provides fault-tolerance and efficiency

Resilient Distributed Dataset

RDD (Resilient Distributed Dataset): Restricted form of DSM

- An immutable, partitioned collection of objects
 - Can only be built through coarse-grained deterministic transformations
-
- **RDD are data structures that:**
 - Either point to a direct data source (e.g. HDFS)
 - Apply some transformations to its parent RDD(s) to generate new data elements

RDD

- An abstraction that encapsulates 3 things
 - Dependencies, Partitions, Compute function
- Dependencies
 - Instruct Spark how an RDD is constructed
 - To reproduce results, Spark can recreate an RDD from these dependencies and replicate operations on it => resiliency
- Partitions
 - Split the work to parallelize computation on partitions across executors
 - Exploit data locality
- Compute function: Partition -> Iterator[t]
 - A function that produces an iterator for data stored in RDD

RDD: Example

- Query: Find average age for each name
 - aggregate all the ages for each name
 - group by name
 - average the ages

```
# In Python  
# Create an RDD of tuples (name, age)  
dataRDD = sc.parallelize([("Brooke", 20), ("Denny", 31), ("Jules", 30),  
    ("TD", 35), ("Brooke", 25)])  
# Use map and reduceByKey transformations with their lambda  
# expressions to aggregate and then compute average  
  
agesRDD = (dataRDD  
    .map(lambda x: (x[0], (x[1], 1)))  
    .reduceByKey(lambda x, y: (x[0] + y[0], x[1] + y[1]))  
    .map(lambda x: (x[0], x[1][0]/x[1][1])))
```

RDD Transformations

- Set of operations that define how to transform an RDD
 - Examples: `map()`, `filter()`, `select()`, `join()`, `orderBy()`, ...
- As in relational algebra, the application of a transformation to an RDD yields a new RDD
 - RDD are immutable
- Transformations are lazily evaluated
 - Computation that performs the transformation is not performed immediately

RDD Actions

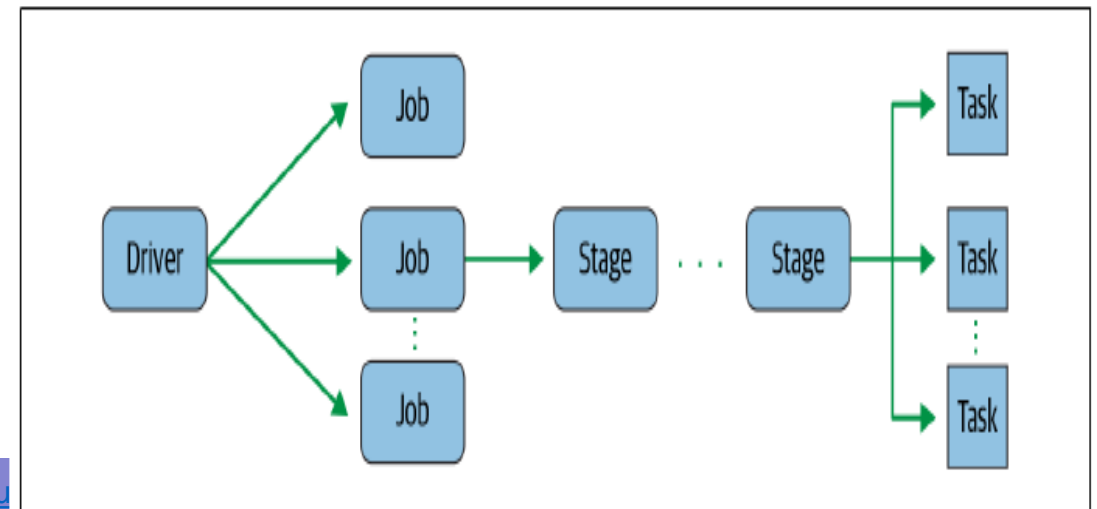
- Actions trigger computation of the chain of transformations
- Some actions only store data to an external data source (e.g. HDFS)
 - Ex: `show()`, `take()`, `count()`, `collect()`, ..
- Others fetch data from the RDD (and its transformation chain) upon which the action is applied, and convey it to the driver
 - `count()` – return the number of elements
 - `take(n)` – return an array of the first *n* elements
 - `collect()` – return an array of all elements
 - ...

RDD & Lazy Execution Demo

- <https://mediaserver.eurecom.fr/permalink/v12641880f54fqht30dc/iframe/#start=4615>

Spark Execution

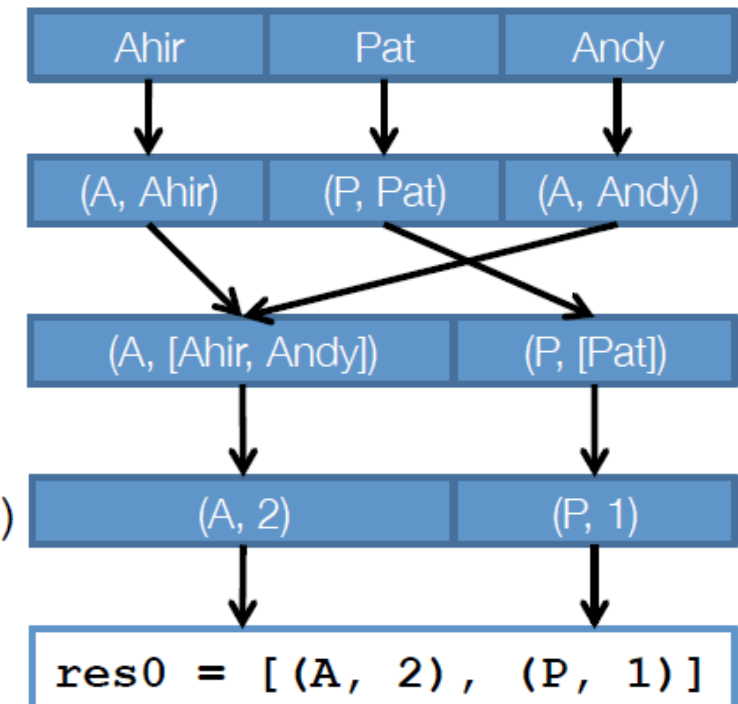
- Spark Driver
 - Converts your Spark application into one or more Spark jobs
 - Transforms each job into a DAG – execution plan
- Job broke into stages
 - Stages are created based on what operations can be performed in parallel
 - Dictate data transfer among Spark executors.
- Stages composed of tasks
 - Task - a unit of execution
 - Maps to a single core
 - Maps to 1 partition of data



Spark DAG execution: An Example

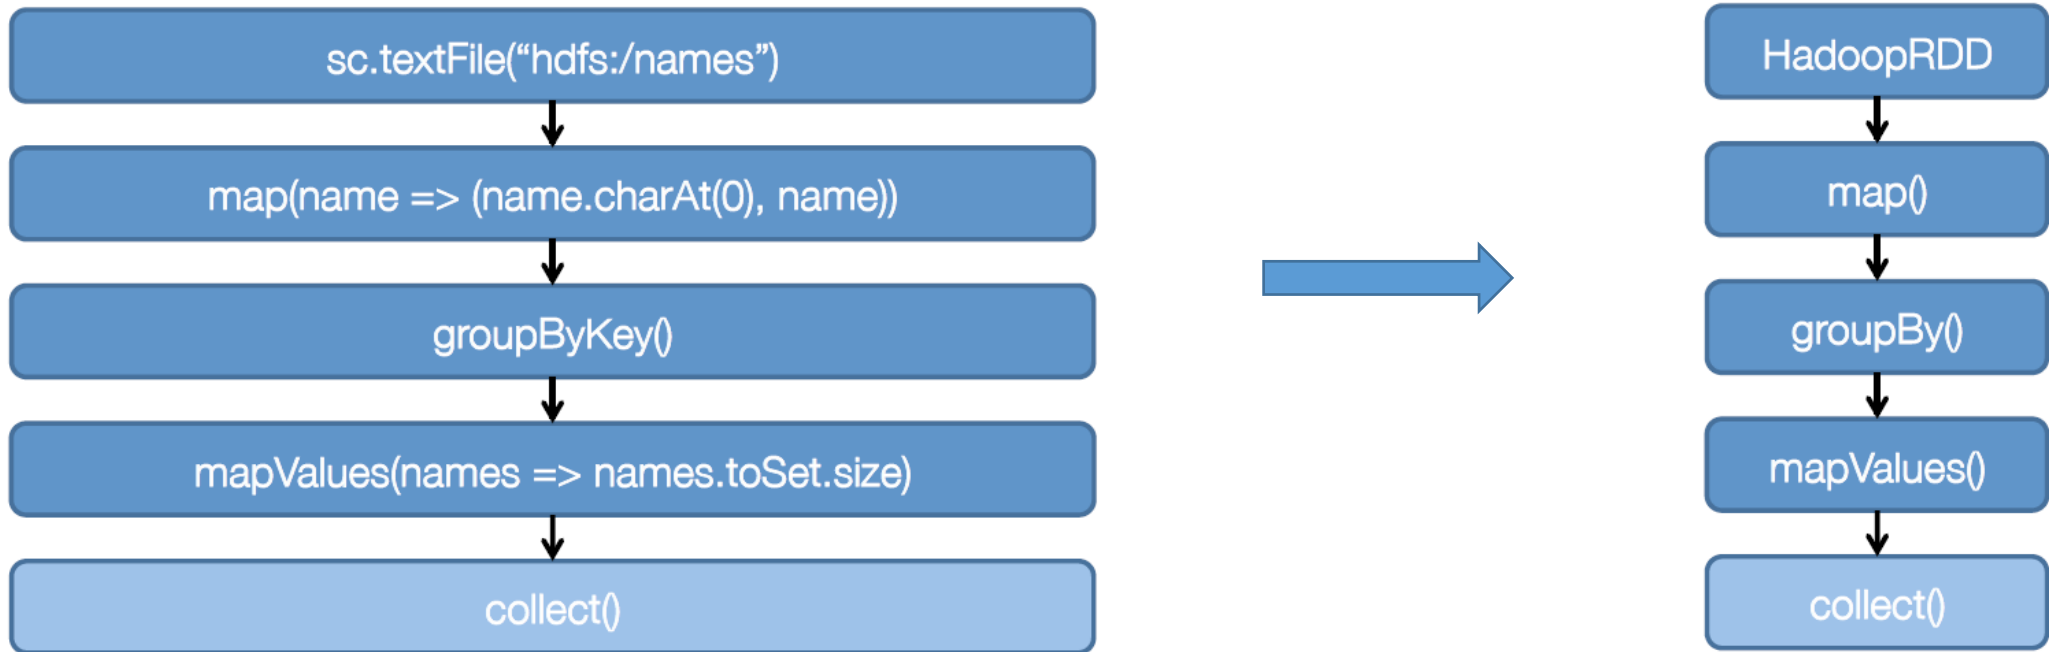
- Goal: Find the number of distinct names per first letter

```
sc.textFile("hdfs:/names")  
  
.map(name => (name.charAt(0), name))  
  
.groupByKey()  
  
.mapValues(names => names.toSet.size)  
  
.collect()
```



Spark Execution (1)

1. Create a DAG of RDDs to represent computation

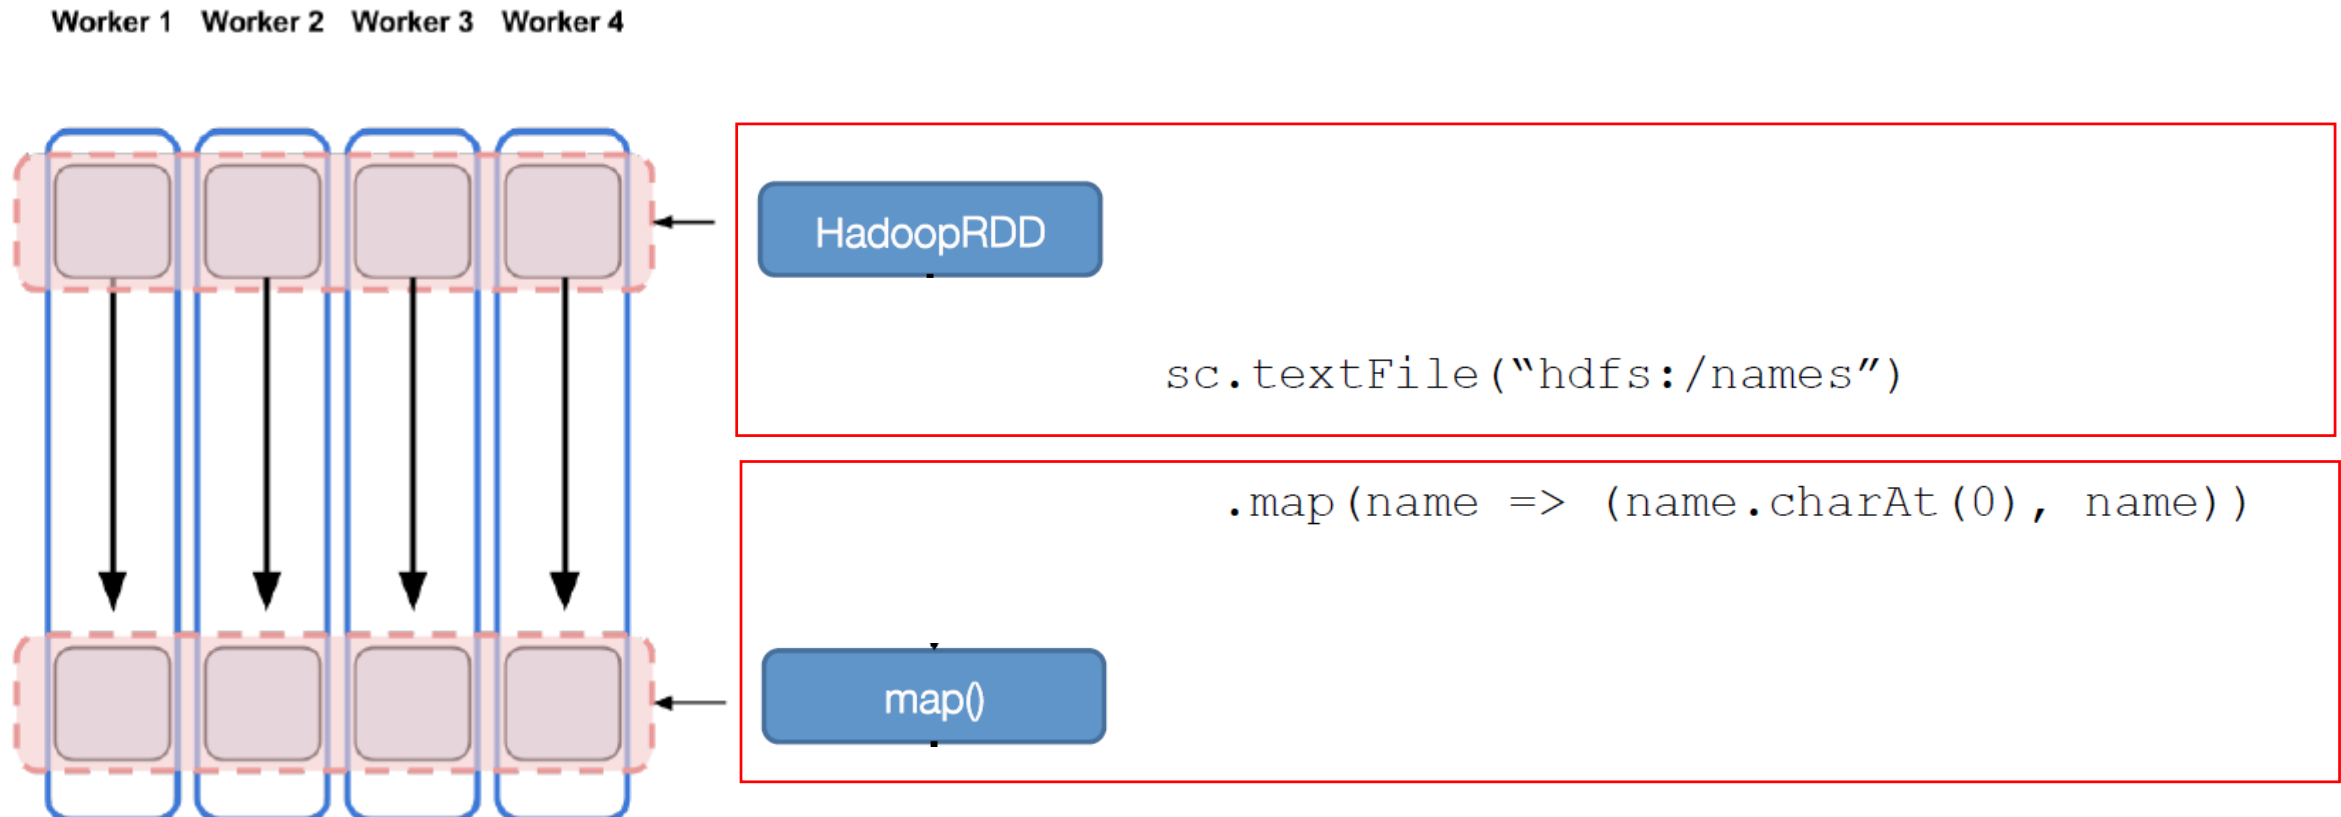


Spark Execution (2)

1. Create a DAG of RDDs to represent computation
2. Create logical execution plan for the DAG
 - Split DAG into “stages” based on dependencies
 - Pipeline as much as possible

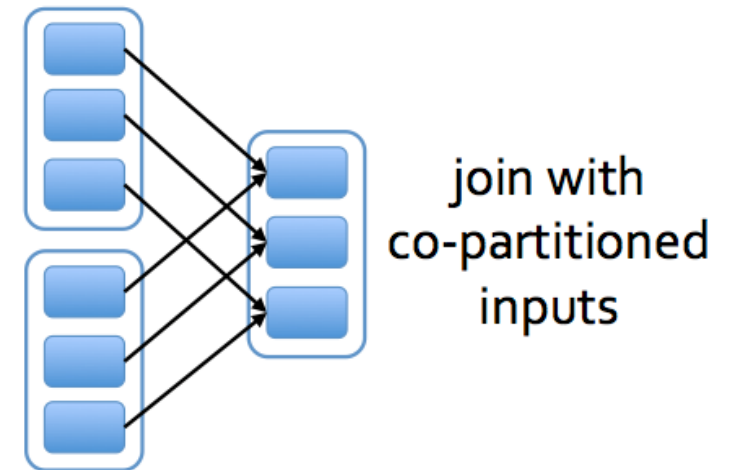
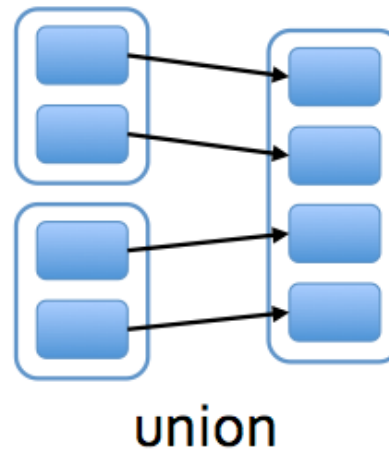
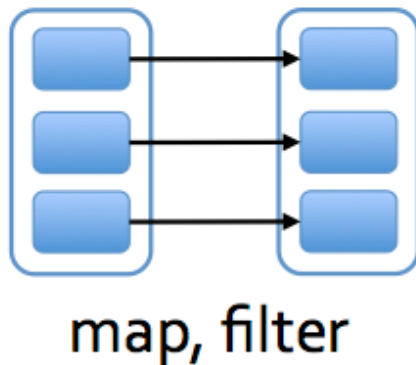
RDD: Data Set vs Partition Views

Much like in Hadoop MapReduce, each RDD is stored physically in multiple nodes as input partitions



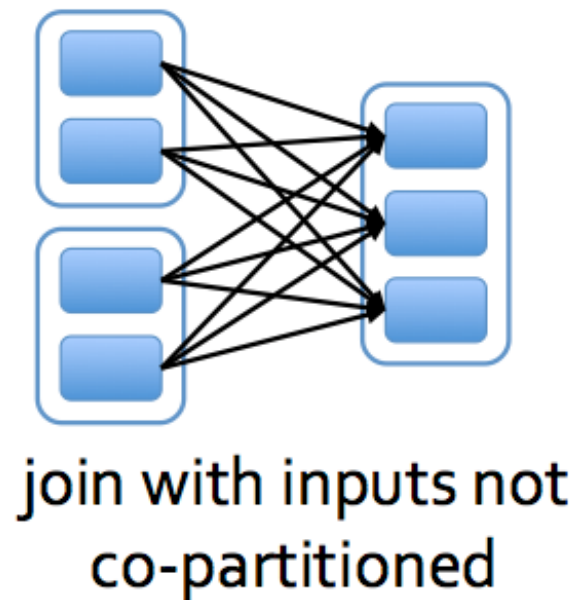
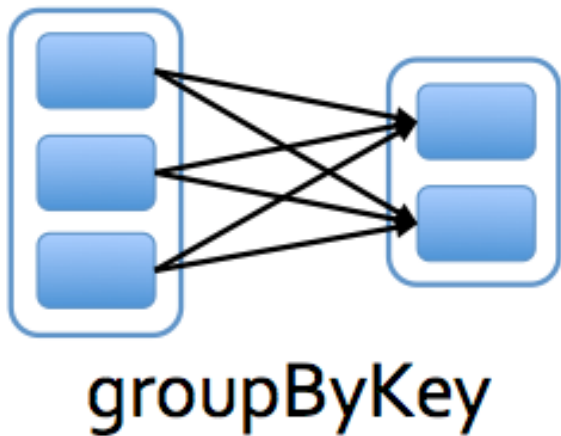
A word about dependencies (1)

- Dependencies determine the need to shuffle data
 - Two types: Narrow and wide
- Narrow dependencies
 - Each partition of the parent RDD is used by at most one partition of the child RDD
 - Task can be executed locally and we don't have to shuffle. (E.g. map, flatMap, filter, sample)



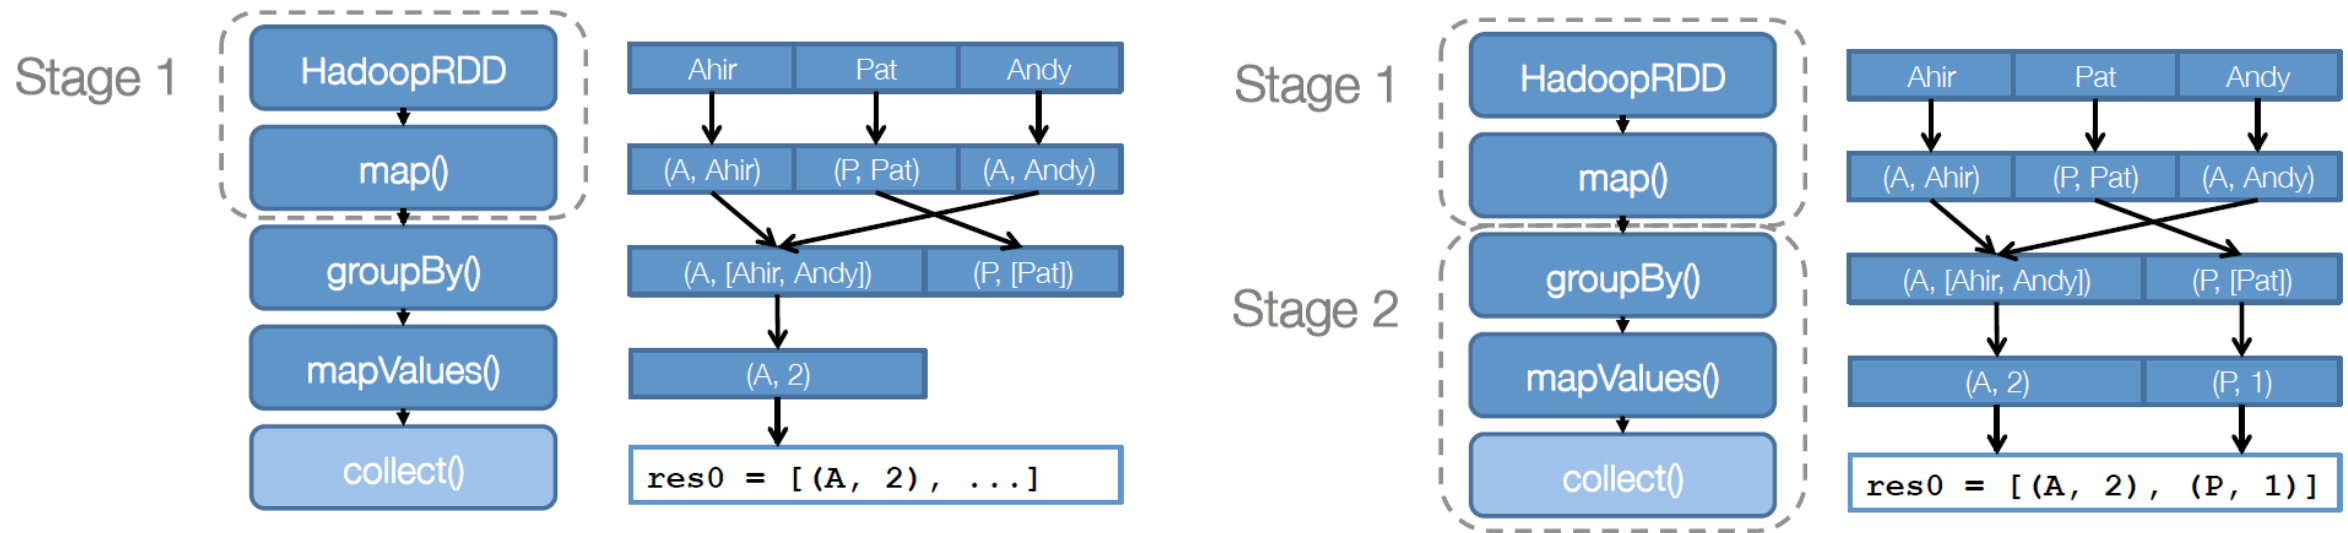
A word about dependencies (2)

- Wide dependencies
 - Multiple child partitions may depend on one partition of the parent RDD
 - We have to shuffle data (E.g. sortByKey, reduceByKey, groupByKey, cogroupByKey, join, cartesian)



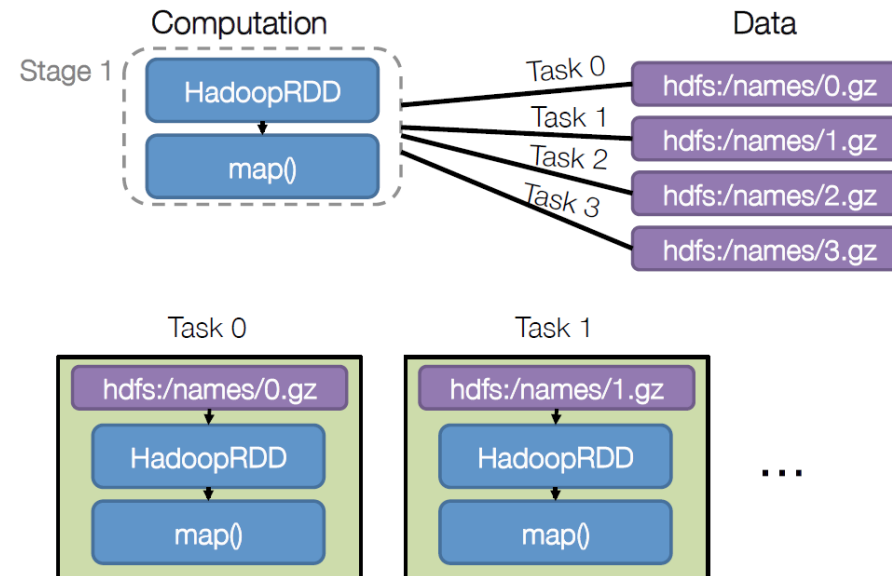
How does Spark execute this job?

1. Create a DAG of RDDs to represent computation
2. Create logical execution plan for the DAG
 - Pipeline as much as possible
 - Split DAG into “stages” based on need to shuffle data



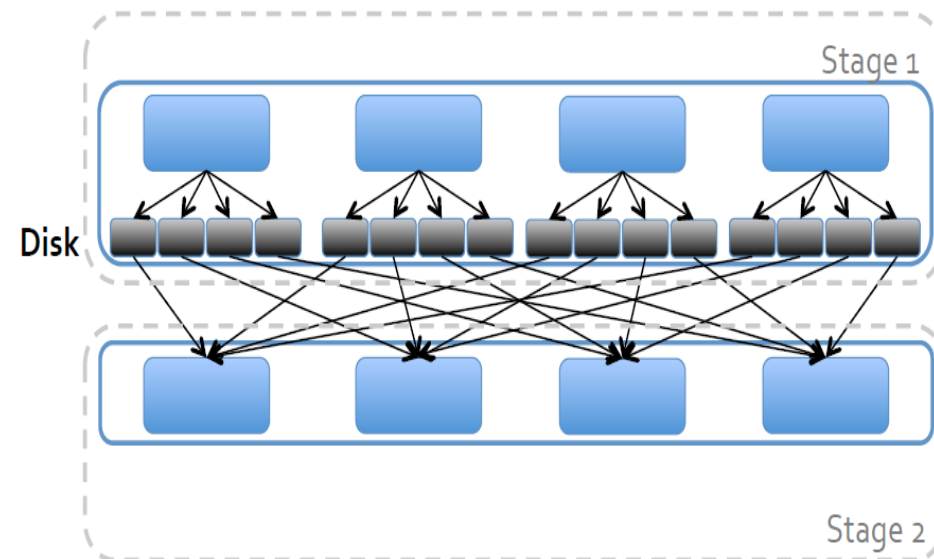
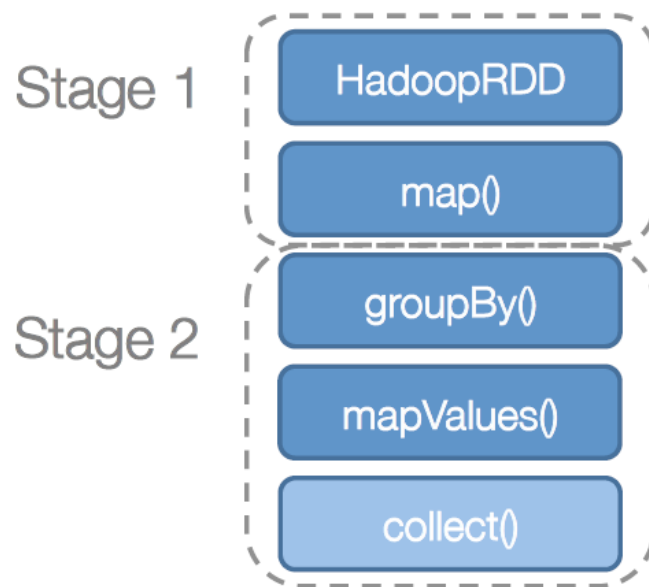
Spark Execution (3)

1. Create a DAG of RDDs to represent computation
2. Create logical execution plan for the DAG
3. Split each stage into tasks and execute tasks stage by stage
 - Task = Data + Computation
 - In this example, all tasks from stage 1 would be executed together first



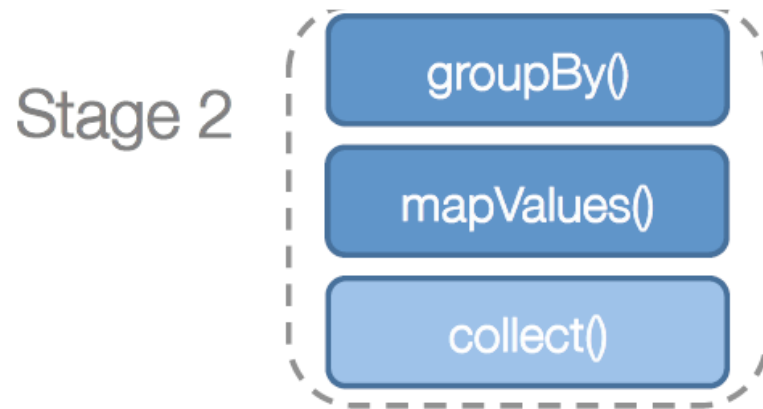
Spark Execution (3)

1. Create a DAG of RDDs to represent computation
2. Create logical execution plan for the DAG
3. Split each stage into tasks and execute tasks stage by stage
 - In this example, all tasks from stage 1 would be executed together first
 - After stage 1, pull-based shuffle occurs (intermediates written to files and pulled)

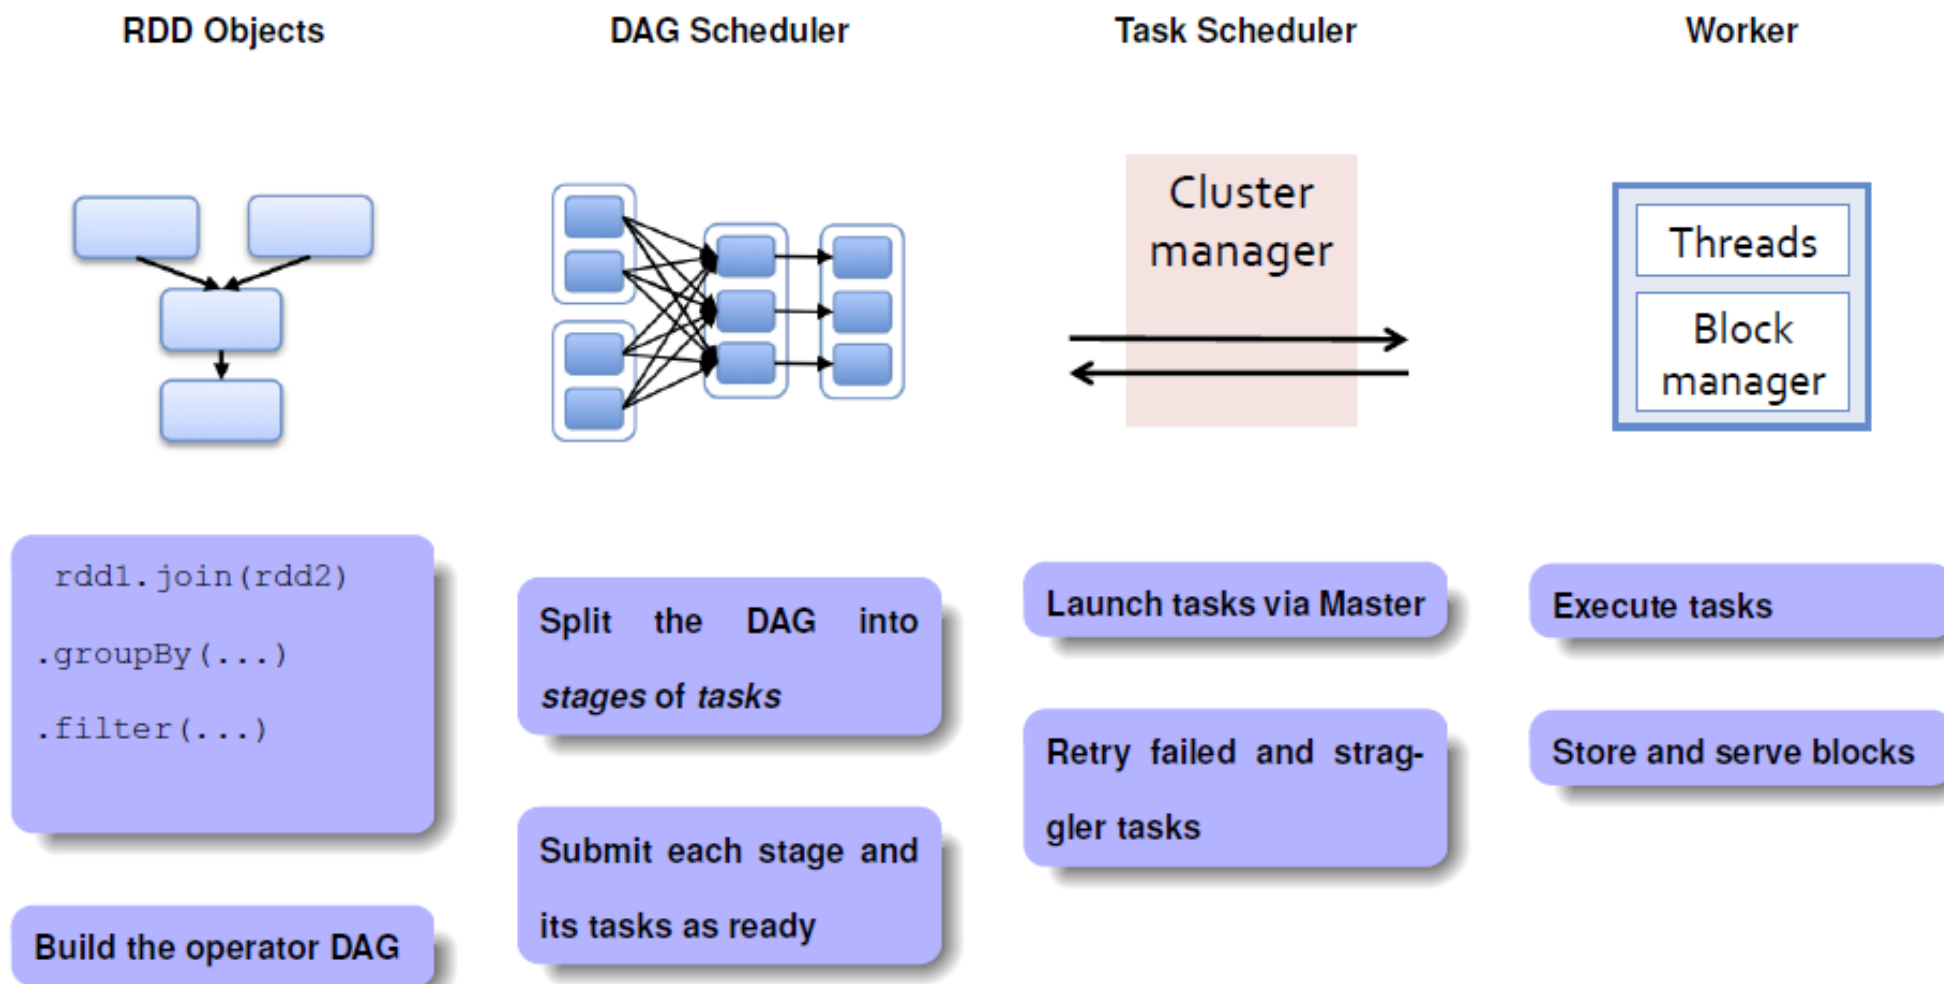


Spark Execution (3)

1. Create a DAG of RDDs to represent computation
2. Create logical execution plan for the DAG
3. Split each stage into tasks and execute tasks stage by stage
 - In this example, all tasks from stage 1 would be executed together first
 - After stage 1, pull-based shuffle occurs (intermediates written to files and pulled)
 - Now, tasks from stage 2 are executed (operators pipelined in each task)



Putting it all together



RDD to Structured API

- Spark can support interactive workloads
 - But working with RDD is procedural
- SQL as a high-level programming language
 - Offers expressiveness, succinctness
 - Enables compatibility with existing tools, e.g. Business Intelligence using JDBC
 - Large pool of engineers proficient in SQL

```
# In Python  
# Create an RDD of tuples (name, age)  
dataRDD = sc.parallelize([("Brooke", 20), ("Denny", 31), ("Jules", 30),  
    ("TD", 35), ("Brooke", 25)])  
# Use map and reduceByKey transformations with their lambda  
# expressions to aggregate and then compute average  
  
agesRDD = (dataRDD  
    .map(lambda x: (x[0], (x[1], 1)))  
    .reduceByKey(lambda x, y: (x[0] + y[0], x[1] + y[1]))  
    .map(lambda x: (x[0], x[1][0]/x[1][1])))
```

```
SELECT name, avg(age)  
FROM people  
GROUP BY name
```

DataFrame: Schema

- **General idea borrowed from Python Pandas**
 - Tabular data with an API
- **Schema to the rescue**
 - A distributed collection of rows organized into named, typed columns
 - Basic types and Structured/Complex types supported
 - Schema defines column names and associated types
 - 3 ways to get schema definition (demo here: <https://mediaserver.eurecom.fr/permalink/v12641880f54fqht30dc/iframe/#start=6211>)
 - (i) Define with struct type, (ii) define with DDL, (iii) auto infer
 - Columns and Rows are objects with APIs

APIs

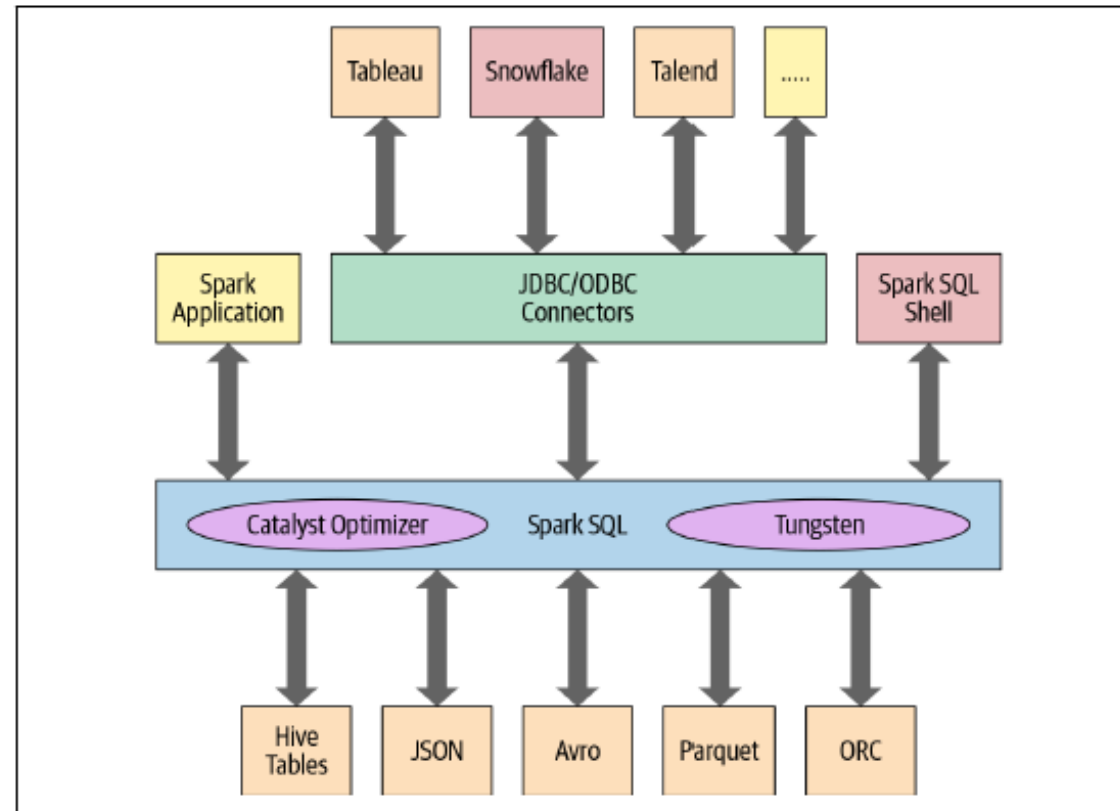
- DataSource API
 - Enables you to read/write data from/to a DataFrame from myriad data sources in formats such as JSON, CSV, Parquet, Text, Avro, ORC, etc.
- Transformations and Actions
 - Relational Projections: done with select() method
 - Relational Selection: done with filter() or where() method
 - Aggregations: groupBy, orderBy, count, ...
 - Descriptive stats: min, max, sum, avg
- Demo:
<https://mediaserver.eurecom.fr/permalink/v12641880f54fqht30dc/iframe/#start=6495>
 - time taken to infer schema vs predeclare
 - Transformations, actions

RDD vs DF

- Use RDD when
 - Want to precisely instruct Spark *how to do* a query
 - Can forgo the code optimization, efficient space utilization, and performance benefits available with DataFrames and Datasets!
 - You can get rdd from df: `df.rdd`
- Basically, save yourself some time and use DF

SparkSQL engine

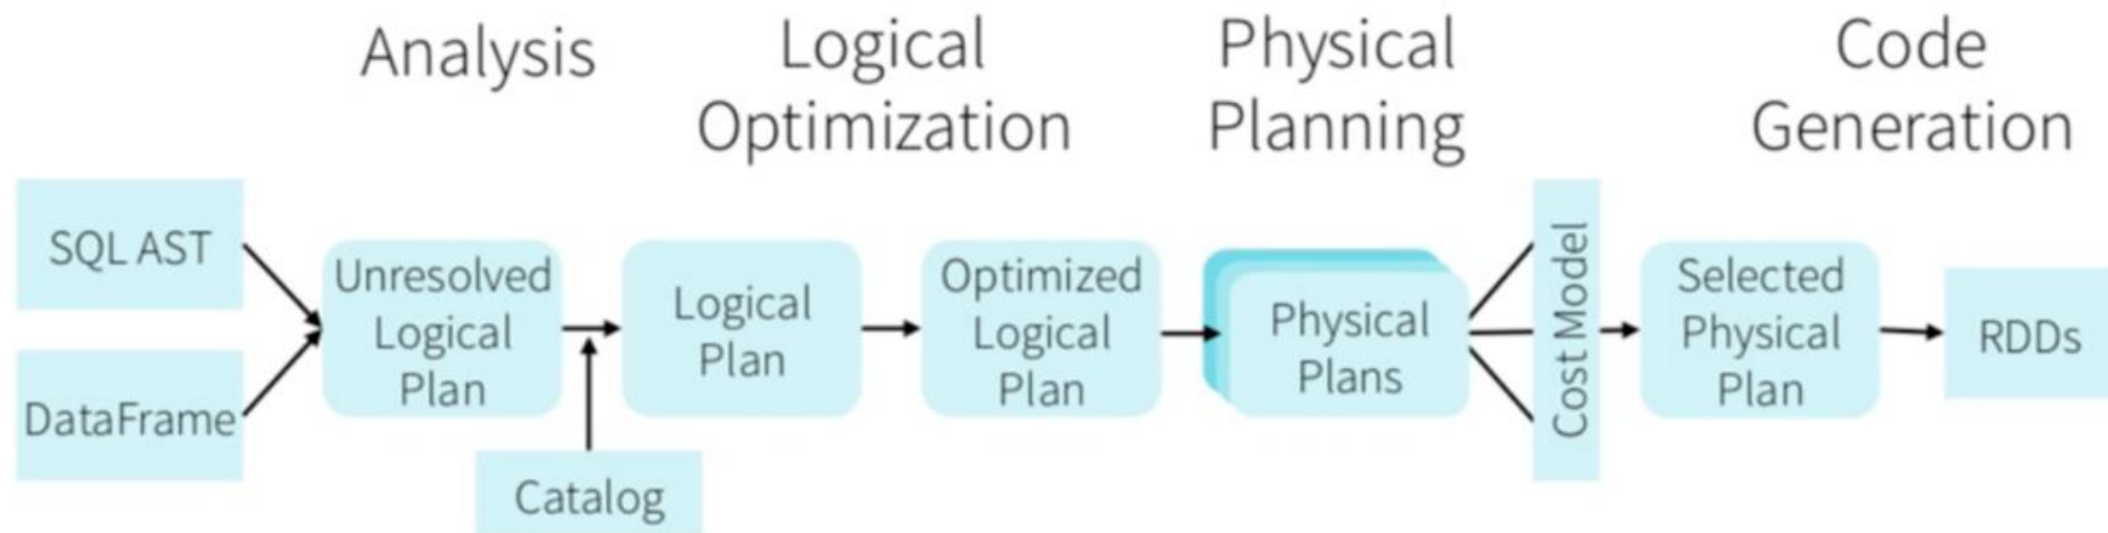
- The substrate on which structured APIs are built
- Core components
 - Catalyst and Tungsten



Catalyst Optimizer

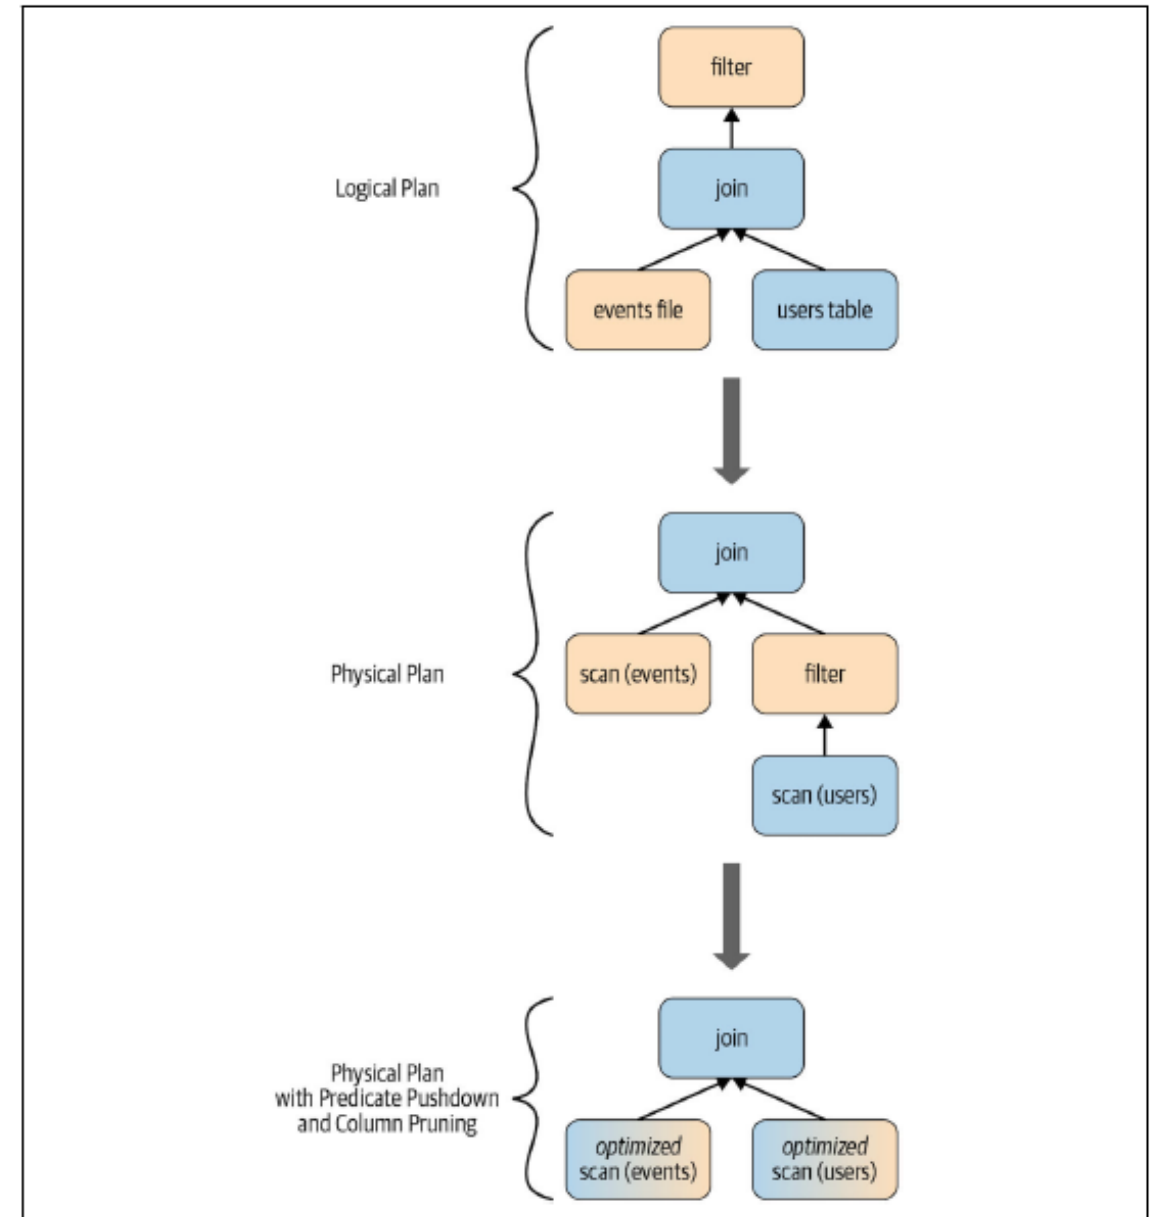
- **Reminiscent of traditional database systems**

- Analysis: SQL to Plan Abstract Syntax Tree
- Logical & physical optimization: Use cost-based optimization to pick optimal plan



Catalyst Example

```
// In Scala
// Users DataFrame read from a Parquet table
val usersDF = ...
// Events DataFrame read from a Parquet table
val eventsDF = ...
// Join two DataFrames
val joinedDF = users
    .join(events, users("id") === events("uid"))
    .filter(events("date") > "2015-01-01")
```



Tungsten & Code Generation

- Take optimized physical plan and do “Full stage code generation”
 - collapses the whole query into a single function
 - getting rid of virtual function calls
 - employing CPU registers for intermediate data
- Demo of Catalyst and Tungsten:
<https://mediaserver.eurecom.fr/permalink/v12641880f54fqht30dc/iframe/#start=7410>

Spark & Caching

- `DataFrame.cache()`
 - store as many of the partitions read in memory across Spark executors as memory
- `Dataframe.persist(StorageLevel)`
 - Control how your data is cached
- `Unpersist`
 - Remove any cached data
- `Cache/persist` are hints
 - `DataFrame` is not fully cached until you invoke an action
- Demo of caching perf:
<https://mediaserver.eurecom.fr/permalink/v12641880f54fqht30dc/iframe/#start=8042>

StorageLevel	Description
MEMORY_ONLY	Data is stored directly as objects and stored only in memory.
MEMORY_ONLY_SER	Data is serialized as compact byte array representation and stored only in memory. To use it, it has to be deserialized at a cost.
MEMORY_AND_DISK	Data is stored directly as objects in memory, but if there's insufficient memory the rest is serialized and stored on disk.
DISK_ONLY	Data is serialized and stored on disk.
OFF_HEAP	Data is stored off-heap. Off-heap memory is used in Spark for storage and query execution; see "Configuring Spark executors' memory and the shuffle service" on page 178.
MEMORY_AND_DISK_SER	Like MEMORY_AND_DISK, but data is serialized when stored in memory. (Data is always serialized when stored on disk.)

Spark & RDBMS: Summary

- Spark: unified analytics engine
 - Quickly adopted RDBMS concepts to optimize SQL analytics
 - Other libraries developed for machine learning (Mlib), graph analytics(GraphX),...
 - RDD: an underlying abstraction that supports several libraries
- DBMSs have also evolved
 - Disk-based to in-memory to NVM
 - One-size-fits-all “OldSQL” DBMS to customized “NewSQL” engines
 - column stores for Business Intelligence
 - highly parallel transaction engines for OLTP
 - Array databases for scientific applications
 - ...
 - NewSQL still king for structured data management