

# Agreement in Distributed Systems: Consensus

Lecture 11

# Recap of last week

- Agreement in distributed systems
  - How do we get all nodes in a distributed system to act in concert despite failures?
- Agreement Requirements
  - **Safety** (correctness)
    - All nodes agree on the same value
    - The agreed value X has been proposed by some node
  - **Liveness** (fault tolerance, availability)
    - If less than some fraction of nodes crash, the rest should still reach agreement

# Recap of last week: Atomic commitment

- **Atomic commitment problem**
  - One type of agreement problem: Participants need to agree on a value, but they have specific constraints on whether they can accept any particular value.
- We looked specifically at atomic commitment in distributed databases: **How to provide atomicity (A of ACID) in the presence of failures?**
  - 2-phase commit: Safe but not live due to blocking
  - Non-blocking 3-phase commit: Live but cannot handle network partition

# Consensus: Formal definition

- Problem
  - A collection of processes,  $P_i$ .
  - They propose values  $V_i$  (e.g., time to attack, client update, lock requests, ...), and send messages to others to exchange proposals.
  - Different processes may propose different values, but they can all accept any of the proposed values.
  - Only one of the proposed values will be “chosen” and eventually (once all failures are addressed) all of the nodes learn that *one chosen value*.
- Requirements:
  - **Consistency:** once a value is chosen, the chosen value of all working processes is the same.
  - **Validity:** the chosen value was proposed by one of the nodes.
  - **Termination:** eventually they agree on a value (a.k.a., a value is “chosen”).

# Consensus vs atomic commitment

- **Consensus** : participants need to agree on a value, but they are willing and capable to accept any value.
  - Ex: A group's decision on where to meet (say, which specific room on campus of those that are of suitable size) can probably be cast as a consensus problem: most likely no one cares where they meet, but they all need to agree on the same value.
- Contrast with atomic commitment: participants need to agree on a value, but they have specific constraints on whether they can accept any particular value.
  - Ex: A group's decision on when to meet is probably an atomic commitment problem, because each participant has his/her own calendar constraints.

# Fischer-Lynch-Paterson Impossibility Result

- What FLP **says**: you can't guarantee both safety and progress when there is even a single fault at an inopportune moment
- What FLP **doesn't say**: in practice, how close can you get to the ideal (always safe and live)?
- Consensus protocols like **Paxos** get close in practice
  - The topic of this lecture

# Paxos

- The most popular fault-tolerant agreement protocol
  - Google Chubby (Paxos-based distributed lock service)
  - Google Spanner: geo-distributed transactional database
  - Yahoo Zookeeper (Paxos-based distributed lock service)
  - Open source: libpaxos (Paxos-based atomic broadcast)
- Paxos' properties: completely-safe and largely-live
  - **Safety**
    - If agreement is reached, everyone agrees on the same value. The value agreed upon was proposed by some node
  - **Fault tolerance** (i.e., as-good-as-it-gets liveness)
    - If less than half the nodes fail, the rest nodes reach agreement *eventually*
  - **No guaranteed termination** (i.e., imperfect liveness)
    - Paxos may not always converge on a value, but only in very degenerate cases that are improbable in the real world


Let us build up Paxos from scratch to understand where many consensus algorithms actually come from.

# Paxos: Starting point

- We assume a client-server configuration, with initially one primary server.
- To make the server more robust, we start with adding a backup server.
- To ensure that all commands are executed in the same order at both servers, the primary assigns unique sequence numbers/timestamps to all commands.
- In Paxos, the primary is called the **leader**.



# Two-server situation




- S: servers, C: client,  $O^n$ : operations/commands
- $\sigma_i^j$ : response from server  $i$  in state  $j$  expressed as sequence of ops carried out

# Handling lost messages: Paxos terminology

- Some Paxos terminology


- The leader sends an **accept** message  $\text{ACCEPT}(o, t)$  to backups when assigning a timestamp  $t$  to command  $o$ .
- A backup responds by sending a **learn** message:  $\text{LEARN}(o, t)$

# Two servers & one crash (1)



- Problem
  - Primary crashes after executing an operation, but the backup never received the accept message.

# Two servers & one crash (2)




- **Solution**


- **Never execute an operation before it is clear that it has been learned**
- When the leader notices that operation  $o$  has not yet been learned, it retransmits  $\text{ACCEPT}(o, t)$  with the original timestamp.

# Failure detection assumption

- Unrealistic assumption: Process can reliably detect crashes
- Only solution for failure detection in async. system: heartbeat
  - Each server sends out message “I’m alive”
  - Other servers set timeout on expected receipt of such messages
- But what happens if heartbeat is delayed?




# Required number of servers




- Observation
  - Paxos needs at least three servers
  - In Paxos with three servers, the leader cannot execute an operation  $o$  until it has received at least one (other) **LEARN**( $o$ ) message

# Required #servers in general case



- Scenario
  - A crash of majority of servers can leave the remaining non-faulty ones inconsistent


# Required #servers in general case



- Scenario
  - A crash of majority of servers can leave the remaining non-faulty ones inconsistent
- In general case to tolerate  $m$  failures, you need  $2m + 1$  servers
  - $m+1$  will give majority



# Scenario: Need for timestamp




- $S_3$  is completely ignorant of any activity by  $S_1$ 
  - Leader crashes after executing  $o^1$
  - $S_3$  never even received  $\text{ACCEPT}(o^1)$
  - $S_2$  received  $\text{ACCEPT}(o^1)$ , detects crash, and becomes leader.
  - $S_2$  sends  $\text{ACCEPT}(o^2)$
  - $S_3$  executes  $o^2 \Rightarrow$  Now,  $S_3$  has done  $o^2$  without doing  $o^1$

# Scenario: Leader crashes after executing $o^1$


- $S_3$  is completely ignorant of any activity by  $S_1$ 
  - $S_3$  never even received  $\text{ACCEPT}(o^1)$ .
  - $S_2$  received  $\text{ACCEPT}(o^1)$ , detects crash, and becomes leader.
  - $S_2$  sends  $\text{ACCEPT}(o^2)$
  - $S_3$  executes  $o^2 \Rightarrow$  Now,  $S_3$  has done  $o^2$  without doing  $o^1$
- Solution
  - Use timestamp to bring inconsistent server up to date
    - $S_2$  received  $\text{ACCEPT}(o^1, 1)$ , detects crash, and becomes leader.
    - $S_2$  sends  $\text{ACCEPT}(o^2, 2)$ 
      - Use a higher timestamp than what was seen before
    - $S_3$  sees unexpected timestamp and tells  $S_2$  that it missed  $o^1$ .
    - $S_2$  retransmits  $\text{ACCEPT}(o^1, 1)$ , allowing  $S_3$  to catch up.

# Need for broadcasting LEARN



- **Scenario:** What happens when  $\text{LEARN}(o^1)$  as sent by  $S_2$  to  $S_1$  is lost?
  - $S_1$  will not execute the operation but  $S_2$  will

# Need for broadcasting LEARN




- **Scenario:** What happens when  $LEARN(o^1)$  as sent by  $S_2$  to  $S_1$  is lost?
- **Solution:**  $S_2$  will also have to wait until it knows that  $S_3$  has learned  $o^1$ .
- **Paxos fundamental rule:** In Paxos, no server  $S$  (not just leader) can execute an operation  $o$  until it has received a  $LEARN(o)$  from a majority of non-faulty servers (2 in the 3 server case)

# Paxos: Detailed protocol

- Each node runs as a *proposer* and *acceptor*
- Proposer (leader) proposes a value and solicit acceptance from acceptors
- A value is *chosen* (a.k.a., consensus is reached) when a majority of acceptors have accepted it.
- A proposer announces a chosen value or tries again if it's failed to converge on a value.
- The protocol
  - guarantees *consistency* (all non-faulty nodes choose the same value)
  - guarantees *validity* (the chosen value was proposed by a proposer)
  - *ensures termination* (eventually, a value is chosen) with high probability but does not guarantee it.

# Paxos Proposal Numbers

- Each proposal has a unique number
  - Higher numbers take priority over lower numbers
  - It must be possible for a proposer to choose a new proposal number higher than anything it has seen/used before
- One simple approach:



- Each server stores maxRound: the largest Round Number it has seen so far
- To generate a new proposal number:
  - Increment maxRound
  - Concatenate with Server Id
- Proposers must persist maxRound on disk: must not reuse proposal numbers after crash/restart

# Paxos operation: node state

- Each node maintains:
  - $my_n$ : my proposal # in the current Paxos
  - $na$ : highest proposal # accepted
  - $va$ : corresponding accepted value
  - $nh$ : highest proposal # seen

# Paxos operation: 3P protocol

- Phase 1 (Propose)

- A node decides to be leader (and propose)
- Leader choose  $my_n > n_h$ 
  - Could be done by simply by incrementing global counter and adding server id
- Leader sends  $\langle \text{prepare}, my_n \rangle$  to all nodes
- Upon receiving  $\langle \text{prepare}, n \rangle$  acceptor does the following

    If  $n < n_h$

        reply  $\langle \text{prepare-reject} \rangle$

    Else

$n_h = n$

        reply  $\langle \text{prepare-ok}, n_a, v_a \rangle$

This node will not accept  
any proposal lower than  $n$



# Paxos operation: 3P protocol

- Phase 2 (Accept):
  - If leader gets prepare-ok from a majority
    - $V$  = non-empty value corresponding to the highest  $n_a$  received
    - If  $V = \text{null}$ , then leader can pick any  $V$
    - Send  $\langle \text{accept}, m_n, V \rangle$  to all nodes
  - If leader fails to get majority prepare-ok
    - Delay and restart Paxos
  - Upon receiving  $\langle \text{accept}, n, V \rangle$  acceptor does following
    - If  $n < n_h$ 
      - reply with  $\langle \text{accept-reject} \rangle$
    - else
      - $n_a = n$ ;  $v_a = V$ ;  $n_h = n$
      - reply with  $\langle \text{accept-ok} \rangle$

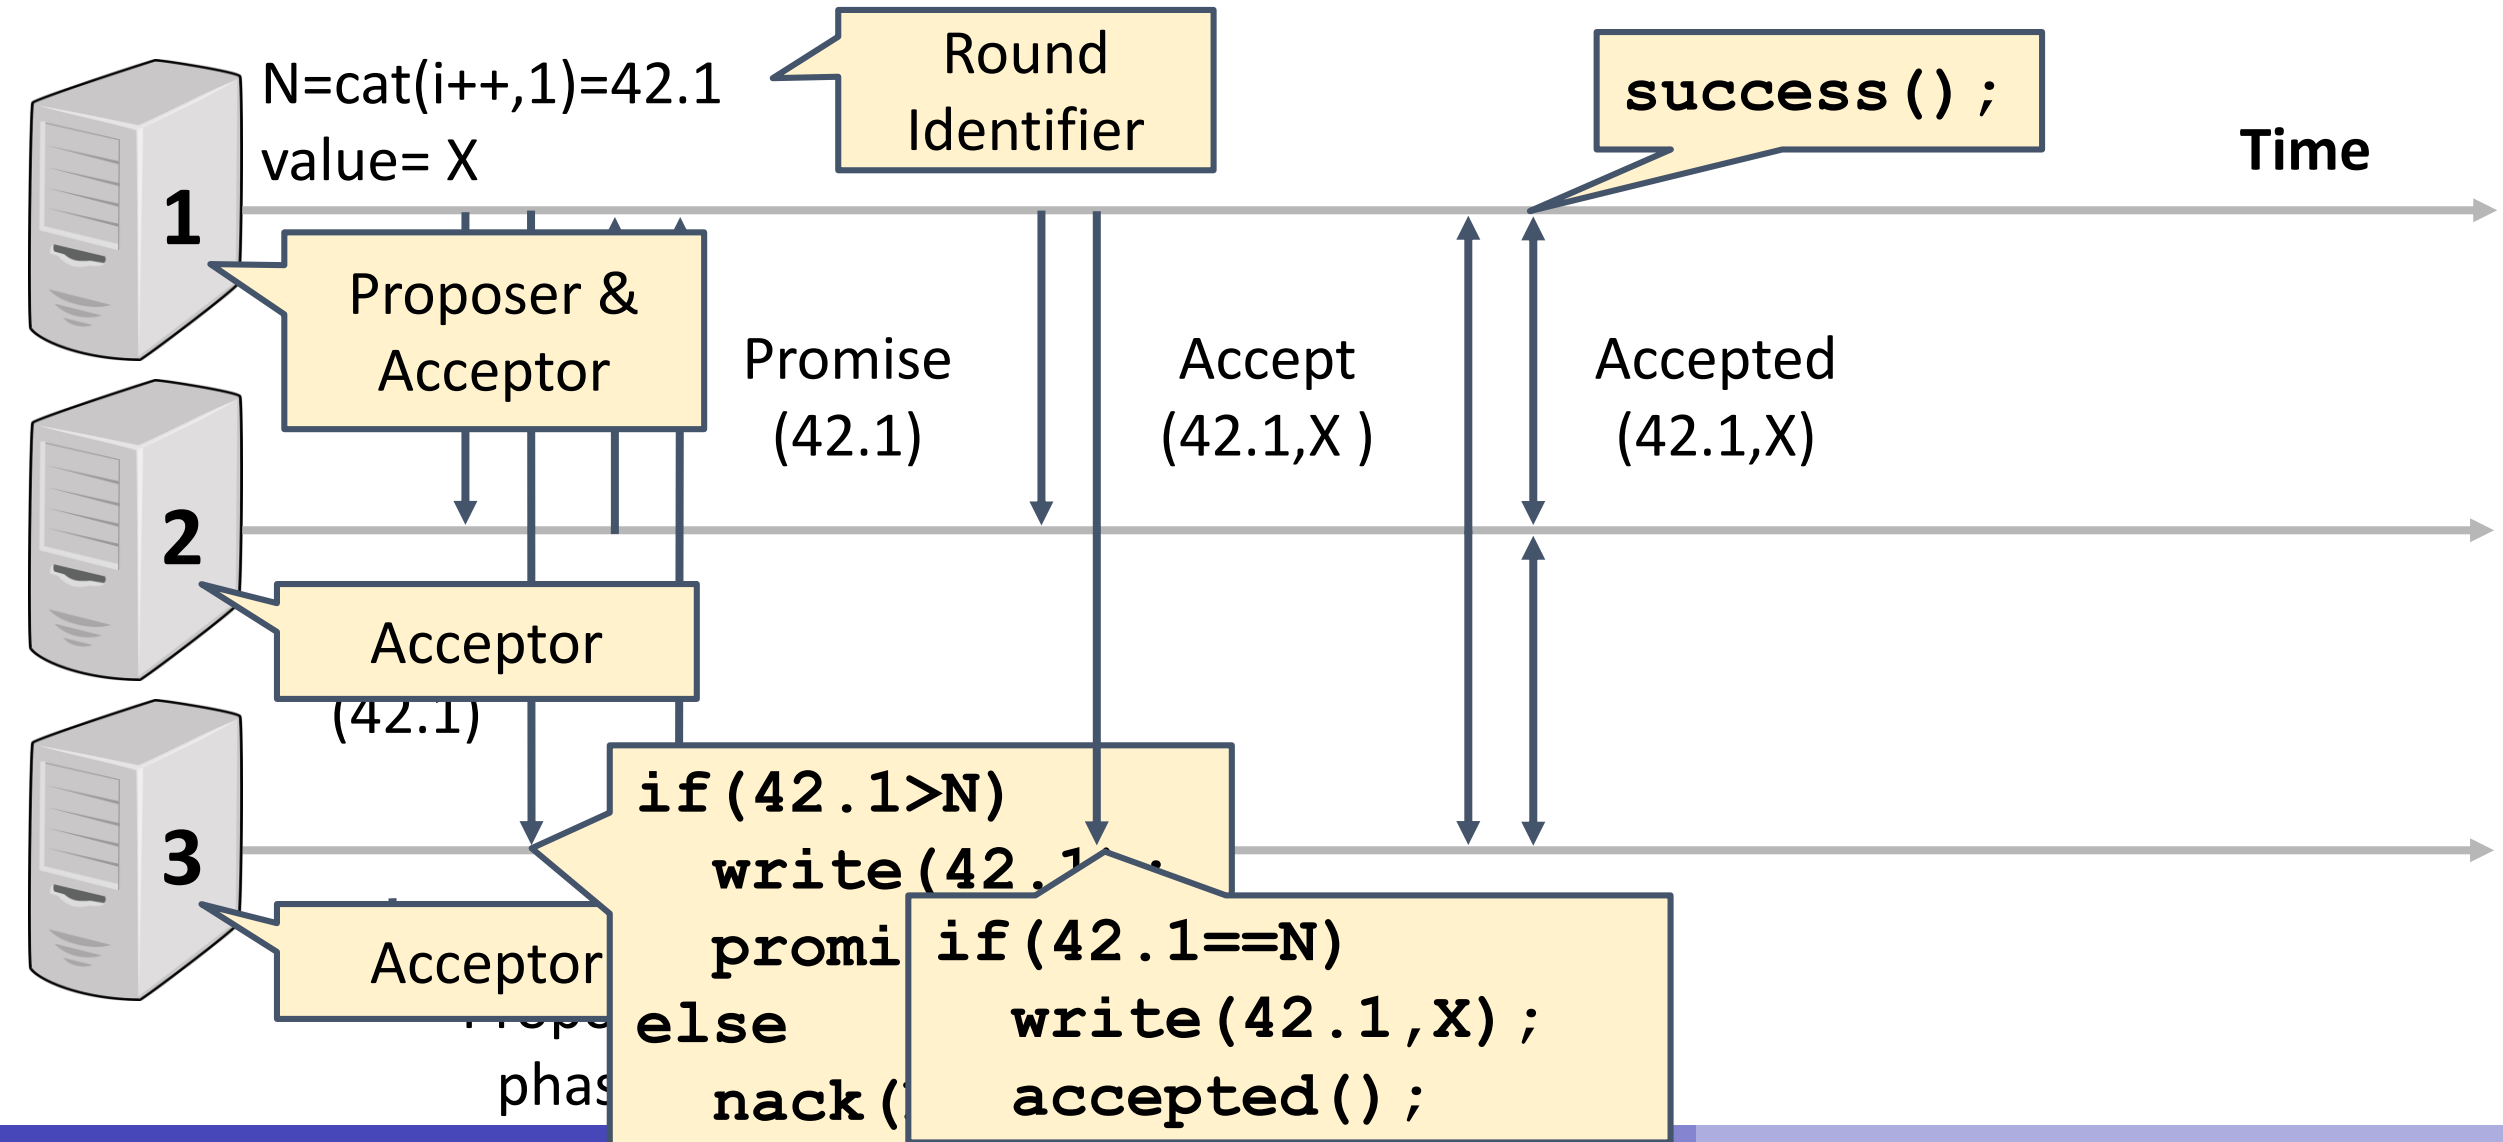
A point in the past, but its proposer didn't quite finish his job, then that value will remain in perpetuity

**So: newer proposers win the rounds, but with old proposers' values!!!**

# Paxos operation: 3P protocol

- Phase 3 (Decide)
  - If leader gets accept-ok from a majority
    - Return Done to client
    - Send  $\langle \text{DECIDE}, V_a \rangle$  to all nodes until you get  $\langle \text{DECIDE}, \text{OK} \rangle$  back
      - This phase is so that nodes close the protocol, and so that nodes that might not have heard previous ACCEPT messages learn the chosen value.
  - If leader fails to get accept-ok from a majority
    - Delay and restart Paxos

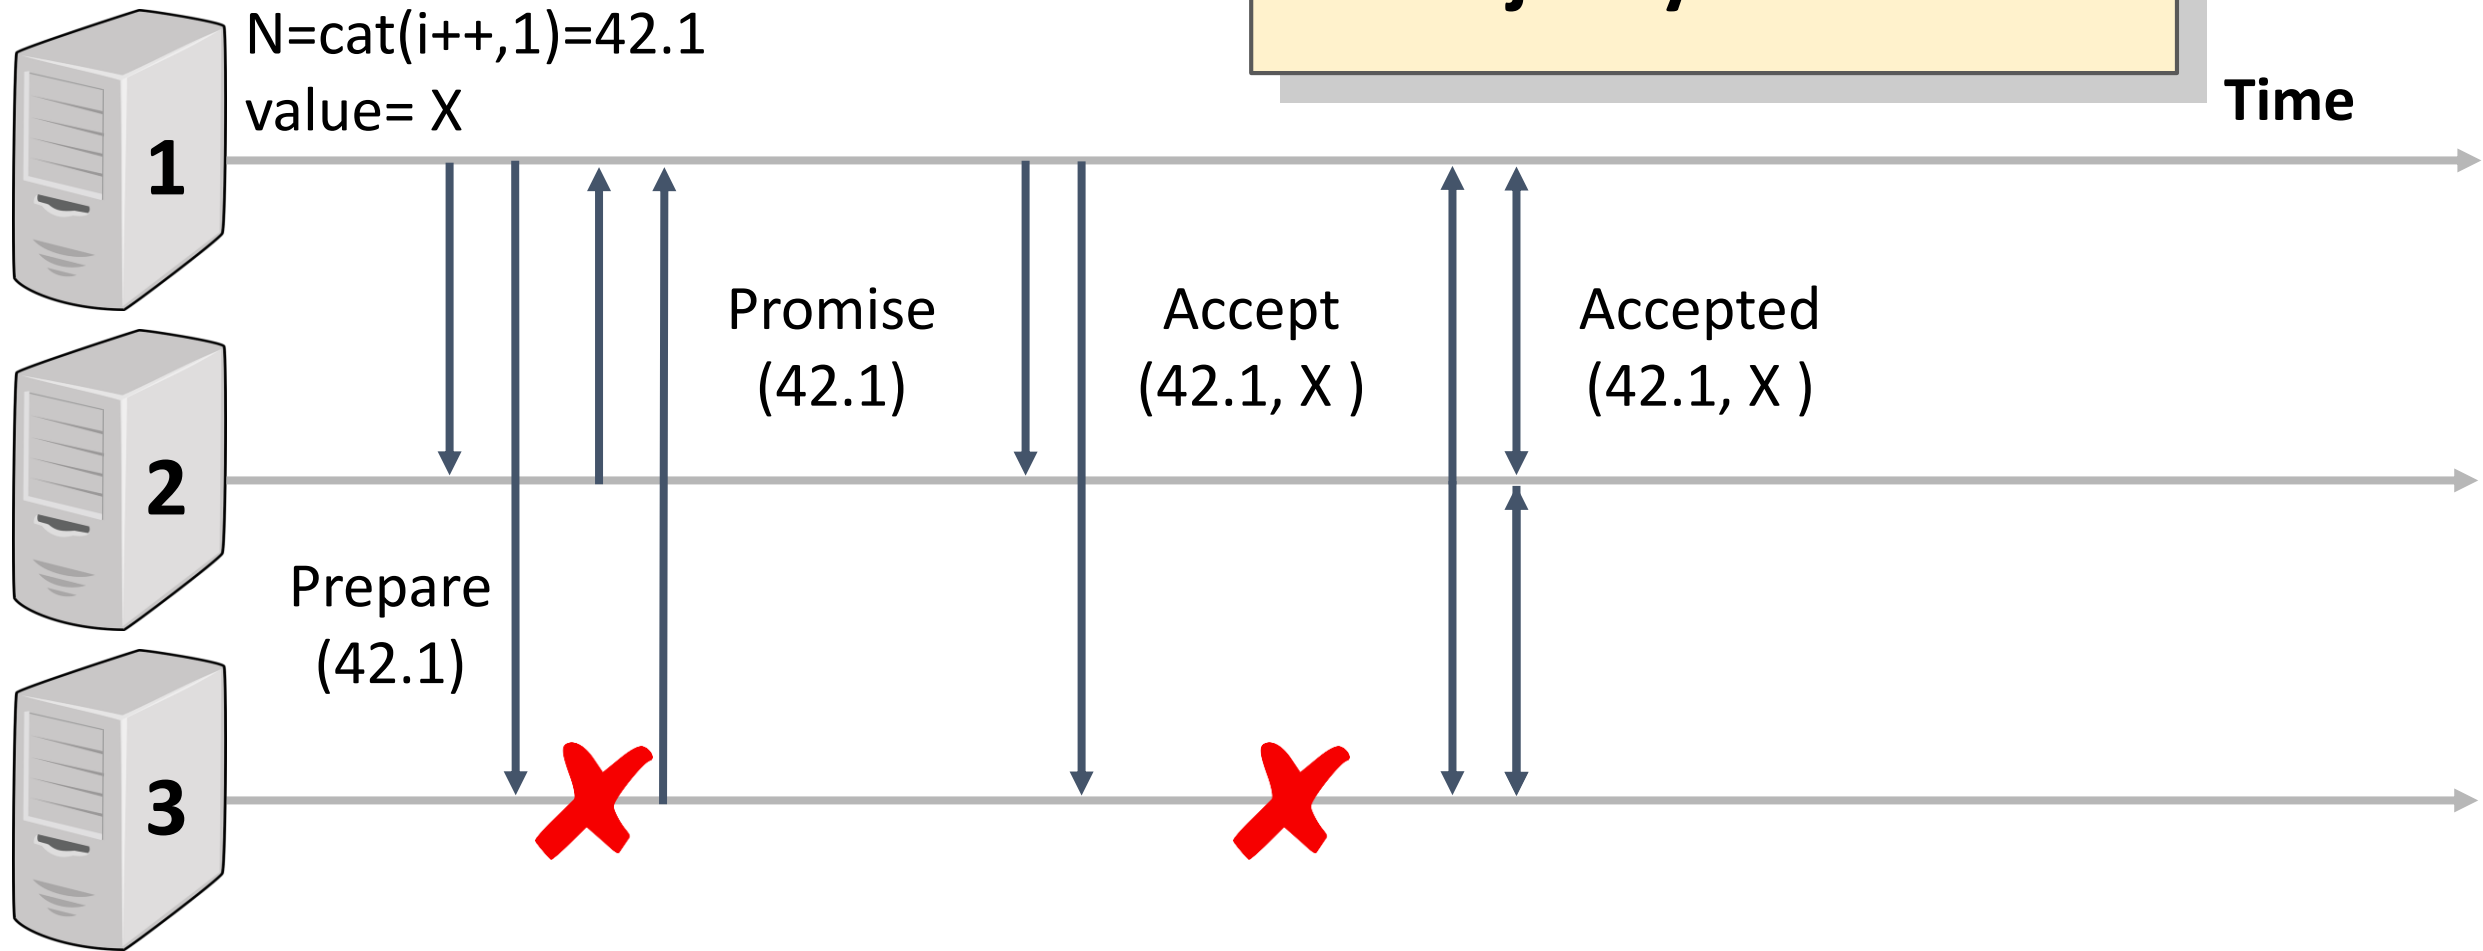
# Basic Paxos (Example without phase 3)




# Failures: Acceptor

**No problem as long as  
majority don't fail**

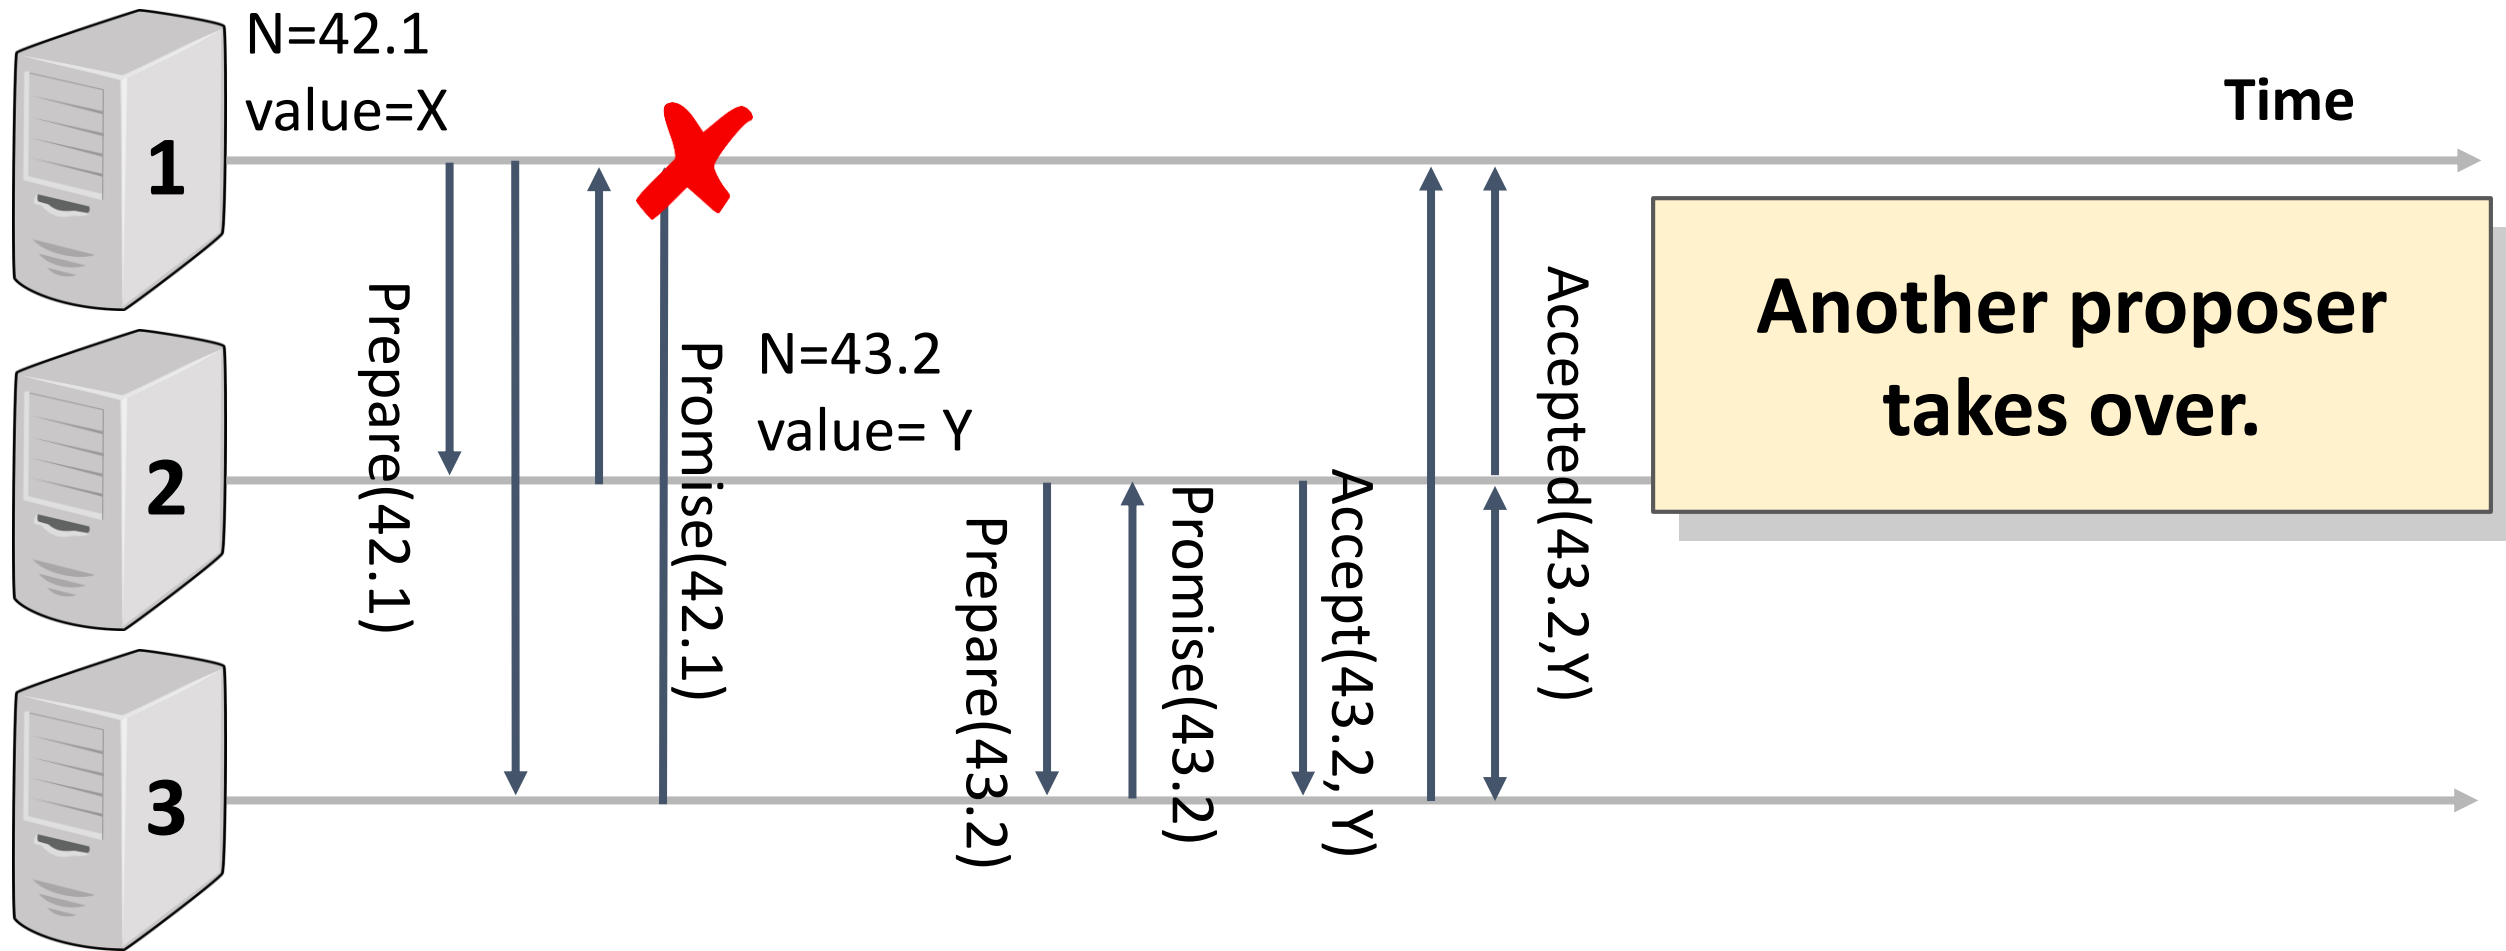
Time




# Failures: Proposer in Prepare Phase



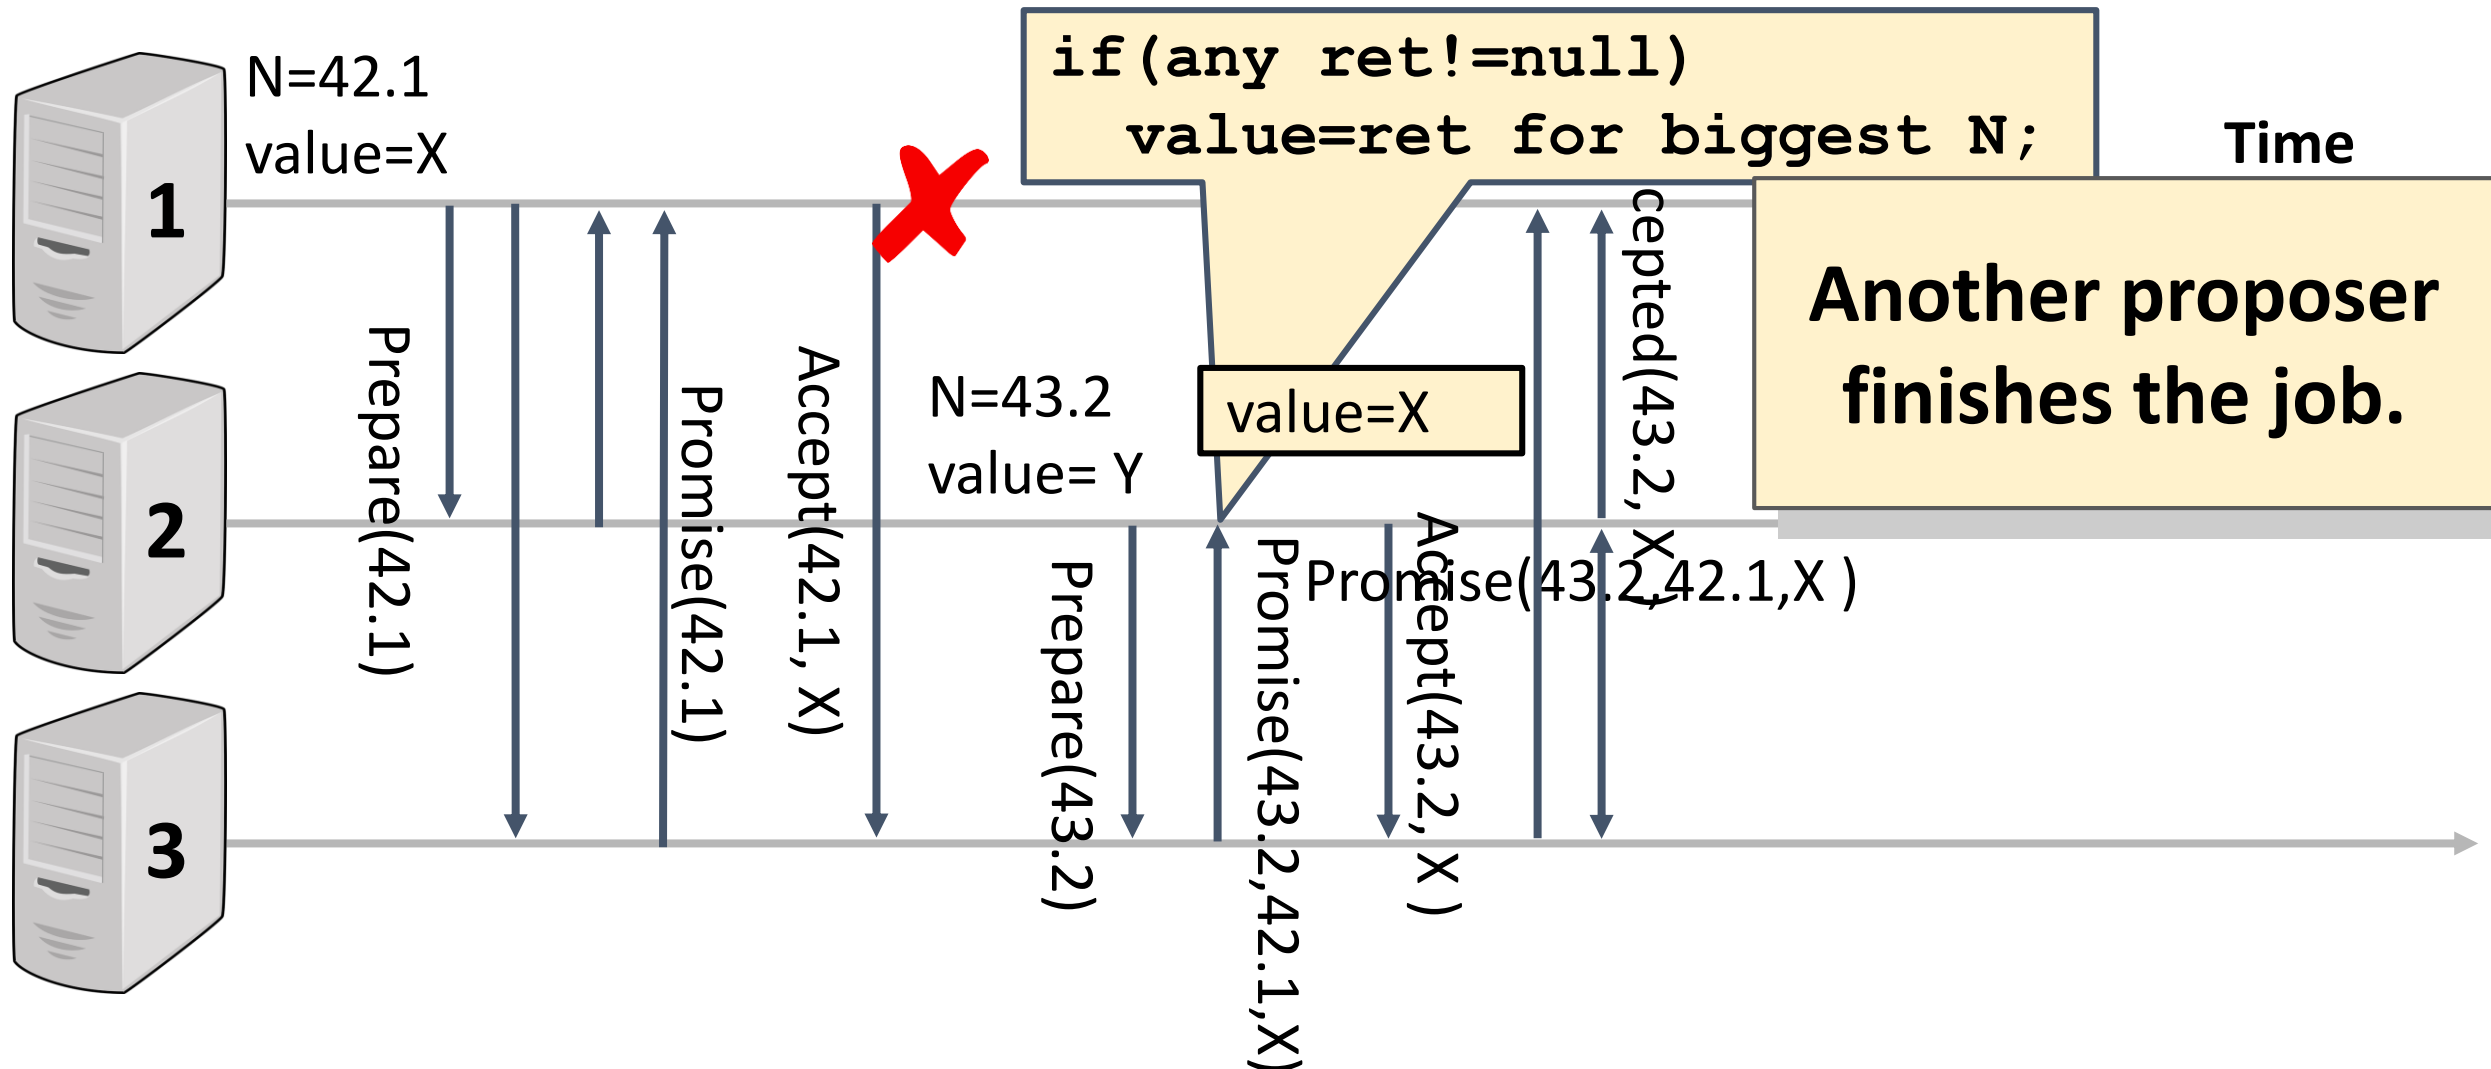
# Failures: Proposer in Prepare Phase



# Failures: Proposer in Accept Phase




# Failures: Proposer in Accept Phase





# What could go wrong?

- One proposer
  - One or more acceptors fails
    - Still works as long as majority are up
  - Proposer fails in prepare phase
    - No-op; another proposer can make progress
  - Proposer fails in accept phase
    - Another proposer overwrites
    - Another proposer finishes the job
- Two or more “simultaneous” proposers
  - A bit more complex and Paxos can come to a grinding halt due to livelock between proposers (no proposer is able to get majority)!



# Paxos in the real world

- Solving livelock with explicit leader election (ex: Bully algorithm)
- Developer: Creating a log of agreements (make more than one decision)
  - Multi-Paxos
- Sysadmin work: Adding and removing nodes from Paxos
  - Naming and cluster membership
- Software tester: Testing and debugging is hard
- Researcher: Do you want to handle byzantine failures?
  - What happens if  $N=\infty$  due to disk corruption?

**Need to implement? Don't! Just use a library!**  
**Or use RAFT**

# Performance

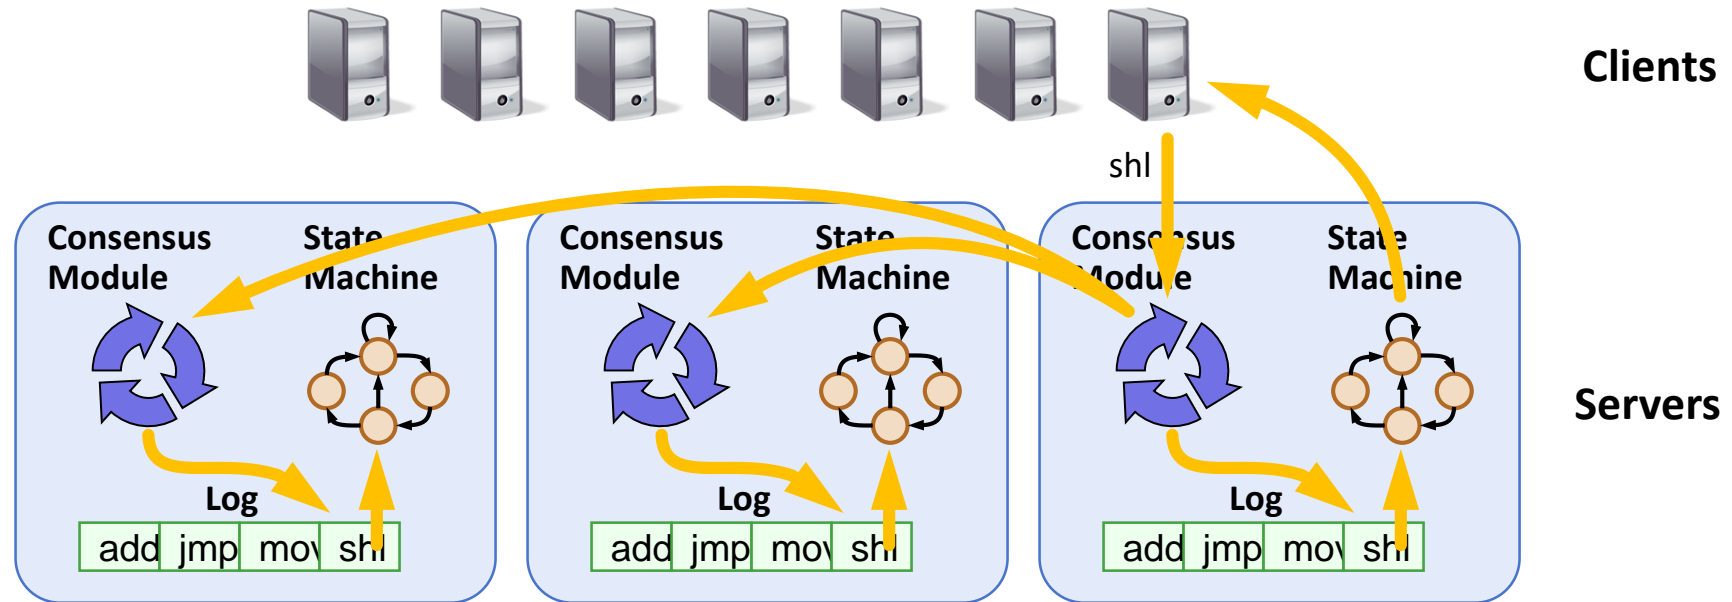
- Basic Paxos latency:
  - 2 network round trips +
  - 2 disk writes
- Multi-Paxos latency:
  - Converges on 1 network round trip + 1 disk write
    - Assuming failures are rare
- Improving reliability: want multiple sites!
  - But that implies cross-rack/datacenter/zone/continent messaging
- Bandwidth can be increased through:
  - Batching operations into one Paxos instance
    - Each paxos round completes multiple transactions
  - Parallel Paxos instances

**Trade-off: more latency for more reliability.**

# Putting together Paxos + 2PC: Sharding & Replication

- We talked about single node database
  - ACID properties of transactions, Isolation with 2PL
- In practice, databases are distributed
  - Data sharded/partitioned for *scalability*
  - We talked about 2PC for atomicity across shards
- In practice, shards are replicated for fault tolerance
  - Each database holding replica typically uses a write-ahead log to guarantee durability (D of ACID)
  - How do you make sure all replicas update the log in the same order?

# Practical Paxos Use Case: Replicated Log



- All servers execute same commands in same order
- Consensus module ensures proper log replication
- System makes progress as long as any majority of servers are up

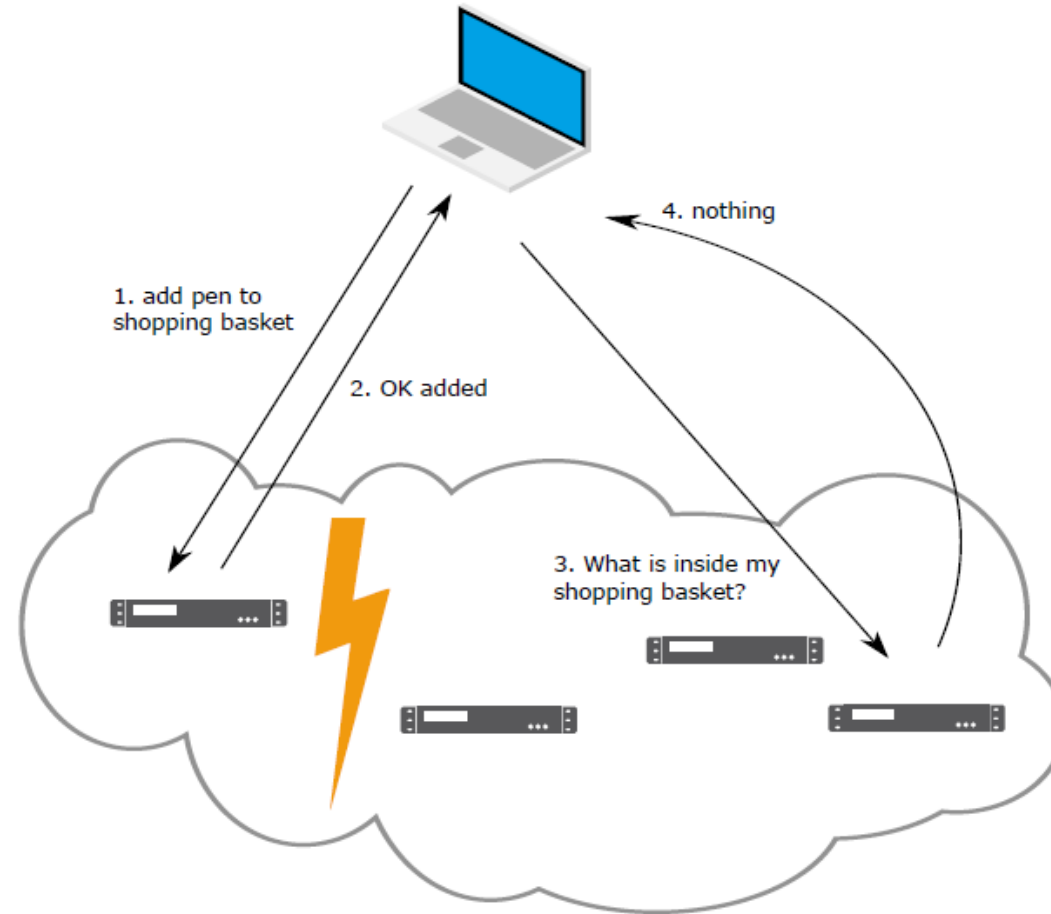
# Example: Google Spanner

- A database system with millions of nodes, petabytes of data, distributed across datacenters worldwide
- Consistency properties:
  - Serializable transaction isolation
  - Many shards, each holding a subset of the data; atomic commit of transactions across shards
- Many standard techniques:
  - Paxos replication within a shard
  - Two-phase locking for serializability
  - Two-phase commit for cross-shard atomicity

# Paxos vs. 2/3PC

- Remember:
  - 2PC was vulnerable to 1-node failures, especially coordinator failures
  - 3PC was vulnerable to network partitions
- Paxos deals with these issues using two mechanisms:
  - **Egalitarian consensus**: no node is special, anyone can take over as coordinator at any time
    - Hence, if one coordinator fails, another one will time out and take over
    - But that requires special ordering and acceptance protocols for proposals
  - **Safe majorities**: instead of requiring all participants to answer Yes, Paxos requires only half + 1 of the nodes
    - Because you **cannot have two simultaneous majorities**, which avoids partitions
- But, If you don't have a majority of non-faulty nodes, Paxos will prioritize consistency over availability
  - Writes will not succeed.

# CAP: Why not all three?



If there is a network partition either C or A will break



# CAP Theorem

- First stated by Eric Brewer (Berkeley) at the PODC 2000 keynote
  - Formally proved by Gilbert and Lynch in 2002
- Consistency (specifically Linearizability)
  - Linearizability = sequential consistency + real-time constraint
- Availability
  - a system is available if every request to a non-failing node always receives a response, eventually
- Partition tolerance
  - The system continues to operate despite an arbitrary number of messages being dropped (or delayed) by the network between nodes
- The theorem says: between Consistency, Availability, Partition tolerance, you can choose only two

# Design Tradeoff

- Network partitions occur outside anyone's control in real life
  - Cannot sacrifice the Partition-Tolerance property
- In the event of a network partition either A or C is maintained: it is the choice of the designer
- Practical distributed systems are CP or AP
  - CP oriented: BigTable, Hbase, MongoDB, Redis, MemCached, ...
  - AP oriented: Amazon Dynamo, CouchDB, Cassandra, ...
- Thus The CAP theorem formally states the trade-offs among different distributed systems properties

# Closing

- We have covered a lot of ground
  - Cloud computing & service models
  - Virtualization, Docker: Infrastructure fundamentals
  - MapReduce, Spark: Programming fundamentals
  - RDMBS, NFS, AFS, GFS: Data management fundamentals
  - Consistency & fault tolerance: Distributed systems fundamentals
- You have all the tools to build a cloud application
- Cloud jobs are in demand. Go put your theoretical knowledge into practice with AWS, Azure, or GCP

## **The top 10 most in-demand hard skills globally**

1. Blockchain
2. Cloud computing
3. Analytical reasoning
4. Artificial intelligence

# Acknowledgements

- This course was built with material from
  - CMU cloud computing, distributed systems courses by Prof. Majd Sakr, Prof. Dave Anderson, Prof. Satya Mahadev
  - Columbia distributed systems course by Prof. Roxana Geambasu
  - NYU distributed systems by Prof. Jinyang Li
  - MIT distributed systems course by Prof. Robert Morris
  - Distributed systems book by Andy Tanenbaum & Maarten van Steen
- The final project were made possible by Eugenio Marinelli
- Finally, this could would not have been possible without you guys! Thank you!

Good luck!  
Maybe your future be cloudy with lots of  
(vegan) meatballs!