

Cloud Data Management: Relational Databases vs MapReduce

Lecture 5

Recap

- MapReduce introduced by Google
 - Simple programming model for building distributed applications that process vast amounts of data
 - Runtime for executing jobs on large clusters in a reliable, fault-tolerant manner
- Hadoop makes MapReduce broadly available
 - HDFS becomes central data repository
 - Becomes Defacto standard for batch processing

New applications, new workloads

- MapReduce originally designed for batch analytics
 - Latency-insensitive: jobs that run for hours
 - Sequential scans of Petabytes of data
 - Built for fault tolerance across thousands of commodity servers
 - Focus on faults during query rather than recovery after updates
- Hadoop starts being used for interactive computations, e.g., ad-hoc analytics
 - Hive and Pig frameworks with SQL-over-Hadoop drive this trend
- But SQL and analytics was the stronghold of relational database engines

“MapReduce: A major step backwards” – Dewitt, Stonebraker

Role of a database system

- Database: **integrated, shared** data collection
- Integrated
 - Eliminate needless redundancy
 - Maintain strong consistency
- Shared
 - Application written by programmers in multiple languages
 - End-users who use applications, forms, CLI to interact
- Database systems shield users from
 - How data is stored (bits & bytes, 1 vs N files, 1 vs N disks...)
 - How data is accessed (btree, hashtable, scan, ...)

What is a data model?

- Collection of application-visible constructs
 - Describe data in application & storage agnostic way
- Constructs to describe structural aspects
 - How do applications perceive the data?
 - Ex: table, graph, associative array...
- Constructs to describe manipulation aspects
 - What operators can applications use?
 - Ex: join, traverse, lookup...
- Constructs to describe data integrity aspects
 - How do we ensure that data manipulation is “correct”?

Relational Model: Structural aspect

- Database = set of named **relations** (or **tables**)
- Each relation has a set of named **attributes** (or **columns**)
- Each **tuple** (or **row**) has a value for each attribute
- Each attribute has a **type** (or **domain**)
 - integer, real, string, file formats (jpeg,...), enumerated and many more

Students

sid	name	login	age	gpa
50000	Dave	dave@cs	19	3.3
53666	Jones	jones@cs	18	3.4
53688	Smith	smit@ee	18	3.2
...

Colleges

name	location	strength
MIT	USA	10000
Oxford	UK	22000
EPFL	CH	9000
...

Relational Model: Structural aspect

- **Relation Schema**: relation name + field names + field domains
 - Students(*sid*: string, *name*: string, *login*: string, *age*: integer, *gpa*: real)
- **Relation Instance**: contents at a given point in time
 - *set* of rows or *tuples*. (all rows are distinct with no specific ordering)
 - Cardinality: # rows, Arity or degree: # attributes
- **Database Schema**: collection of relation schemas
- **Database Instance**: collection of relation instances

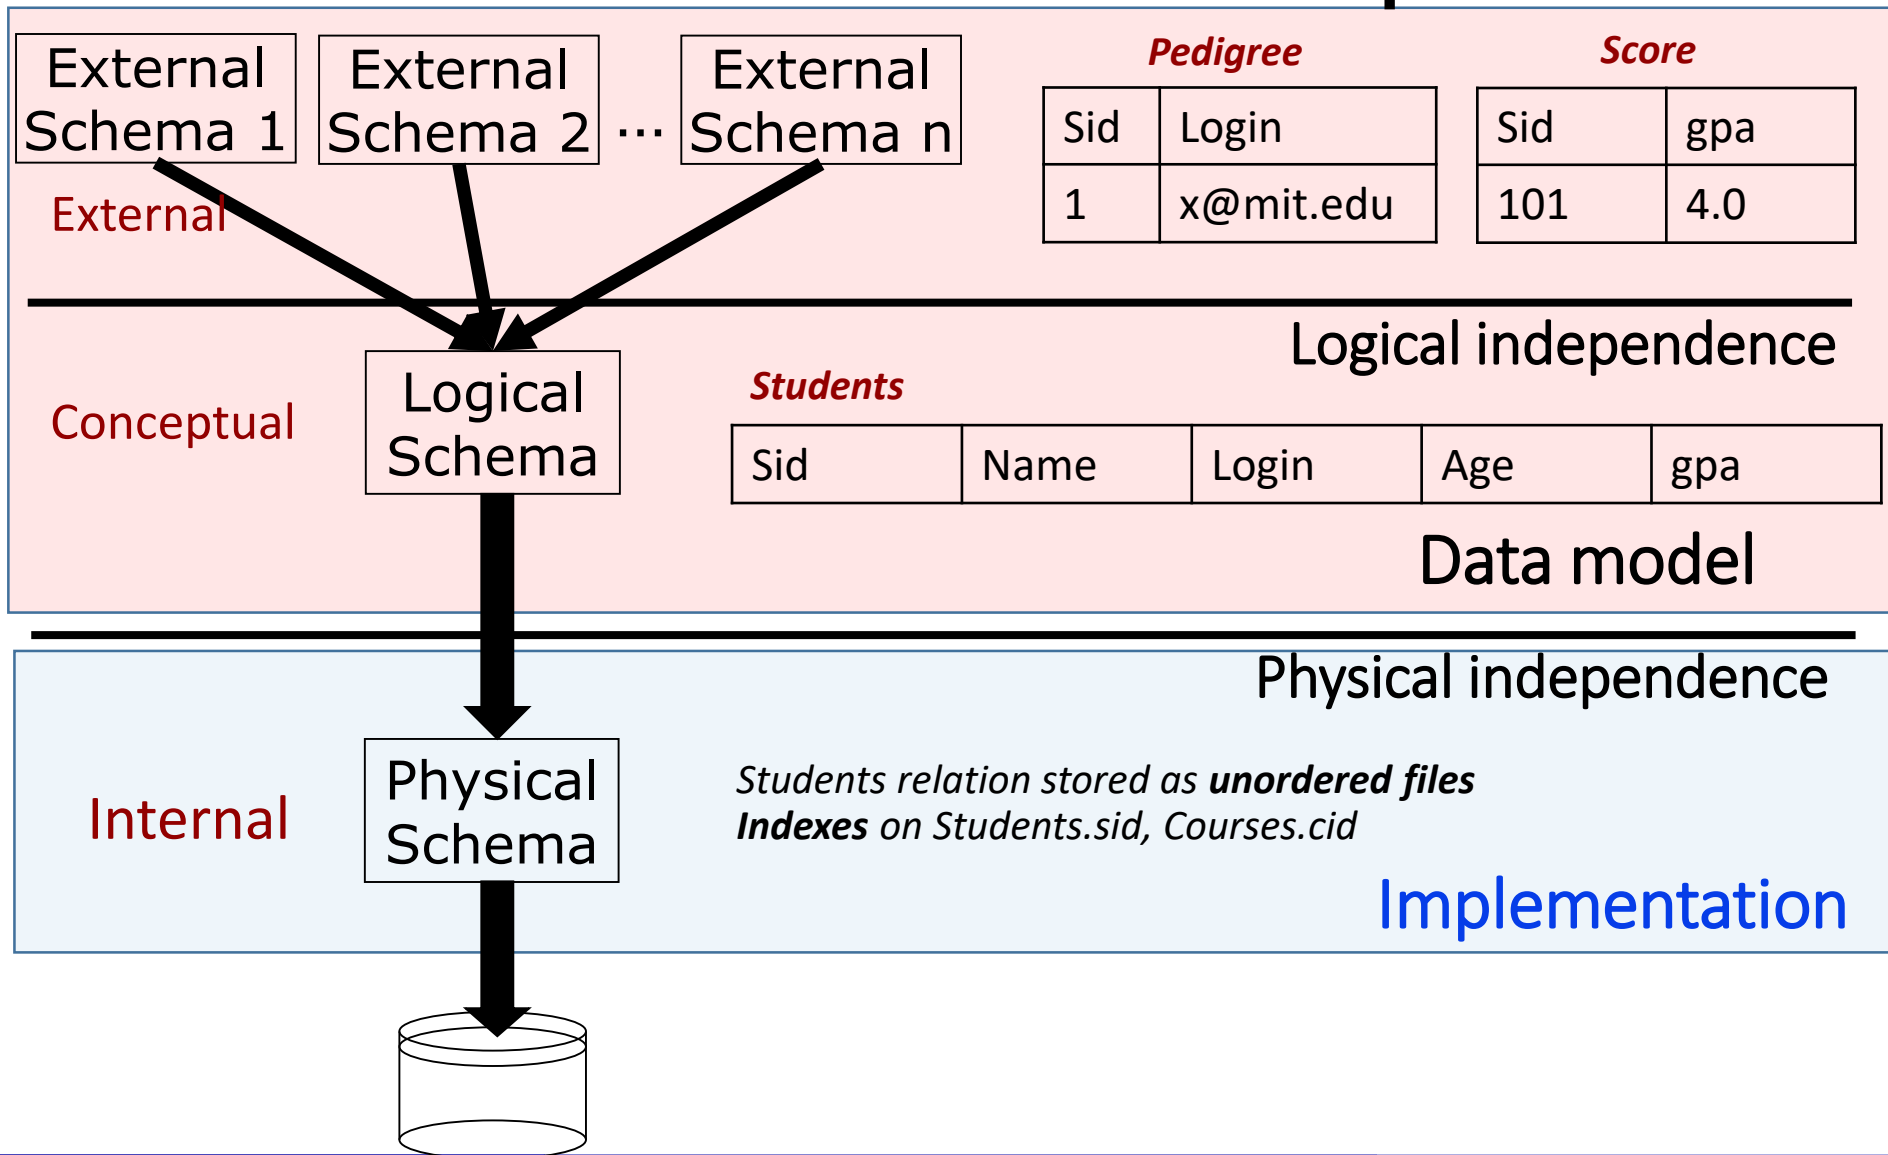
Students

sid	name	login	age	gpa
50000	Dave	dave@cs	19	3.3
53666	Jones	jones@cs	18	3.4
53688	Smith	smit@ee	18	3.2
...

Colleges

name	location	strength
MIT	USA	10000
Oxford	UK	22000
EPFL	CH	9000
...

Relational model & data independence



Relational Model: Integrity Aspect

- Relational model provides **Integrity Constraints**
 - condition specified on schema that restricts the data that can be stored in **any** instance
 - ICs are specified when schema is defined.
 - ICs are checked when relations are modified.
- A **legal** instance of a relation is one that satisfies all specified ICs
 - DBMS should not allow illegal instances.
- With ICs, stored data is more faithful to real-world meaning
 - Avoids data entry errors, too!

Relational Model: **Keys**

- Attribute whose value is unique in each tuple
- Or set of attributes whose combined values are unique
- Keys specify **key constraint**
 - Enforced when tuples are inserted/updated

Students

sid	name	login	age	gpa
50000	Dave	dave@cs	19	3.3
53666	Jones	jones@cs	18	3.4
53688	Smith	smit@ee	18	3.2
...

Colleges

name	location	strength
MIT	USA	10000
Oxford	UK	22000
EPFL	CH	9000
...

Relational Model: **Foreign Keys**


- Set of fields in one relation that `refer' to a tuple in another relation (like a pointer)
- Foreign keys specify **Foreign Key Constraint**
 - FK must correspond to the primary key of the other relation
- If all foreign key constraints are enforced, **referential integrity** is achieved (i.e., no dangling references.)

Students

sid	name	login	age	gpa
50000	Dave	dave@cs	19	3.3
53666	Jones	jones@cs	18	3.4
53688	Smith	smit@ee	18	3.2
...

Enrolled

cid	sid	grade
Carnatic101	53666	C
Raggae203	50000	B
Topology112	53666	A
...



Relational Model: Manipulation Aspect

- **Query languages:** Allow manipulation and **retrieval of data** from a database.
- Relational model supports simple, powerful QLs:
 - Strong formal foundation based on logic.
 - Allows for much optimization.
- Two mathematical Query Languages form the basis for “real” languages (e.g. SQL), and for implementation:
 - **Relational Algebra:** More **operational**, very useful for representing execution plans.
 - **Relational Calculus:** Lets users describe what they want, rather than how to compute it. (**Non-procedural, declarative.**)

Preliminaries

- A query is applied to *relation instances*, and the result of a query is also a relation instance.
 - *Schemas of input* relations for a query are *fixed* (but query will run over any legal instance)
 - The *schema for the result* of a given query is also *fixed*. It is determined by the definitions of the query language constructs.

Example Schema and Instances

- Boats(bid: integer, bname: string, color: string)
- Sailors(sid: integer, sname: string, rating: integer, age: real)
- Reserves(sid: integer, bid: integer, day: date)

Boats

<u>bid</u>	bname	color
101	Interlake	blue
102	Interlake	red
103	Clipper	green
104	Marine	red

R1

<u>sid</u>	<u>bid</u>	day
22	101	10/10/96
58	103	11/12/96

S1

<u>sid</u>	sname	rating	age
22	Dustin	7	45.0
31	Lubber	8	55.5
58	Rusty	10	35.0

S2

<u>sid</u>	sname	rating	age
28	yuppy	9	35.0
31	Lubber	8	55.5
44	guppy	5	35.0
58	Rusty	10	35.0

Relational Algebra: 5 Basic Operations

- **Selection** (σ) Selects a subset of *rows* from relation (horizontal).
- **Projection** (π) Retains only wanted *columns* from relation (vertical).
- **Cross-product** (\times) Allows us to combine two relations.
- **Set-difference** ($-$) Tuples in r_1 , but not in r_2 .
- **Union** (\cup) Tuples in r_1 and/or in r_2 .

Since each operation returns a relation, *operations can be composed!* (Algebra is “closed”).

Selection Operator: (σ)

- Selects rows that satisfy *selection condition*.
- **Output schema** of result is same as that of the input relation **S2**

<u>sid</u>	sname	rating	age
28	yuppy	9	35.0
31	Lubber	8	55.5
44	guppy	5	35.0
58	Rusty	10	35.0

$\sigma_{rating < 9}(S2)$

Output

<u>sid</u>	sname	rating	age
31	Lubber	8	55.5
44	guppy	5	35.0

S2

<u>sid</u>	sname	rating	age
28	yuppy	9	35.0
31	Lubber	8	55.5
44	guppy	5	35.0
58	Rusty	10	35.0

$\sigma_{rating < 9 \wedge age > 50}(S2)$

Output

<u>sid</u>	sname	rating	age
31	Lubber	8	55.5

Projection Operator (π)

- Retains only attributes that are in the *projection list*.
- **Output schema** is exactly the fields in the projection list, with the same names that they had in the input relation.

S2

<u>sid</u>	sname	rating	age
23	yuppy	9	35.0
31	Lubber	8	55.5
44	guppy	5	35.0
53	Rusty	10	35.0

$\pi_{sname, rating}(S2)$

Output

sname	rating
yuppy	9
Lubber	8
guppy	5
Rusty	10

Projection Operator (π): Duplicate Elimination

- Relational algebra is set based while SQL is bag (multiset) based
- Projection operator *eliminates duplicates*

S2

<u>sid</u>	sname	rating	age
23	yuppy	9	35.0
31	Lubber	8	55.5
44	guppy	5	35.0
53	Rusty	10	35.0

$\pi_{age}(S2)$

Output

age
35.0
55.5

Composing multiple operators

- Output of one operator can become input to another operator

S2

<u>sid</u>	sname	rating	age
23	yuppy	9	35.0
31	Lubber	8	55.5
44	guppy	5	35.0
53	Rusty	10	35.0

Output

sname	rating
yuppy	9
Rusty	10

$$\pi_{sname, rating} \left(\sigma_{rating > 8}(S2) \right)$$

Union and Set-Difference

- All of these operations take two input relations, which must be *union-compatible*:
 - Same number of fields.
 - “Corresponding” fields have the same type.

Union operator (U)

$S1$

<u>sid</u>	sname	rating	age
22	Dustin	7	45.0
31	Lubber	8	55.5
58	Rusty	10	35.0

$S2$

<u>sid</u>	sname	rating	age
28	yuppy	9	35.0
31	Lubber	8	55.5
44	guppy	5	35.0
58	Rusty	10	35.0

$S1 \cup S2$

<u>sid</u>	sname	rating	age
22	Dustin	7	45.0
31	Lubber	8	55.5
58	Rusty	10	35.0
44	guppy	5	35.0
28	yuppy	9	35.0

Set Difference Operator (-)

S1

<u>sid</u>	sname	rating	age
22	Dustin	7	45.0
31	Lubber	8	55.5
58	Rusty	10	35.0

S2

<u>sid</u>	sname	rating	age
28	yuppy	9	35.0
31	Lubber	8	55.5
44	guppy	5	35.0
58	Rusty	10	35.0

S1 - S2

<u>sid</u>	sname	rating	age
22	Dustin	7	45.0

S2 - S1

<u>sid</u>	sname	rating	age
28	yuppy	9	35.0
44	guppy	5	35.0

Cross-Product (\times)

- $S1 \times R1$: Each row of $S1$ paired with each row of $R1$.
- **Result schema** has one field per field of $S1$ and $R1$, with field names “inherited” if possible.
 - *May have a naming conflict*: Both $S1$ and $R1$ have a field with the same name.
 - In this case, can use the *renaming operator*:

$$\rho(C(1 \rightarrow sid1, 5 \rightarrow sid2), S1 \times R1)$$

Call C the result of $S1 \times R1$ and respectively rename the 1st & 5th fields of C to $sid1$ & $sid2$

Cross-Product Example

S1

<u>sid</u>	sname	rating	age
22	Dustin	7	45.0
31	Lubber	8	55.5
58	Rusty	10	35.0

R1

<u>sid</u>	<u>bid</u>	day
22	101	10/10/96
58	103	11/12/96

S1* × *R1

sid	sname	rating	age	sid	bid	day
22	Dustin	7	45.0	22	101	10/10/96
22	Dustin	7	45.0	58	103	11/12/96
31	Lubber	8	55.5	22	101	10/10/96
31	Lubber	8	55.5	58	103	11/12/96
58	Rusty	10	35.0	22	101	10/10/96
58	Rusty	10	35.0	58	103	11/12/96

$\rho_{1 \rightarrow sid1, 5 \rightarrow sid2}(S1 \times R1)$

sid1	sname	rating	age	sid2	bid	day
22	Dustin	7	45.0	22	101	10/10/96
22	Dustin	7	45.0	58	103	11/12/96
31	Lubber	8	55.5	22	101	10/10/96
31	Lubber	8	55.5	58	103	11/12/96
58	Rusty	10	35.0	22	101	10/10/96
58	Rusty	10	35.0	58	103	11/12/96

Compound Operator: Join

- Joins are compound operators involving cross product, selection, and (sometimes) projection.
- Most common type of join is a **natural join** (often just called “join”). $R \bowtie S$ conceptually is:
 - Compute $R \times S$
 - Select rows where attributes that appear in both relations have equal values
 - Project all unique attributes and one copy of each of the common ones.
- Note: Usually done much more efficiently than this.
- Useful for putting “normalized” relations back together.

Natural Join Example

$$\pi_{S1.sid, sname, \dots} (\sigma_{S1.sid=R1.sid} (S1 \times R1))$$

S1

<u>sid</u>	sname	rating	age
22	Dustin	7	45.0
31	Lubber	8	55.5
58	Rusty	10	35.0

R1

<u>sid</u>	<u>bid</u>	day
22	101	10/10/96
58	103	11/12/96

sid	sname	rating	age	sid	bid	day
22	Dustin	7	45.0	22	101	10/10/96
22	Dustin	7	45.0	58	103	11/12/96
31	Lubber	8	55.5	22	101	10/10/96
31	Lubber	8	55.5	58	103	11/12/96
58	Rusty	10	35.0	22	101	10/10/96
58	Rusty	10	35.0	58	103	11/12/96

S1 ⋈ R1

sid	sname	rating	age	bid	day
22	Dustin	7	45.0	101	10/10/96
58	Rusty	10	35.0	103	11/12/96

Condition Join or Theta-Join

$$R \bowtie_C C = \sigma_C(R \times S)$$

- **Output schema** same as that of cross-product.
- May have fewer tuples than cross-product.

S1

<u>sid</u>	sname	rating	age
22	Dustin	7	45.0
31	Lubber	8	55.5
58	Rusty	10	35.0

R1

<u>sid</u>	<u>bid</u>	day
22	101	10/10/96
58	103	11/12/96

$$S1 \bowtie_{S1.sid < R1.sid} R1$$

sid	sname	rating	age	sid	bid	day
22	Dustin	7	45.0	58	103	11/12/96
31	Lubber	8	55.5	58	103	11/12/96

Equi-Join

- **Special case of theta-join**: condition c contains only conjunction of *equalities*.
- Find **all pairs** of sailors in $S2$ who have same age.

$S2$

<u>sid</u>	sname	rating	age
28	yuppy	9	35.0
31	Lubber	8	55.5
44	guppy	5	35.0
58	Rusty	10	35.0

• $S1 \bowtie_{S1.age=S2.age} S2$

• $\sigma_{sid1 \neq sid2} \left(S1 \bowtie_{S1.age=S2.age} (S2) \right)$

Grouping and Aggregation

- Grouping and Aggregation: $\gamma_X (R)$
 - Given a relation R , partition its tuples according to their values in one set of attributes G
 - The set G is called the **grouping attributes**
 - Then, for each group, aggregate the values in certain other attributes
 - Aggregation functions: SUM, COUNT, AVG, MIN, MAX, ...
- In the notation, X is a list of elements that can be:
 - A grouping attribute
 - An expression $\theta(A)$, where θ is one of the (five) aggregation functions and A is an attribute **NOT** among the grouping attributes

Grouping and Aggregation: Example

- Let's work with an example
 - Imagine that a social-networking site has a relation `Friends (User, Friend)`
 - The tuples are pairs (a, b) such that b is a friend of a
 - *Query: compute the number of friends each member has*
- $\gamma_{User, COUNT(Friend)}(Friends)$
 - This operation groups all the tuples by the value in their first component
 - There is one group for each user
 - Then, for each group, it counts the number of friends

Renaming Operator (ρ)

- Renames the list of attributes specified in the form of **oldname** \rightarrow **newname** or **position** \rightarrow **newname**
- **Output schema** is same as input except for the renamed attributes.
- Returns same tuples as input
- Can also be used to rename the name of the output relation

Boats

<u>bid</u>	bname	color
101	Interlake	blue
102	Interlake	red
103	Clipper	green
104	Marine	red

$\rho_{bname \rightarrow boatname, color \rightarrow boatcolor}(Boats)$

<u>bid</u>	boatname	boatcolor
101	Interlake	blue
102	Interlake	red
103	Clipper	green
104	Marine	red

$\rho_{2 \rightarrow boatname, 3 \rightarrow boatcolor}(Boats)$

Relational Algebra: Summary

Formal foundation for real query languages

- Helps represent and reason about execution plans

5 basic operators forming a closed algebra

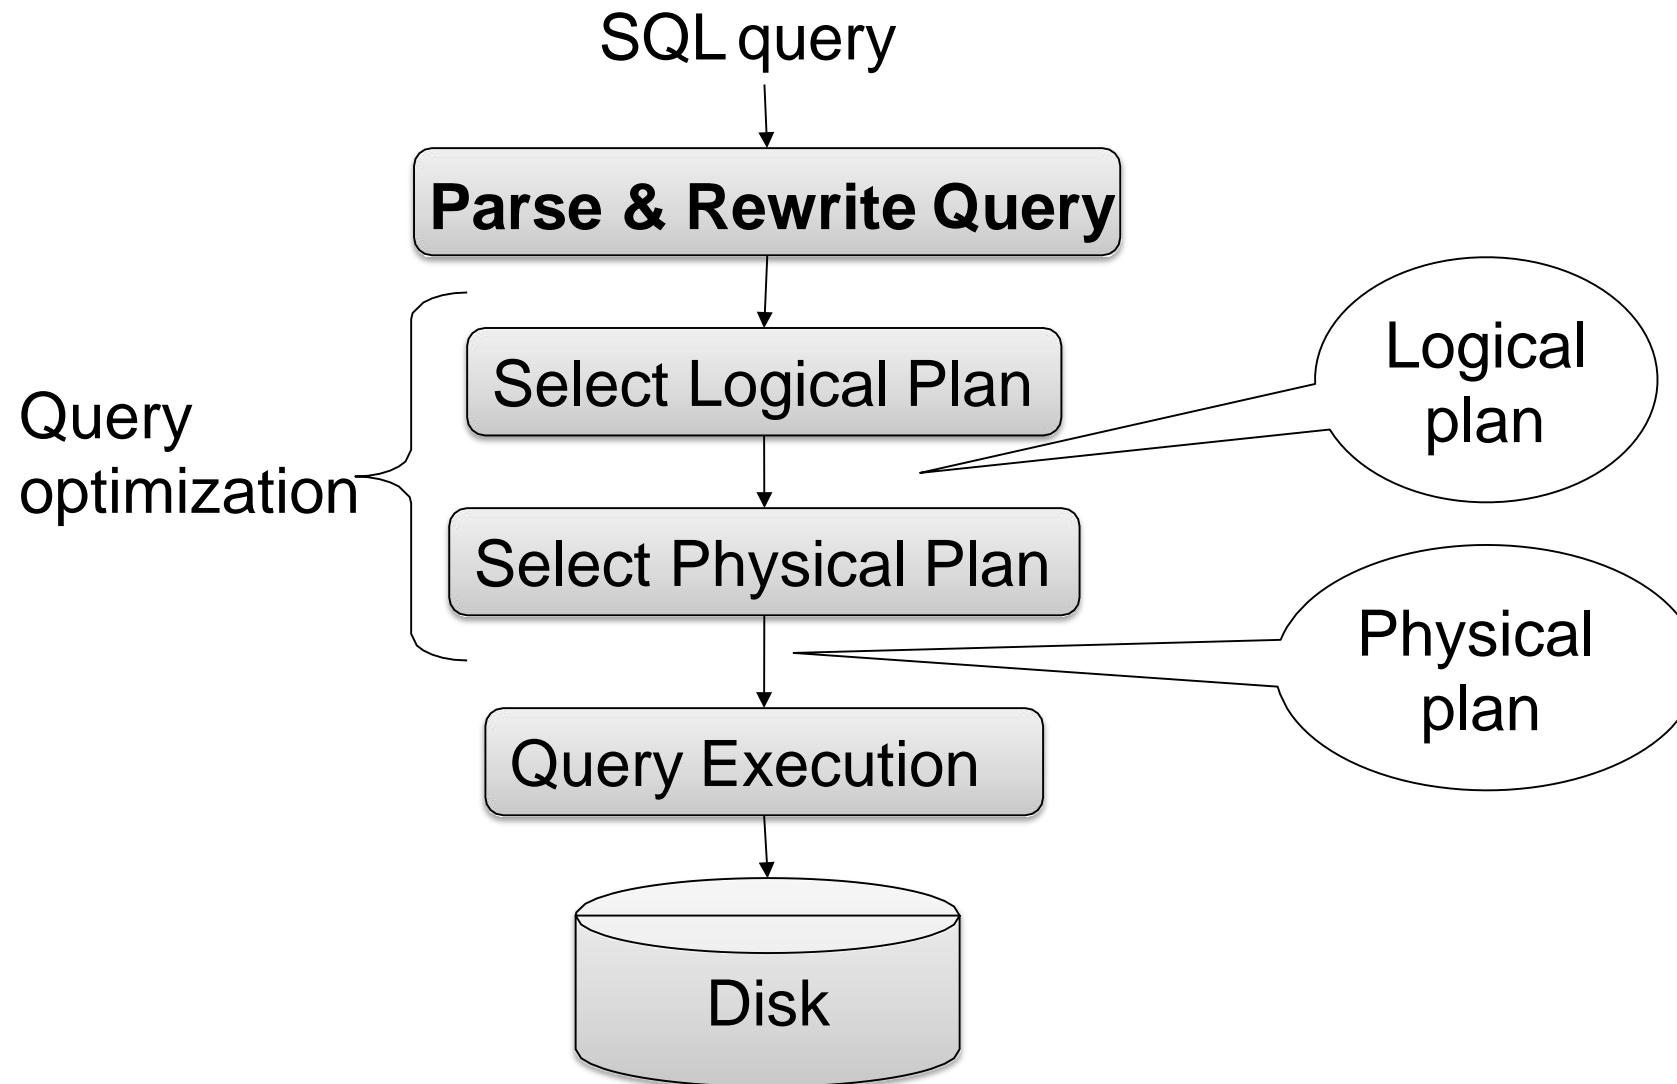
- Selection, projection, cross-product, union, set difference

Compound operators

- Useful shorthands like join and division
- Can be expressed with basic operators
- But enable faster query execution

Query Processing

Steps in Query Processing



Query parsing & transformation

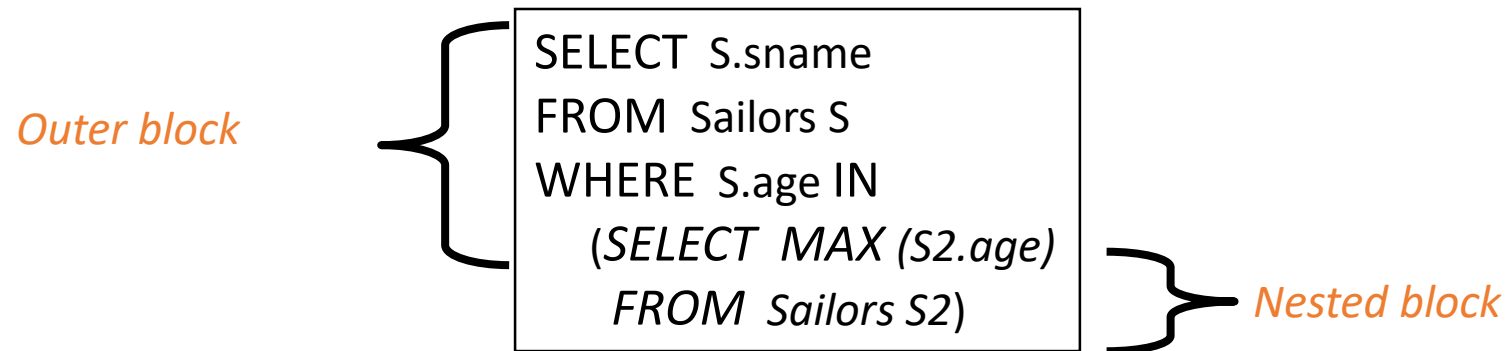
A Query:

```
SELECT S.sname  
FROM Reserves R, Sailors S  
WHERE R.sid=S.sid AND  
      R.bid=100 AND S.rating>5
```

1. Query first broken into “blocks”
2. Each block converted to relational algebra

Step 1: Break query into Query Blocks

- Query block = unit of optimization
- Nested blocks are usually treated as calls to a subroutine, made once per outer tuple
(This is an over-simplification, but serves for now)



Step 2: Converting query block into relational algebra expression

```
SELECT S.sid  
FROM Sailors S, Reserves R, Boats B  
WHERE S.sid = R.sid AND R.bid = B.bid AND B.color = "red"
```

$$\pi_{S.sid}(\sigma_{B.color = \text{"red"}}(Sailors \bowtie Reserves \bowtie Boats))$$

Relational Algebra Equivalences

- Selections: $\sigma_{c_1 \wedge \dots \wedge c_n}(R) \equiv \sigma_{c_1} \left(\dots \left(\sigma_{c_n}(R) \right) \right)$ (*Cascade*)
 $\sigma_{c_1} \left(\sigma_{c_2}(R) \right) \equiv \sigma_{c_2} \left(\sigma_{c_1}(R) \right)$ (*Commute*)
- Projections: $\pi_{a_1}(R) \equiv \pi_{a_1} \left(\dots \left(\pi_{a_n}(R) \right) \right)$ (*Cascade*)
 a_i is a set of attributes of R and $a_i \subseteq a_{i+1}$ for $i = 1 \dots n - 1$
- These equivalences allow us to ‘push’ selections and projections ahead of joins.

Another Equivalence

- A projection commutes with a selection that only uses attributes retained by the projection

$$\pi_{\text{age, rating, sid}} (\sigma_{\text{age} < 18 \wedge \text{rating} > 5} (\text{Sailors})) \\ \longleftrightarrow \sigma_{\text{age} < 18 \wedge \text{rating} > 5} (\pi_{\text{age, rating, sid}} (\text{Sailors}))$$

Equivalences Involving Joins

$$R \bowtie (S \bowtie T) \equiv (R \bowtie S) \bowtie T \quad (\textit{Associative})$$

$$(R \bowtie S) \equiv (S \bowtie R) \quad (\textit{Commutative})$$

- These equivalences allow us to choose different join orders

Examples ...

$$\sigma_{\text{age} < 18 \wedge \text{rating} > 5} (\text{Sailors})$$

$$\longleftrightarrow \sigma_{\text{age} < 18} (\sigma_{\text{rating} > 5} (\text{Sailors}))$$

$$\longleftrightarrow \sigma_{\text{rating} > 5} (\sigma_{\text{age} < 18} (\text{Sailors}))$$

~~$$\pi_{\text{age}, \text{rating}} (\text{Sailors}) \longleftrightarrow \pi_{\text{age}} (\pi_{\text{rating}} (\text{Sailors})) \quad (??)$$~~

$$\pi_{\text{age}, \text{rating}} (\text{Sailors}) \longleftrightarrow \pi_{\text{age}, \text{rating}} (\pi_{\text{age}, \text{rating}, \text{sid}} (\text{Sailors}))$$

Mixing Joins with Selections & Projections

- Converting selection + cross-product to join

$$\sigma_{S.sid = R.sid} (Sailors \times Reserves)$$

$$\leftrightarrow Sailors \bowtie_{S.sid = R.sid} Reserves$$

- Selection on just attributes of S commutes with $R \bowtie S$

$$\sigma_{S.age < 18} (Sailors \bowtie_{S.sid = R.sid} Reserves)$$


$$\leftrightarrow (\sigma_{S.age < 18} (Sailors)) \bowtie_{S.sid = R.sid} Reserves$$

- We can also “push down” projection (*but be careful...*)

$$\pi_{S.sname} (Sailors \bowtie_{S.sid = R.sid} Reserves)$$

$$\leftrightarrow \pi_{S.sname} (\pi_{sname, sid} (Sailors) \bowtie_{S.sid = R.sid} \pi_{sid} (Reserves))$$

Steps in Query Processing



We know...

Supplier(sno,sname,scity,sstate)

Part(pno,pname,psize,pcolor)


Supply(sno,pno,price)

For each SQL query....


```
SELECT S.sname
FROM Supplier S, Supply U
WHERE S.scity='Seattle' AND S.sstate='WA'
AND S.sno = U.sno
AND U.pno = 2
```

There exist many logical query plans...

Example Query: Logical Plan 1



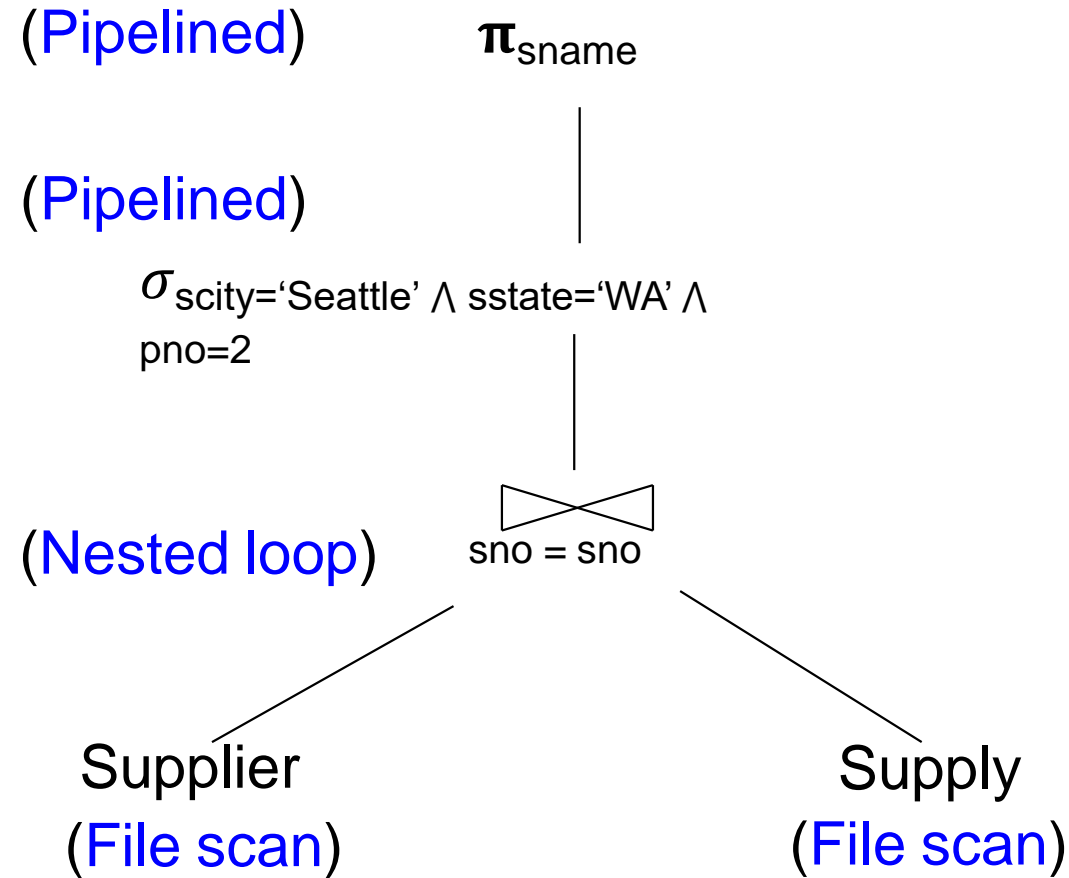
Example Query: Logical Plan 2



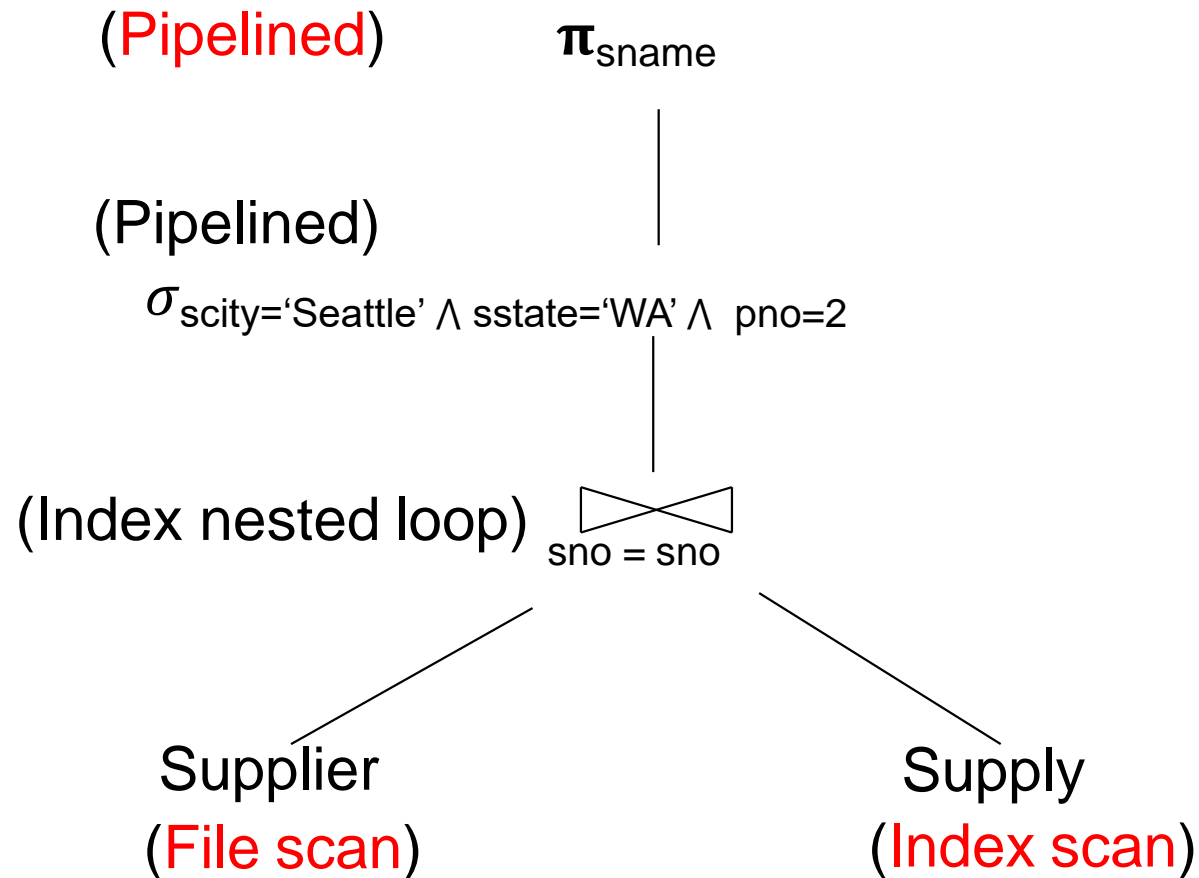
What We Also Know

- For each logical plan...
- There exist many physical plans

Example Query: Physical Plan 1



Example Query: Physical Plan 2




Query Optimization

1. Transformation produces relational algebra expression per “block”
2. Then, for each block, several alternative **query plans** are considered
3. Plan with lowest **estimated cost** is selected

```
SELECT S.sname  
FROM Reserves R, Sailors S  
WHERE R.sid=S.sid AND  
      R.bid=100 AND S.rating>5
```


$\pi_{(sname)} \sigma_{(bid=100 \wedge rating > 5)} (Reserves \bowtie Sailors)$



Query Optimizer Overview

- **Input:** A logical query plan
- **Output:** A good physical query plan
- **Basic query optimization algorithm**
 - Enumerate alternative plans (logical and physical)
 - Compute estimated cost of each plan
 - Compute number of I/Os
 - Optionally take into account other resources
 - Choose plan with lowest cost
 - This is called cost-based optimization

Steps in Query Processing



Query Execution Models

- A DBMS's processing model defines how the system executes a query plan.
 - Different trade-offs for workloads
- Approach #1: Interpreted execution
- Approach #2: Compiled execution


Interpreted execution with Volcano model

Each operator implements an iterator interface

- **open()**
 - Initializes operator state
 - Sets parameters such as selection condition
- **next()**
 - Operator invokes `get_next()` recursively on its inputs
 - Performs processing and produces an output tuple
- **close()**: clean-up state

Pipelined Execution


Tuples generated by an operator are immediately sent to the parent




Pipelined Execution

- Tuples generated by an operator are immediately sent to the parent
- Benefits:
 - Pull based: No operator synchronization issues
 - Saves cost of writing intermediate data to disk
 - Saves cost of reading intermediate data from disk
 - Enables implementation of parallelization as an operator


Exchange parallelization



```
SELECT A id, B val ue
FROM A JOIN B
ON A id = B id
WHERE A val ue < 99
AND B val ue > 100
```




Exchange parallelization




```


SELECT A id, B val ue
FROM A JOIN B
      ON A id = B id
WHERE A val ue < 99
      AND B val ue > 100
  
```




Exchange parallelization



```
SELECT A id, B val ue
FROM A JOIN B
ON A id = B id
WHERE A val ue < 99
AND B val ue > 100
```




Exchange parallelization




```


SELECT A id, B value
FROM A JOIN B
      ON A id = B id
WHERE A value < 99
      AND B value > 100
  
```




Exchange parallelization



```
SELECT A.id, B.value
FROM A JOIN B
ON A.id = B.id
WHERE A.value < 99
AND B.value > 100
```




Exchange parallelization




```

SELECT A.id, B.value
FROM A JOIN B
ON A.id = B.id
WHERE A.value < 99
AND B.value > 100
  
```




Exchange parallelization

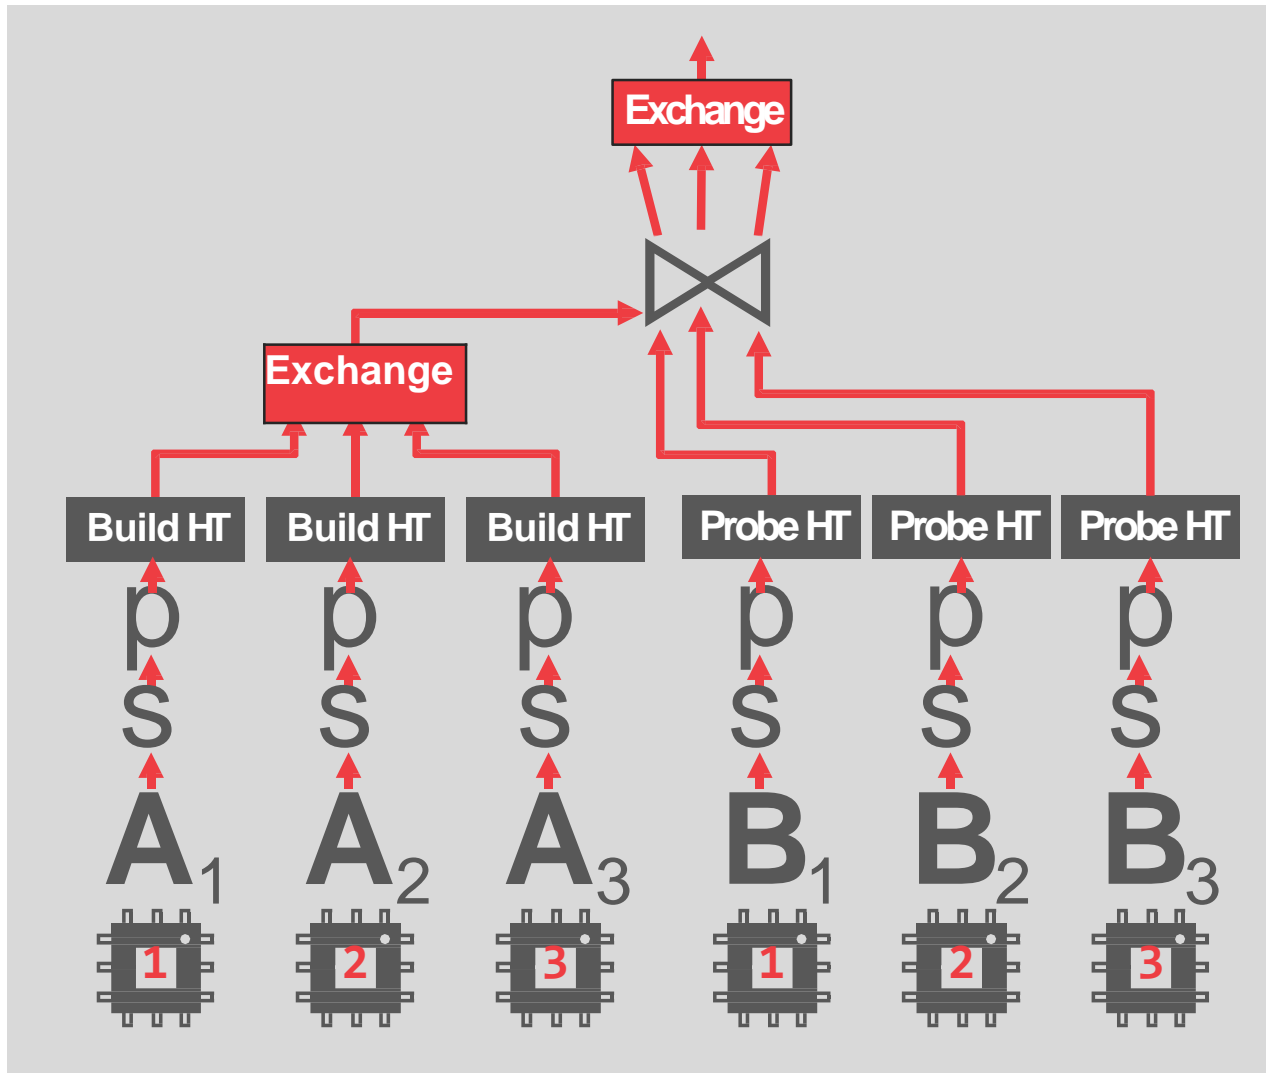


```

SELECT A id, B. value
FROM A JOIN B
ON A id = B id
WHERE A value < 99
AND B. value > 100
  
```




Exchange parallelization



```

SELECT A.id, B.value
FROM A JOIN B
ON A.id = B.id
WHERE A.value < 99
AND B.value > 100

```



Pipelined Execution

- Tuples generated by an operator are immediately sent to the parent
- Benefits:
 - Pull based: No operator synchronization issues
 - Saves cost of writing intermediate data to disk
 - Saves cost of reading intermediate data from disk
 - Enables implementation of parallelization as an operator
- Drawback
 - High function call overhead
 - Difficult to perform SIMD-style vectorization

Compiled execution

- Compile queries in-memory into native code
- Organizes query processing in a way to keep a tuple in CPU registers for as long as possible.
 - Push-based vs. Pull-based
 - Data Centric vs. Operator Centric
- LLVM typically used for compilation
 - Collection of modular and reusable compiler and toolchain technologies.
 - Core component is a low-level programming language (IR) that is like assembly.


Query interpretation vs compilation: Example

```
CREATE TABLE A (  
  id INT PRIMARY KEY,  
  val INT  
);
```

```
CREATE TABLE B (  
  id INT PRIMARY KEY,  
  val INT  
);
```

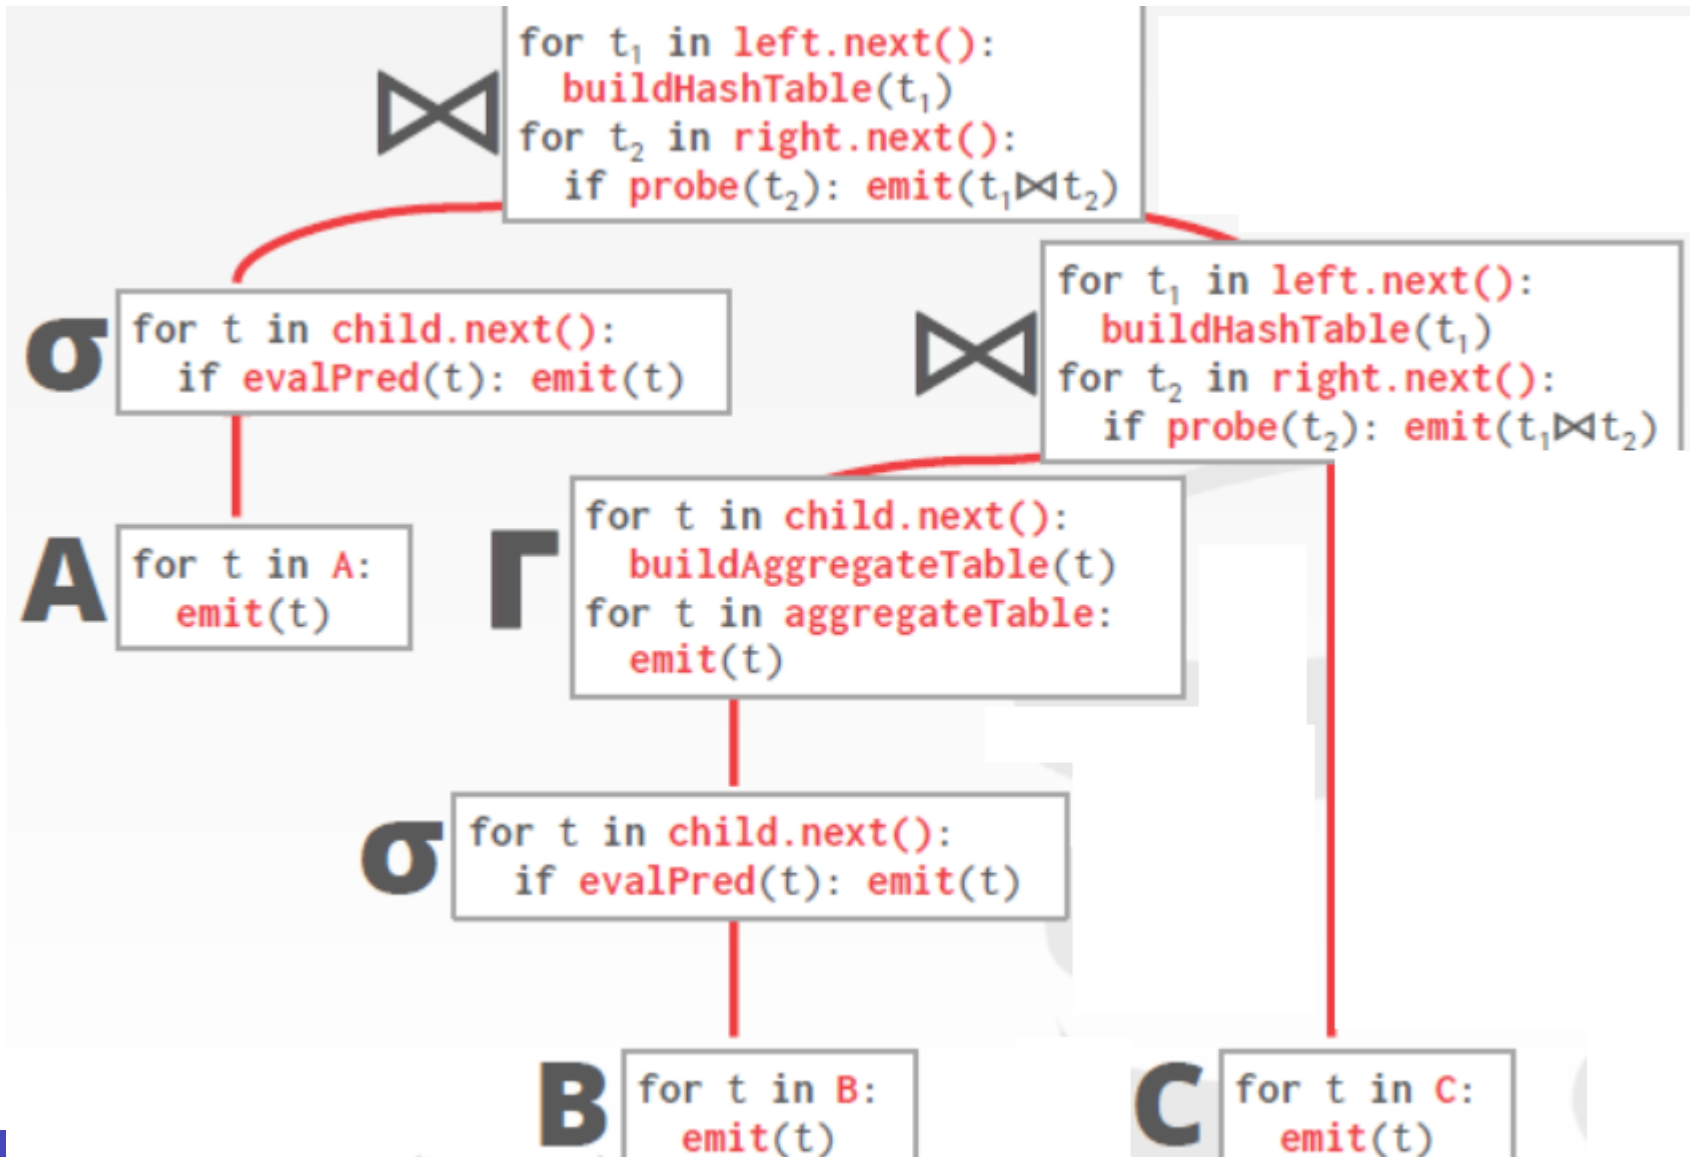
```
CREATE TABLE C (  
  a_id INT REFERENCES A(id),  
  b_id INT REFERENCES B(id),  
  PRIMARY KEY (a_id, b_id)  
);
```

```
SELECT *  
FROM A, C,  
  (SELECT B.id, COUNT(*)  
   FROM B  
   WHERE B.val = ? + 1  
   GROUP BY B.id) AS B  
WHERE A.val = 123  
AND A.id = C.a_id  
AND B.id = C.b_id
```




Query interpretation vs compilation: Example

```
SELECT *  
FROM A, C,  
      (SELECT B.id, COUNT(*)  
       FROM B  
       WHERE B.val = ? + 1  
       GROUP BY B.id) AS B  
WHERE A.val = 123  
      AND A.id = C.a_id  
      AND B.id = C.b_id
```



Query interpretation vs compilation: Example

```
SELECT *
FROM A, C,
  (SELECT B.id, COUNT(*)
   FROM B
   WHERE B.val = ? + 1
   GROUP BY B.id) AS B
WHERE A.val = 123
AND A.id = C.a_id
AND B.id = C.b_id
```




```
#1 { for t in A:
      if t.val == 123:
        Materialize t in HashTable  $\bowtie(A.id=C.a_id)$ 

#2 { for t in B:
      if t.val == <param> + 1:
        Aggregate t in HashTable  $\Gamma(B.id)$ 

#3 { for t in  $\Gamma(B.id)$ :
      Materialize t in HashTable  $\bowtie(B.id=C.b_id)$ 

#4 { for t3 in C:
      for t2 in  $\bowtie(B.id=C.b_id)$ :
        for t1 in  $\bowtie(A.id=C.a_id)$ :
          emit(t1  $\bowtie$  t2  $\bowtie$  t3)
```

On-disk Storage




DBMS on-disk storage

- The DBMS stores a database as one or more files on disk typically in a proprietary format.
 - The OS doesn't know anything about the contents of these files.
- DBMS storage manager is responsible for maintaining a database's files
 - It organizes the files as a collection of pages.
 - Tracks data read/written to pages.
 - Tracks the available space.

DBMS Page

- A page is a fixed-size block of data.
 - It can contain tuples, meta-data, indexes, log records...
- There are three different notions of "pages" in a DBMS:
 - Hardware Page (usually 4KB)
 - OS Page (usually 4KB)
 - Database Page (512B-16KB)
- We won't talk about page organization
 - Heap/tree/sorted/hashing files
 - Pages organization



Storage Models

- The relational model does not specify how a DBMS must store all a tuple in a page.
- Storage models
 - Dictate how tuples are stored within a page
 - Different models optimized for different workloads

Database Workloads

- **On-Line Transaction Processing (OLTP)**

- Fast operations that only read/update a small amount of data each time.

- **On-Line Analytical Processing (OLAP)**

- Complex queries that read a lot of data to compute aggregates.

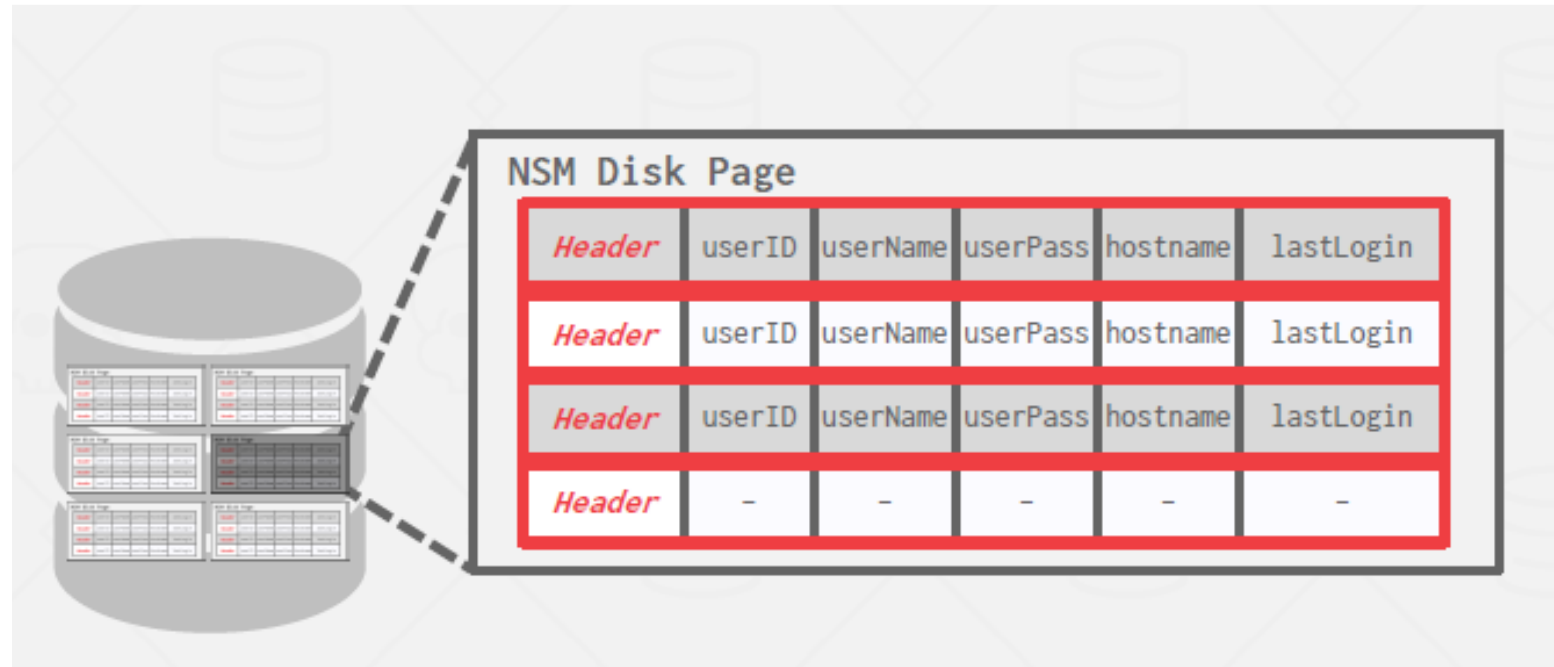
- **Hybrid Transaction + Analytical Processing**

- OLTP + OLAP together on the same database instance

N-ary storage model

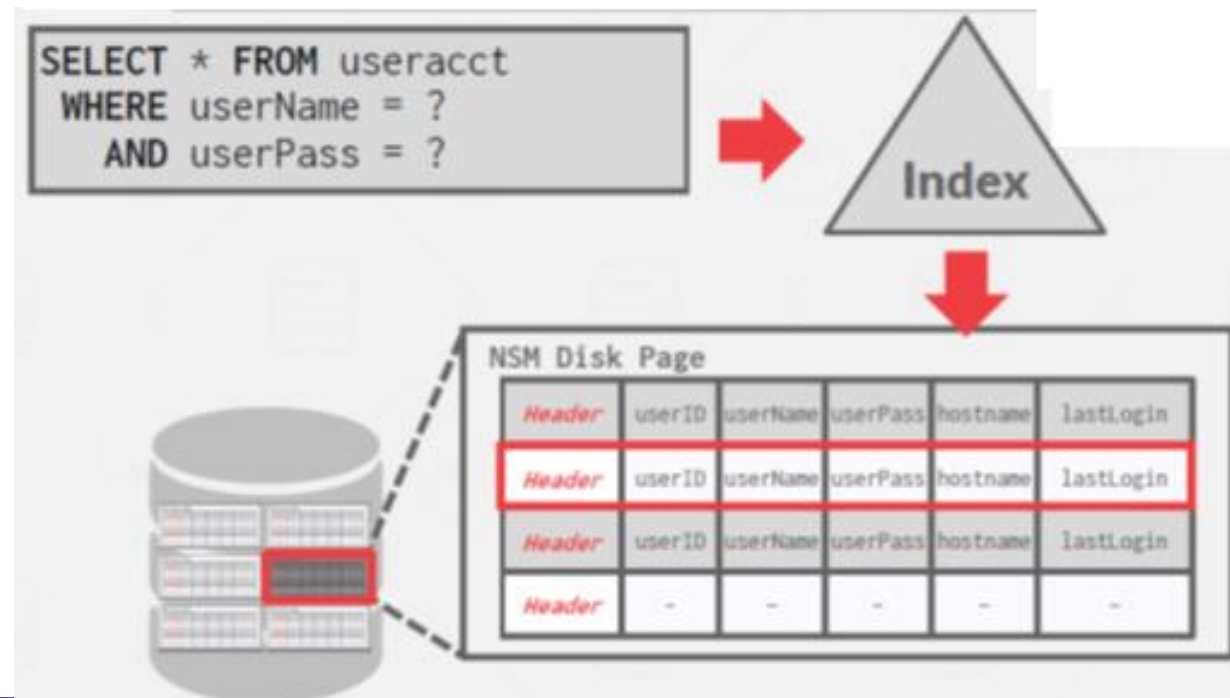
- DBMS stores all attributes for a single tuple contiguously in a page.

```
CREATE TABLE useracct (  
  userID INT PRIMARY KEY,  
  userName VARCHAR UNIQUE,  
  :  
);
```



N-ary storage model

- The DBMS stores all attributes for a single tuple contiguously in a page.
- Ideal for OLTP workloads where queries tend to operate only on an individual entity and insert heavy workloads.



N-ary storage model

- **Advantages**

- Fast inserts, updates, and deletes.
- Good for queries that need the entire tuple.

- **Disadvantages**

- Not good for scanning large portions of the table and/or a subset of the attributes.

```
SELECT COUNT(U.lastLogin),  
       EXTRACT(month FROM U.lastLogin) AS month  
FROM useracct AS U  
WHERE U.hostname LIKE '%.gov'  
GROUP BY EXTRACT(month FROM U.lastLogin)
```



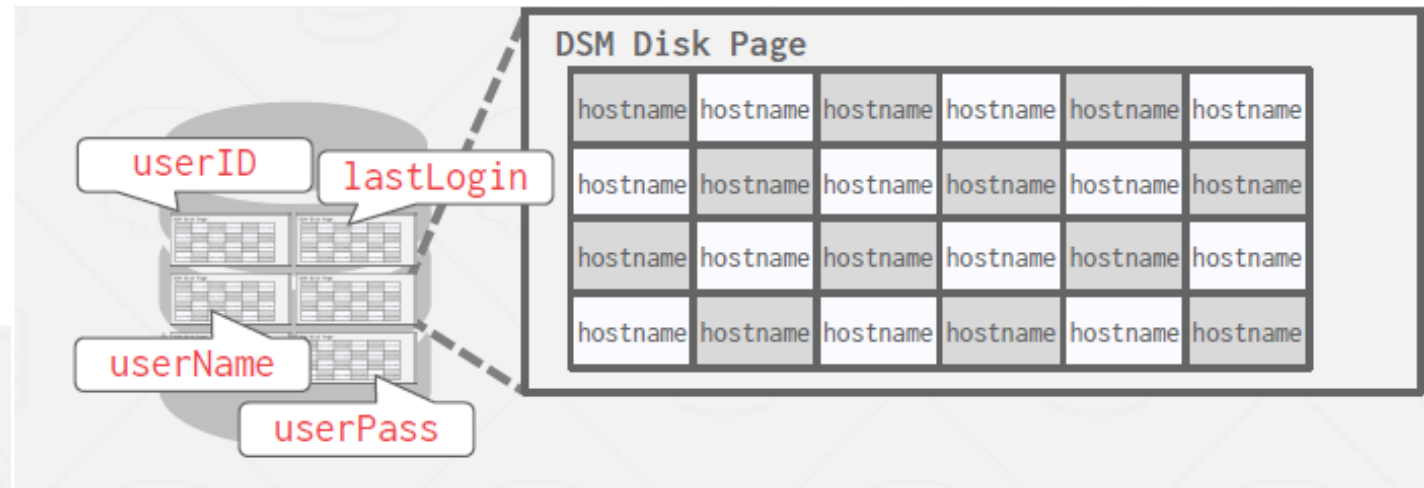
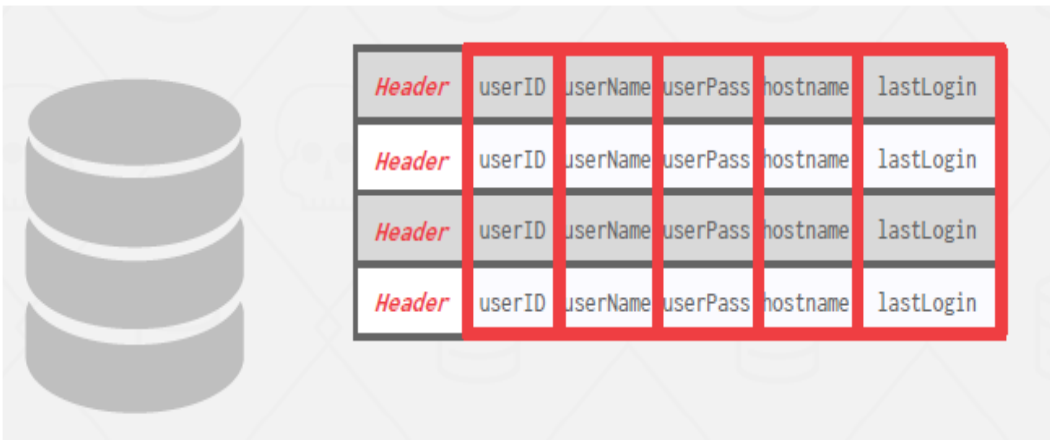
NSM Disk Page					
Header	userID	userName	userPass	hostname	lastLogin
Header	userID	userName	userPass	hostname	lastLogin
Header	userID	userName	userPass	hostname	lastLogin
Header	userID	userName	userPass	hostname	lastLogin

Useless Data

Decomposition storage model (DSM)

- The DBMS stores the values of a single attribute for all tuples contiguously in a page.
- Also known as a "column store"

```
CREATE TABLE useracct (  
  userID INT PRIMARY KEY,  
  userName VARCHAR UNIQUE,  
  :  
);
```



Decomposition storage model (DSM)

- Ideal for OLAP workloads where read-only queries perform large scans over a subset of the table's attributes.

```
SELECT COUNT(U.lastLogin),  
       EXTRACT(month FROM U.lastLogin) AS month  
FROM useracct AS U  
WHERE U.hostname LIKE '%.gov'  
GROUP BY EXTRACT(month FROM U.lastLogin)
```



DSM Disk Page

hostname	hostname	hostname	hostname	hostname	hostname
hostname	hostname	hostname	hostname	hostname	hostname
hostname	hostname	hostname	hostname	hostname	hostname
hostname	hostname	hostname	hostname	hostname	hostname

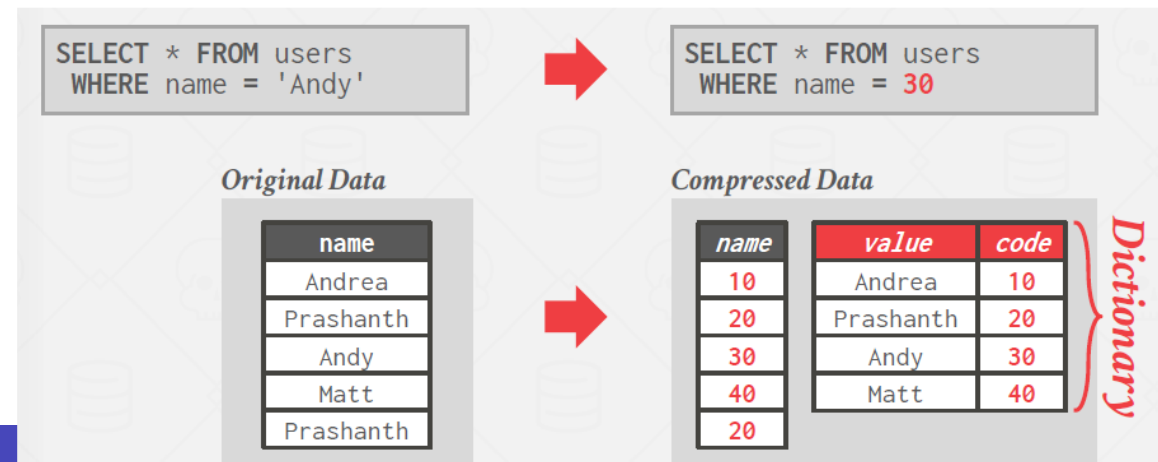
Decomposition storage model (DSM)

- **Advantages**

- Reduces the amount wasted I/O because the DBMS only reads the data that it needs.
- Better query processing and data compression

- **Dictionary compression example**

- Most used scheme in DBMSs enables queries over compressed data
- Build a data structure that maps variable-length values to integer identifier.
- Replace those values with their identifier in the dictionary data structure.



Storage Models: Summary

- N-ary storage model
 - **Advantages**
 - Fast inserts, updates, and deletes.
 - Good for queries that need the entire tuple.
 - **Disadvantages**
 - Not good for scanning large portions of the table and/or a subset of the attributes.
- Decomposition storage model
 - **Advantages**
 - Reduces the amount wasted I/O because the DBMS only reads the data that it needs.
 - Better query processing and data compression
 - **Disadvantages**
 - Slow for point queries, inserts, updates, and deletes because of tuple splitting/stitching.

It's all about them tables

- Relational databases
 - Relational model: Logical data independence
 - Relational algebra: Algebraic optimization, declarative querying
 - Optimized access paths: Indexing, materialized views, ...
 - Transactional semantics: ACID guarantees
- What is the link with MapReduce?

MR DBMS Comparison


“MapReduce and Parallel DBMS: A comparison of approaches to large-scale data analysis” – Andy Pavlo '09

- Tested Systems
 - Hadoop (MR)
 - Vertica (Columnar DBMS)
 - DBMS-X (Rowstore)


Benchmark 1: Grep Task

- Grep task: Shows overhead of data loading in DBMS
 - Search 3 byte pattern across 10Billion records
 - Each record: 100 bytes (10-byte key, 90-byte value)
 - 1TB across 25,50 or 100 nodes

DBMS slow during data Loading




DBMS fast during execution



Analytical Task

- Web processing: Shows the benefit of query optimization
 - 600k html documents
 - 155Million uservisits records, 18 million rankings records
 - Task: Find sourceIP that generated most revenue with avg. pagerank
 - DBMS: Complex SQL join query, MR: 3 separate MR programs



MapReduce: Not a silver bullet

- MapReduce does not fit many cases.
 - Interactive computations
 - E.g. not a user-facing web site back-end.
 - Small updates to big data (***more on OLTP and this later***)
 - Add a few documents to a big index
 - Small data, since overheads are high.


Relational operations over MR: Natural Join

- **Let's look at two relations $R(A, B)$ and $S(B, C)$**
 - We must find tuples that agree on their B components
 - We shall use the B-value of tuples from either relation as the key
 - The value will be the other component and the name of the relation
 - That way the reducer knows each tuple's relation
- **Map:**
 - For each tuple (a, b) of R emit the key/value pair $(b, ('R', a))$
 - For each tuple (b, c) of S emit the key/value pair $(b, ('S', c))$
- **Reduce:**
 - Each key b will be associated to a list of pairs that are either $('R', a)$ or $('S', c)$
 - Emit key/value pairs of the form $(b, [(a_1, b, c_1), (a_2, b, c_2), \dots, (a_n, b, c_n)])$

Shuffle overhead

Each Reducer:

- Handles $1/R$ of the possible key values
- Fetches its file from each of M mappers => **Shuffle**
- Sorts all of its entries to group values by keys




- Shuffle is an all-to-all communication that can overload the network
 - Paper's root switch: 100 to 200 gigabits/second
 - 1800 machines, so 55 megabits/second/machine.
 - Small, e.g. much less than disk (~ 50 - 100 MB/s at the time) or RAM speed.

Combiner optimization

- Often a map task will produce many pairs of the form (k, v_1) , (k, v_2) , ... for the same key k
 - E.g., popular words in Word Count
- Save network time by pre-aggregating at mapper with ***combiner function***
 - Decreases size of intermediate data transferred during shuffle
 - Reduce network load

Combiner example: Word count



Combiner: Algorithmic correctness

- The use of combiners must be thought carefully
 - the correctness of the algorithm cannot depend on computation (or even execution) of the combiners
- Commutative and Associative computations
 - Reducer and Combiner code may be interchangeable
 - This is not true in the general case
- Counter example: Mean
 - We have a large dataset where input keys are strings and input values are integers
 - Dataset can be a log from a website, where the keys are user IDs and values are some measure of activity
 - Compute the mean of all integers associated with the same key

Baseline approach for mean

- We use an **identity mapper**, which groups and sorts appropriately input key-value pairs
 - Reducers keep track of running sum and the number of integers encountered
 - The mean is emitted as the output of the reducer, with the input string as the key

```
1: class MAPPER
2:   method MAP(string  $t$ , integer  $r$ )
3:     EMIT(string  $t$ , integer  $r$ )

1: class REDUCER
2:   method REDUCE(string  $t$ , integers  $[r_1, r_2, \dots]$ )
3:      $sum \leftarrow 0$ 
4:      $cnt \leftarrow 0$ 
5:     for all integer  $r \in$  integers  $[r_1, r_2, \dots]$  do
6:        $sum \leftarrow sum + r$ 
7:        $cnt \leftarrow cnt + 1$ 
8:      $r_{avg} \leftarrow sum / cnt$ 
9:     EMIT(string  $t$ , integer  $r_{avg}$ )
```

Mean with combiners: Caution

- Note: operations are not distributive
 - $\text{Mean}(1,2,3,4,5) \neq \text{Mean}(\text{Mean}(1,2), \text{Mean}(3,4,5))$
 - Hence: a combiner cannot output partial means and hope that the reducer will compute the correct final mean
- Rule of thumb:
 - Combiners are optimizations, the algorithm should work even when “removing” them


Mean with combiners

```
1: class MAPPER
2:   method MAP(string t, integer r)
3:     EMIT(string t, pair (r, 1))
1: class COMBINER
2:   method COMBINE(string t, pairs [(s1, c1), (s2, c2) . . .])
3:     sum ← 0
4:     cnt ← 0
5:     for all pair (s, c) ∈ pairs [(s1, c1), (s2, c2) . . .] do
6:       sum ← sum + s
7:       cnt ← cnt + c
8:     EMIT(string t, pair (sum, cnt))
1: class REDUCER
2:   method REDUCE(string t, pairs [(s1, c1), (s2, c2) . . .])
3:     sum ← 0
4:     cnt ← 0
5:     for all pair (s, c) ∈ pairs [(s1, c1), (s2, c2) . . .] do
6:       sum ← sum + s
7:       cnt ← cnt + c
8:     ravg ← sum/cnt
9:     EMIT(string t, integer ravg)
```


MapReduce: Not a silver bullet

- MapReduce does not fit many cases.
 - Interactive computations
 - E.g. not a user-facing web site back-end.
 - Small updates to big data (more on OLTP and this later)
 - Add a few documents to a big index
 - Small data, since overheads are high.
 - **Unpredictable reads (neither Map nor Reduce can choose input) & load balancing issues**
 - Critical to scaling -- bad for N-1 servers to wait for 1 to finish.
 - Might be bug, flaky hardware, or poor partitioning
 - Leads to **Stragglers**: Tasks that take long time to execute

Better load balancing (1): Custom partitioning



- Mapper Operation
 - Reads input file blocks
 - Generates pairs $\langle K, V \rangle$
 - Writes to local file $h(K)$
- Hash Function h partitions intermediate key space K
 - Default h : Maps each key K to integer i such that $0 \leq i < R$
- Can also specify a customized partitioning function
 - Ex: output keys are URLs, we want all entries for a single host to end up in the same output file.
 - Can use “ $\text{hash}(\text{Hostname}(\text{urlkey})) \bmod R$ ” as h



Better load balancing (2): Task duplication

- Assign many more tasks than workers.
 - Master hands out new tasks to workers who finish previous tasks.
 - So no task is so big it dominates completion time (hopefully).
 - So faster servers do more work than slower ones, finish about the same time.

MapReduce: Not a silver bullet

- MapReduce does not fit many cases.
 - Interactive computations
 - E.g. not a user-facing web site back-end.
 - Small updates to big data (more on OLTP and this later)
 - Add a few documents to a big index
 - Small data, since overheads are high.
 - Unpredictable reads (neither Map nor Reduce can choose input) & load balancing issues
 - Critical to scaling -- bad for N-1 servers to wait for 1 to finish.
 - Might be bug, flaky hardware, or poor partitioning
 - Leads to **Stragglers**: Tasks that take long time to execute
 - **Iterative computations (*more on this in the next lecture*)**
 - **Algorithms with multiple rounds, e.g. k-means, page rank**

MapReduce: Not a silver bullet

- MapReduce does not fit many cases.
 - Interactive computations
 - E.g. not a user-facing web site back-end.
 - Small updates to big data (more on OLTP and this later)
 - Add a few documents to a big index
 - Small data, since overheads are high.
 - Unpredictable reads (neither Map nor Reduce can choose input) & load balancing issues
 - Critical to scaling -- bad for N-1 servers to wait for 1 to finish.
 - Might be bug, flaky hardware, or poor partitioning
 - Leads to **Stragglers**: Tasks that take long time to execute
 - Iterative computations (*more on this in the next lecture*)
 - Algorithms with multiple rounds, e.g. k-means, page rank
- **“MapReduce: A major step backwards” – Dewitt, Stonebraker**

RDBMS vs MapReduce: Summary

- Systems designed to meet different requirements
- Traditional relational databases
 - Interactive SQL analytics
 - Optimized for point queries (random access) and range queries (scans)
 - Built for enterprises (dedicated DB admin, few DB servers)
 - No need to scale to 1,000 or more nodes
 - Proprietary and paid products
 - Fine-grained updates to shared data
 - Guaranteeing ACID properties despite concurrent access and failures
- MapReduce
 - Latency-insensitive batch analytics
 - Sequential scans of Petabytes of data
 - Built for the cloud: Fault tolerance across commodity servers
 - Focus on faults during query rather than recovery after updates
 - Open source and “One person” deployment
 - Turn any Java developer into a distributed analytics engineer